

# 1 Connectionism

## 1.1 Perceptron

**Threshold Unit**  $f[w, b](x) = \text{sign}(x \cdot w + b)$  with **Decision Boundary**  $x \cdot w + b = 0$ .  $-\frac{b}{\|w\|}$  is the signed distance of the hyperplane from 0.

**Geometric Margin**  $\gamma[w, b](x, y) = \frac{y(x \cdot w + b)}{\|w\|}$ .

**Maximum Margin Classifier**

$(w^*, b^*) \in \arg\max_{w, b} \gamma[w, b](\mathcal{S})$ , with margin  $\gamma[w, b](\mathcal{S}) := \min_{(x, y) \in \mathcal{S}} \gamma[w, b](x, y)$ .

### Perception Learning

if  $f[w, b](x) \neq y$ :  $w \leftarrow w + yx, b \leftarrow b + y$ .

Aims to find some solution; assumes version space is non-empty. Does not aim for small err.  $w_0 \in \text{span}(x_1, \dots, x_s) \Rightarrow w_t \in \text{span}(x_1, \dots, x_s) \forall t$

### Convergence - Novikoff's Theorem

1. If  $\exists w^*, \|w^*\| = 1$ , s.t.  $\gamma[w^*](\mathcal{S}) = \gamma > 0 \Rightarrow w_t \cdot w^* \geq \gamma$ .

2. Let  $R = \max_{x \in \mathcal{S}} \|x\|$ . Then  $\|w_t\| \leq R\sqrt{\gamma}$ .  
 $\cos \angle(w^*, w_t) = \frac{w^* \cdot w_t}{\|w^*\| \|w_t\|} \geq \frac{\gamma}{\sqrt{\gamma}R} = \sqrt{\frac{\gamma}{R}} \leq 1 \Rightarrow t \leq \frac{R^2}{\gamma^2}$ .

### Cover's Theorem for $\mathcal{S} \subset \mathbb{R}^n, |\mathcal{S}| = s$

$C(\mathcal{S}, n)$ : # of ways to separate  $\mathcal{S}$  in  $n$  dimensions. Position of pts does not matter (as long as they are in general position).  $C(s+1, n) = 2 \sum_{i=0}^{n-1} \binom{s}{i}$ ,  $C(s, n) = 2^s$  for  $s \leq n$ . Phase transition at  $s = 2n$ . For  $s < 2n$  empty version space is the exception, otherwise the rule.

## 1.2 Hopfield Networks

$E(X) = -\frac{1}{2} \sum_{i \neq j} w_{ij} X_i X_j + \sum_i b_i X_i$ , where  $X_i \in \{\pm 1\}$ .  $w_{ij} = w_{ji}, w_{ii} = 0$ .

### Hebbian Learning

Choose patterns  $\{x^t\}_{t=1}^s \in \{\pm 1\}^n$ , compute weights once using them:  $w_{ij} = \frac{1}{n} \sum_{t=1}^s x_i^t x_j^t$ ;  $w_{ii} = 0$ . For inference, update  $X$  iteratively:  $X_i^{t+1} = \text{sign}\left(\sum_j w_{ij} X_j^t + b_i\right)$  asynchronously. Capacity for random, uncorrelated patterns:  $s_{\max} \approx 0.138n$ . Requiring pattern to be retrieved with high probability:  $s \leq \frac{n}{2 \log_2 n}$ .

If  $X = \text{diag}(1, \dots, 1)$ , no reconstruction happens. Under async update step, any Hopfield network is guaranteed to converge.

# 2 Feedforward Networks

## 2.1 Linear Models

### Linear regression (MSE)

$$L[w](X, y) = \frac{\|Xw - y\|^2}{2n}, \quad \nabla_w L = \frac{X^\top Xw - X^\top y}{n}.$$

### Moore-Penrose inverse solution

$w^* = X^*y \in \arg\min_w L[w](X, y)$ , where  $X^* = \lim_{\delta \rightarrow 0} (X^\top X + \delta I)^{-1} X^\top$  Moore-Penrose inverse.

**Stochastic gradient descent update**  $w_{t+1} = w_t + \eta(y_{i_t} - w_t^\top x_{i_t})x_{i_t}, i_t \sim \mathcal{U}([1, n])$ .

**Gaussian noise model**  $y_i = w^\top x_i + \varepsilon_i, \varepsilon_i \sim \mathcal{N}(0, \sigma^2)$ , LSQ equivalent to NLL of gaussian noise model.

**Ridge regression**  $h_\lambda[w] = h[w] + \frac{\lambda}{2} \|w\|^2, w^* = (X^\top X + \lambda I)^{-1} X^\top y$ .

### Logistic function

$$\sigma(z) = \frac{1}{1+e^{-z}}, \sigma(z) + \sigma(-z) = 1.$$

$$\sigma' = \sigma(1 - \sigma), \sigma'' = \sigma(1 - \sigma)(1 - 2\sigma)$$

### Cross entropy loss

$$\ell(y, z) = -y \log \sigma(z) - (1 - y) \log(1 - \sigma(z))$$

$$\ell(y, z) = -y \log \sigma(z) - (1 - y) \log(1 - \sigma(z))$$

### Logistic regression with Cross Entropy loss:

$$L[w] = \frac{1}{n} \sum_{i=1}^n \ell_i(y_i, w^\top x_i), \nabla \ell_i = [\sigma(w^\top x_i) - y_i] x_i.$$

## 2.2 Feedforward Networks

### Generic feedforward layers

$$F: \underbrace{\mathbb{R}^{m(n+1)}}_{\text{parameters}} \times \underbrace{\mathbb{R}^n}_{\text{input}} \rightarrow \underbrace{\mathbb{R}^m}_{\text{output}}, F[\theta](x) = \varphi(Wx + b).$$

### Layer composition

$$G = F^L[\theta^L] \circ \dots \circ F^1[\theta^1].$$

### Layer activations

$$x^l = F^l[\theta^l] \circ \dots \circ F^1[\theta^1](x) = F^l(x^{l-1}), x^0 = x, x^L = F(x).$$

**Softmax** ( $z$ ) $_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$ , CE-loss in terms of logits  $\ell(y, z) = \frac{1}{\ln(2)} \left[ -y_z + \ln \sum_j e^{z_j} \right]$ . CE between two pmfs:  $l(p; q) = -\sum_i p_i \log q_i$ . CE with hard labels is NLL-loss.

**Residual layer**  $F[W, b](x) = x + (\varphi(Wx + b) - \varphi(0))$ , therefore  $F[0, 0] = \text{id}$ . Link that propagates  $x$  forward is called a **skip connection**. Composing residual layers: number of paths grows exponentially, can include projections for flexibility of changing dimensionality.

## 2.3 Sigmoid Networks

**Sigmoid activation**  $\sigma(z) = \frac{1}{1+e^{-z}}$ .

**Hyperbolic tangent activation**

$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}} = 2\sigma(2z) - 1.$$

$$\tanh'(z) = 1 - \tanh^2(z).$$

### Smooth function approximation

Polynomials, ridge functions ( $\varphi(a^\top x + b)$ ) and MLPs with  $C^\infty$  activations are universal approximators.

**Weierstrass**: Polynomials are universal approximators of  $C(\mathbb{R})$  on any given compact  $I$ .

**Barron's Theorem**: For  $f: \mathbb{R}^d \rightarrow \mathbb{R}$  with  $C_f = \int \|\omega\| |\hat{f}(\omega)| d\omega < \infty$ ,  $\exists$  width- $m$  MLP  $g_m$  s.t.:  $\int_B |f - g_m|^2 dx \leq O(\frac{1}{m})$ .

## 2.4 ReLU( $z$ ) = $\max(0, z)$ networks

### Zalavsky's Thoerem: Activation patterns

$m$  ReLU neurons in  $\mathbb{R}^n$ . Each neuron's hyperplane  $\{w_i^\top x = 0\}$  partitions  $\mathbb{R}^n$  into  $R(m)$  connected regions of constant activation pattern.  $R(m) \leq \sum_{i=0}^{\min(n, m)} \binom{m}{i} \ll 2^m$ .

**Montufar: Connected linear regions in ReLU network**  $R(m, L) \geq R(m) \lfloor \frac{m}{n} \rfloor^{n(L-1)}$ ,  $L$ : layers,  $m$ : width.

**Shekhtman**: Piecewise linear functions are dense

in  $C([0, 1])$ . **Lebesgue**: Piecewise linear function with  $m$  pieces can be written  $g(x) = ax + b + \sum_{i=1}^{m-1} c_i (x - x_i)_+$ ,  $m+1$  parameters,  $a, b, c_i$ .

### ReLU networks with 1 hidden layer are universal approximators.

**Wang and Sun**: Every continuous piecewise linear function  $g: \mathbb{R}^n \rightarrow \mathbb{R}$  can be written as a signed sum of  $k$ -Hinges with  $k \leq n+1$ . A  $k$ -Hinge is a function  $g(x) = \max_{j=1}^k \{w_j^\top x + b_j\}$ , generalizes ReLU, known as Maxout unit.

**Linear Autoencoder**: Optimal  $A = DE$ , s.t. frobenius norm reconstruction err of  $AX$  is minimized, is  $D = U_k, E = U_k^\top$ , not jointly convex in  $E$  and  $D$ , but individually.  $\hat{X}^* = \arg\min_{\text{rank } (\hat{X})=k} \|X - \hat{X}\|_F^2 = U \Sigma_k V^\top$  SVD.

Implies  $\lambda_i(\nabla^2 f(x)) \geq \mu$  for all EVs  $\lambda_i$  of  $\nabla^2 f(x)$ .

For  $f: \mathbb{R} \rightarrow \mathbb{R}$ , these become:  $L \geq f''(x) \geq \mu \quad \forall x$ .

If  $f$   $\mu$ -strongly convex,  $L$ -smooth, GD iterates  $x_t$  with  $0 < \eta \leq \frac{1}{L}$  converge to unique minimizer  $x^*$  at rate  $\|x_t - x^*\|^2 \leq (1 - \eta\mu)^k \|x_0 - x^*\|^2$ . If  $f$  convex, diff and  $L$ -smooth, with  $\eta \leq \frac{1}{L}$ ,  $f(x_t) - f(x^*) \leq \frac{1}{2\eta t} \|x_0 - x^*\|^2$ . **Non-Convex case**: If  $f$  diff,  $L$ -smooth, with minimum  $f^*$ , GD iterates with  $\eta \leq \frac{1}{L}$  satisfy  $\min_{i=0}^t \|\nabla f(x_i)\|^2 \leq \frac{2(f(x_0) - f^*)}{\eta(t+1)}$ .

If  $f$  diff and  $L$ -smooth:  $f(x) - f(x^*) \geq \frac{1}{2L} \|\nabla f(x)\|^2$ .

### Polyak-Lojasiewicz condition

$$\frac{1}{2} \|\nabla f(x)\|^2 \geq \mu(f(x) - \min f) \text{ (forall } x).$$

$\mu$ -strong convex  $\Rightarrow \mu$ -PL.

### GD Convergence Rates & Learning Rates

**L-smooth only**:  $\eta^* = \frac{1}{L}$ . To reach  $\varepsilon$ -stationary point ( $\|\nabla f\| \leq \varepsilon$ ) needs at most  $\frac{2L}{\varepsilon^2}(f(x_0) - \min f)$  steps.

**$\mu$ -PL + L-smooth**: Use  $\eta^* = \frac{2}{L+\mu}$ . Convergence:  $f(x_t) - f(x^*) \leq (1 - \frac{\mu}{L})^t (f(x_0) - f(x^*))$ .

## 3 Gradient-Based Learning

Forward mode is more memory efficient, but backward mode is more runtime efficient. Fwd is  $O(\# \text{params})$ , reverse is  $O(d_{\text{out}})$ .

**Numerator layout**:

For  $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$ ,

$$(\frac{\partial y}{\partial x})_{ij} = \frac{\partial y_i}{\partial x_j} \in \mathbb{R}^{m \times n}$$

and  $f: \mathbb{R}^{n_1 \times n_2} \rightarrow \mathbb{R}$ ,

$$\nabla f(X)_{ij} = \frac{\partial f}{\partial X_{ij}} \in \mathbb{R}^{n_1 \times n_2}$$

CE between two pmfs:  $l(p; q) = -\sum_i p_i \log q_i$ .

**3.1 Backpropagation**

$x^\ell = \varphi(W^\ell x^{\ell-1} + b^\ell)$ ,

$$\frac{\partial \mathcal{L}}{\partial b^\ell} = \delta^\ell, \delta^\ell = \frac{\partial \mathcal{L}}{\partial x^\ell} \odot \varphi'(W^\ell x^{\ell-1} + b^\ell), \frac{\partial \mathcal{L}}{\partial x^\ell} = (W^{\ell+1})^\top \delta^{\ell+1}$$

$\frac{\partial L}{\partial x} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial x}$ . For  $x \in \mathbb{R}^n$  different  $z$ :  $\frac{\partial (Wx)}{\partial x} = W$ , element wise  $f$  gives:  $\frac{\partial f(x)}{\partial x} = \text{diag}(f'(x))$ ,  $\frac{\partial \hat{y} - y}{\partial y} = 2(\hat{y} - y)^\top, \frac{\partial L}{\partial y} = h \cdot \frac{\partial L}{\partial \hat{y}}$ .

$$\frac{d}{dx_j} \text{softmax}(x)_i = \text{sm}(x)_i (\delta_{ij} - \text{sm}(x)_j)$$

**3.2 Gradient Descent**

**Update**:  $x_{t+1} = x_t - \eta \nabla f(x_t)$ .

**Gradient flow ODE**  $\frac{dx}{dt} = -\nabla f(x)$  gives ideal trajectory to be approximated by gradient descent.

**Newton's method** gives optimal step for quadratic model:  $\Delta x = -[\nabla^2 f(x)]^{-1} \nabla f(x)$ .

$$\nabla_x^2 [x^\top Ax + b^\top x + c] = A + A^\top$$

**Minibatch SGD**: Variance  $\downarrow$  by  $\propto r$ . Can  $\uparrow \eta \propto r$ .

**Var. Reduction with SVRG** w/ occasional snapshot  $\bar{\theta}$ :  $\theta_{t+1} = \theta_t - \eta [\nabla f_i(\theta_t) - \nabla f_i(\bar{\theta}) + \nabla f(\bar{\theta})]$ .

## 3.4 Acceleration and Adaptivity

### Heavy ball momentum update

$$\theta_{t+1} = \theta_t - \eta \nabla f(\theta_t) + \beta(\theta_t - \theta_{t-1})$$

### Nesterov acceleration

$$\tilde{\theta}_{t+1} = \theta_t + \beta(\theta_t - \theta_{t-1})$$

$$\theta_{t+1} = \tilde{\theta}_{t+1} - \eta \nabla h(\tilde{\theta}_{t+1})$$

More theoretical grounding than heavy ball.

### AdaGrad updates

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\gamma_t + \varepsilon}} \odot \nabla f(\theta_t),$$

$$\gamma^t = \gamma^{t-1} + \nabla f(\theta_t) \odot \nabla f(\theta_t)$$

### Adam updates

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) \nabla f(\theta_t), \hat{m}_t = \frac{m_t}{1 - \beta_1^t}$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) (\nabla f(\theta_t))^2, \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

$$\theta_{t+1} = \theta_t - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t + \varepsilon}}$$

**RMSprop**: Adam without momentum term.

**signSGD**:  $\theta_{t+1} = \theta_t - \eta \text{ sign}(\nabla f_{I_k}(\theta_t))$ .

$x^4$  is strictly convex but not strongly convex, since near 0 the growth of  $x^4$  is slower than  $x^2$ , violating the uniform lower bound on curvature. With  $f$   $L$ -smooth and  $\mu$ -PL, GD with optimal step-size  $\arg\min_\eta f(\theta_t - \eta \nabla f(\theta_t))$  converges globally at linear rate.  $f(w) = \frac{\|w-y\|^2}{2} + \lambda \|w\|^2$  satisfies PL-condition.

**Muon**: Orthogonalize gradient, should increase the scale of other rare directions which have small magnitude in update but are important.

$\Delta W = -\|\nabla L(W)\|_* \cdot \frac{d_{\text{out}}}{d_W} \cdot U^\top V$  ( $\nabla L(W) = U \Sigma V^\top$ ) minimizes RHS of:  $L(W + \Delta W) \leq L(W) + \langle \nabla W L(W), \Delta W \rangle_F + \frac{1}{2} \frac{d_{\text{out}}}{d_W} \|\Delta W\|_F^2$ .

GD Trajectory always orthogonal to level set.

# 4 Convolutional Networks

Convolution ( $f * g)(u) = (g * f)(u)$

$$= \int_{-\infty}^{\infty} f(t) g(u-t) dt = \int_{-\infty}^{\infty} f(u-t) g(t) dt.$$

$(f * g)$  = Toeplitz-Matrix( $g$ ) $f$ .

### Fourier transform convolution property

$$\mathcal{F}(f * g) = \mathcal{F}(f) \cdot \mathcal{F}(g).$$

### Cross-correlation:

$$(g * f)[u] = \sum_t g[t] f[u+t].$$

$$(f(t) * g(t))(y) = (f(-t) * g(t))(y).$$

$$(f(t) * g(t))(-y) = (f(-t) * g(t))(y) = (f(t) * g(-t))(y).$$

Translation-Equivariant Operators = Convolutions.

For  $x \in \mathbb{R}^D, h \in \mathbb{R}^K$ ,  $x * h = W_h x$ . **Toeplitz Matrix**:  $W_h \in \mathbb{R}^{(D-K+1) \times D}$

$$(W_h)_{k,j} = \begin{cases} h_{k-j+k-1} & \text{if } k \leq j \leq k+K-1 \\ 0 & \text{otherwise} \end{cases} \text{ for } h \in \mathbb{R}^K.$$

$\nabla_w v^\top \text{vec}(\sigma(x * w)) = \text{flip}[x] * [\text{mat}(v) \odot \sigma'(x * w)]$ , flip rows and columns.

Normal convolution of image  $x \in \mathbb{R}^{d \times d}$ , kernel  $w \in \mathbb{R}^{q \times q} x * w$  requires  $\mathcal{O}((d-q)^2 q^2)$ . If  $w$  separable st  $w = uv^\top, u \in \mathbb{R}^q, v \in \mathbb{R}^p, \mathcal{O}(dq(d-q))$

Output size:  $H_{\text{out}} = \lfloor \frac{H_{\text{in}}+2P-K}{S} \rfloor + 1$ , Height, Input size, padding, Kernel size, stride.

## 4.1 Convolutional Networks

Conventions for Padding: Add zeros around input.

### ConvNets for Images (r out channel, u in channel)

$$y[r][s, t] = \sum_u \sum_{i,j} w[r, u][i, j] * x[u][s+i, t+j].$$

### Number of parameters of a convolutional layer

$$(|r| \times |u|) \cdot (|i| \times |j|) : \text{fully connected} \times \text{patch-size}.$$

## 4.2 Word2Vec

Per word  $\omega$ , have input embedding  $x_\omega$  and output embedding  $y_\omega$ .

Predict context word  $\nu$  given center word  $\omega$ :  
 $P(\nu | \omega) = \frac{\exp(x_\omega^\top y_\nu)}{\sum_u \exp(x_\omega^\top y_u)}$ .

NLL loss:  $\ell_{\omega,\nu} = -x_\omega^\top y_\nu + \ln \sum_\mu \exp(x_\omega^\top y_\mu)$ .  
 Total:  $h(\{x_\omega\}, \{y_\nu\}) = \sum_{(\omega, \nu)} \ell_{\omega,\nu}$  over observed pairs. Use only input embeddings after training.

## 5 Geometric Deep Learning

**Group** is set  $G$  with a binary operation s.t.:  
 1)  $(gh)f = g(hf)$ , 2)  $\exists e \in G$  s.t.  $ef = fe = f$ ,  
 3)  $\forall g \exists g^{-1} \in G$  s.t.  $gg^{-1} = g^{-1}g = e$ , 4)  $gh \in G \forall g, h$ . Abelian if  $gh = hg$ .

### 5.1 Sets and Points

**Order-invariance property:**

$$f(x_1, \dots, x_M) = f(x_{\pi(1)}, \dots, x_{\pi(M)}) \quad (\text{perturbations}).$$

**(Permutation) Equivariance property:**

$$\begin{aligned} f(x_1, \dots, x_M) &= (y_1, \dots, y_M) \\ f(x_{\pi(1)}, \dots, x_{\pi(M)}) &= (y_{\pi(1)}, \dots, y_{\pi(M)}) \end{aligned}$$

**Deep Sets model** (invariant layer):  
 $f(x_1, \dots, x_M) = \rho(\sum_{m=1}^M \varphi(x_m))$ .

**Equivariant map construction:**

$$\rho : \mathbb{R} \times \mathbb{R}^N \rightarrow Y, (x_m, \sum_{k=1}^M \varphi(x_k)) \mapsto y_m$$

### 5.2 Graph Convolutional Networks

**Feature and adjacency matrices**

$$X = \text{mat}(x_1^\top; \dots; x_M^\top), A = (a_{nm}) \text{ with } a_{nm} = 1 \Leftrightarrow \{v_n, v_m\} \in E.$$

**Permutation matrix constraints**

$$P \in \{0, 1\}^{M \times M} \text{ with single 1 in each row and col.}$$

**Graph invariance definition**

$$f(X, A) \stackrel{!}{=} f(PX, PAP^\top), \forall P \in \Pi_M.$$

**Graph equivariance definition**

$$f(X, A) = Pf(PX, PAP^\top), \forall P \in \Pi_M.$$

**Node neighborhood features**

$$X_m = \{\{x_n : \{v_n, v_m\} \in E\}\}, \quad \{\cdot\} = \text{multiset}$$

**Message passing scheme**

$$\varphi(x_m, X_m) = \varphi(x_m, \bigoplus_{x \in X_m} \psi(x)), \quad \oplus \text{ is some permutation-invariant operation.}$$

**Normalized adjacency matrix**

$$\bar{A} = D^{-\frac{1}{2}}(A + I)D^{-\frac{1}{2}}, \quad D = \text{diag}(d_1, \dots, d_M), d_m = 1 + \sum_{n=1}^M a_{nm}.$$

**One GCN layer**

$$X^+ = \sigma(\bar{A}XW), W \in \mathbb{R}^{M \times N}.$$

### 5.2.1 Spectral Graph Theory

**Laplacian operator**

$$\Delta f = \sum_{n=1}^N \frac{\partial^2 f}{\partial x_n^2}, f : \mathbb{R}^N \rightarrow \mathbb{R}.$$

**Graph Laplacian**

$$L = D - A, (Lx)_n = \sum_{m=1}^M a_{nm}(x_n - x_m). \quad x^\top Lx = \frac{1}{2} \sum_u \sum_v A_{uv}(x_u - x_v)^2 \geq 0 \text{ (psd).}$$

**Normalized Laplacian**

$$\tilde{L} = I - D^{-\frac{1}{2}}AD^{-\frac{1}{2}} = D^{-\frac{1}{2}}(D - A)D^{-\frac{1}{2}}.$$

**Graph Fourier transform**

$$L = D - A = U\Lambda U^\top, \hat{f} = U^\top f, f = U\hat{f}. \quad \Lambda := \text{diag}(\lambda_1, \dots, \lambda_M), \lambda_i \geq \lambda_{i+1}.$$

**Convolution:**  $x * y = U((U^\top x) \odot (U^\top y))$ .

**Filtering operation**

$$G_\theta(L)x = UG_\theta(\Lambda)U^\top x$$

**Polynomial kernels**

$$U(\sum_{k=0}^K \alpha_k \Lambda^k)U^\top = \sum_{k=0}^K \alpha_k L^k$$

**Polynomial kernel network layer**

$$\begin{aligned} x_i^{l+1} &= \sum_j p_{ij}(L) x_j^l + b_i, \\ p_{ij}(L) &= \sum_{k=0}^K \alpha_{ijk} L^k \end{aligned}$$

GNNs cannot distinguish between certain graphs that are topologically different. Unconstrained set architectures are more powerful. If **WL-test** says graphs are different, they are; but if it says they're the same, they might still be different.

## 6 Theory of DNNs

### 6.1 Statistical Learning Theory

**Risk decomposition**

$f^*$ : optimal predictor over all functions,

$\hat{f}_H$ :  $\underset{f \in H}{\text{argmin}} \mathcal{R}(f)$ ,  $\hat{f}_H$ : learned from finite data.

$$\mathcal{R}(\hat{f}_H) - \mathcal{R}(f^*) = \underbrace{\mathcal{R}(\hat{f}_H) - \mathcal{R}(f_H)}_{\text{excess risk}} + \underbrace{\mathcal{R}(f_H) - \mathcal{R}(f^*)}_{\text{estimation error}} + \underbrace{\mathcal{R}(f^*) - \mathcal{R}(\hat{f}_H)}_{\text{approximation error}}$$

**Rademacher Complexity**  $G = \{g_h \mid h \in H\}$ :

For  $\sigma \in \{-1, 1\}$ , measures how well  $G$  can fit random noise:  $\hat{\mathcal{R}}_{D_n}(G) = \mathbb{E}_\sigma [\sup_{g \in G} \frac{1}{n} \sum_{i=1}^n \sigma_i g(z_i)]$ .

$$\mathbb{E}[\sup_{h \in H} \mathcal{R}(h) - \hat{\mathcal{R}}_{D_n}(h)] \leq 2\hat{\mathcal{R}}_{D_n}(G), \quad \mathbb{E}[\mathcal{R}(\hat{h}_H)] \leq \mathcal{R}(h^*) + 2\hat{\mathcal{R}}_{D_n}(G).$$

**Double descent:**

Beyond the interpolation point, models eventually may level out at a lower generalization error.

**Implicit bias towards min norm solutions:**

Any convergent algorithm with iterates in  $\text{span}\{x_1, \dots, x_n\}$  finds the minimum norm solution.

#### 6.1.1 A PAC-Bayesian result

$P$  prior distribution over functions before seeing data,  $Q$  posterior after training.

**PAC-Bayesian theorem**

Bounds generalization gap for stochastic

classifiers ( $f \sim Q$ ):  $E_Q[\mathcal{R}(f)] - E_Q[\hat{\mathcal{R}}_n(f)] \leq \sqrt{\frac{2}{n} [\text{KL}(Q||P) + \ln(2\sqrt{n}/\varepsilon)]}$

•  $P$ : prior,  $Q$ : posterior (learned). Rate  $\tilde{O}(1/\sqrt{n})$

•  $\text{KL}(Q||P)$ : “information cost” of moving  $P \rightarrow Q$

• Insight: generalization depends on **distance moved**, not parameter count

**PAC-Bayesian for DNNs**

$$P = \mathcal{N}(0, \lambda I), Q = \mathcal{N}(\theta, \text{diag}(\sigma_i^2)), \quad \text{KL}(Q||P) = \sum_i \left[ \log \frac{\lambda}{\sigma_i} + \frac{\sigma_i^2 + \theta_i^2 - 1}{2\lambda^2} - \frac{1}{2} \right]$$

Minimize directly:  $E_Q[\hat{\mathcal{R}}] + \sqrt{\frac{2}{n} [\text{KL}(Q||P) + \dots]}$

$\Rightarrow$  encourages wide/flat minima (perturbations  $\theta + \varepsilon$  must also perform well)

**Implementation:** Reparameterization:  $\tilde{\theta} = \theta + \text{diag}(\sigma_i)\eta, \eta \sim \mathcal{N}(0, I)$ , Backprop through  $\theta$  and  $\sigma$ .

### 6.2 Linearized DNNs and NTK

Training neural network  $f(\theta)(x)$  can be approximated by **linearizing** around initialization  $\theta_0$  when parameters change slowly.

**Linearization  $\rightarrow$  Kernel Regression:**

Taylor approximation:  $h(\beta)(x) = f(\theta_0)(x) + \beta \cdot \nabla f(\theta_0)(x), \quad \beta = \theta - \theta_0$ .

With residuals  $\tilde{y}_i = y_i - f(\theta_0)(x_i)$ , training becomes **linear regression** with features  $\nabla f(\theta_0)(x_i)$ :  $\min \|\tilde{y}_i - \beta \cdot \nabla f(\theta_0)\|$

**Neural Tangent Kernel (NTK):**

Definition:  $k_\theta(x, x') := \nabla f(\theta)(x) \cdot \nabla f(\theta)(x')$ .

**Dual representation:**

$$h(\alpha)(x) = f(\theta_0)(x) + \sum_{i=1}^n \alpha_i k_{\theta_0}(x_i, x).$$

**Optimization problem:**  $\min_\alpha \frac{1}{2} \|K_{\theta_0} \alpha - \tilde{y}\|^2$

**Optimal solution (kernel regression):**

$$\alpha^* = K_{\theta_0}^\dagger (y - f(\theta_0)), h^*(x) = k_{\theta_0}(x)^\top \alpha^*$$

**Functional Gradient Flow**

Training dynamics in function space:

$$\dot{f}(\theta) = K(\theta)(y - f(\theta))$$

• If  $K(\theta)$  constant  $\rightarrow$  **linear ODE** with closed-form solution

• If  $K(\theta)$  evolves  $\rightarrow$  **nonlinear dynamics**, feature learning

**Infinite-Width Limit**

Initialization:  $w_{ij}^{(\ell)} \sim \frac{\sigma_w}{\sqrt{m_\ell}} \mathcal{N}(0, 1)$ . Result: As width  $m \rightarrow \infty$ :  $k_{\theta(t)} \rightarrow k_\infty$  (constant during training).

• Kernel becomes **deterministic** (depends only on architecture/init scheme)

• Training = kernel regression with frozen  $k_\infty$

• No feature learning

**Finite-width:**  $\|K(\theta_0) - K(\theta(t))\| = \mathcal{O}(\frac{1}{m})$

**Why Kernel Stays Constant:**

Kernel grad  $\nabla K = \nabla^2 f(x) \nabla f(z) + \nabla^2 f(z) \nabla f(x)$ ,

at  $m = \infty$ :  $\nabla^2 f \rightarrow 0 \Rightarrow \nabla K \rightarrow 0 \rightarrow$  kernel frozen.

Lazy Training (NTK Regime)	Feature Learning (Rich Regime)
$m \rightarrow \infty$ , small LR	Finite width, normal LR
$K$ const $\rightarrow$ linear dy.	$K$ evolves $\rightarrow$ nonlinear nematics
No feature learning	Learns representations
Theoretically tractable	SOTA performance

For  $y = f(x) + \varepsilon, f(x) = x^\top w, \varepsilon \sim \mathcal{N}(0, \sigma^2)$ ,  $p(y \mid X, w) = \mathcal{N}(X^\top w, \sigma^2 I)$ . If  $w \sim \mathcal{N}(0, \Sigma)$ ,  $p(w \mid X, X) = \mathcal{N}(\Sigma_{\text{post}}^{-\frac{1}{2}} X^\top w, \Sigma_{\text{post}})$ , where  $\Sigma_{\text{post}}^{-1} = \Sigma^{-1} + \frac{1}{\sigma^2} X X^\top$ . Maximizing log  $p(w \mid X, X)$  is same as minimizing least-squares with  $\ell_2$  penalty  $\frac{1}{2} w^\top \Sigma^{-1} w$ , ridge when  $\Sigma = \lambda^{-1} I$ . **Predictive:**  $f_* \mid x_*, y, X = x_*^\top w \sim \mathcal{N}(x_*^\top \mu_{\text{post}}, x_*^\top \Sigma_{\text{post}} x_*)$ .  $y_* \mid x_*, y, X$  adds  $\sigma^2$ .  $f = x_*^\top w$  with  $w \sim \mathcal{N}(0, \Sigma)$  is GP w  $k(x, x') = x_*^\top \Sigma x'$ .

### 6.3.2 NNGPs

**Setup:** Random 1-hidden-layer NN with  $m$  units:  $f(x) = v_0 + \frac{1}{\sqrt{m}} \sum_{j=1}^m v_j \varphi(\theta_j^\top x)$ . Random init:  $v_0 \sim \mathcal{N}(0, \sigma_0^2), \mathbb{E}[v_j^\top] = \sigma_v^2, \text{Cov}(\theta_j) = \Sigma_\theta$ .

**Result:** As  $m \rightarrow \infty$ ,  $f(\cdot) \rightarrow$  GP( $0, k$ ) where  $k(x, x') = \sigma_v^2 + \sigma_v^2 \mathbb{E}_\theta[\varphi(\theta^\top x) \varphi(\theta^\top x')]$

**Monte Carlo approximation:** Sample  $B$  random NNs  $\{f_b\}_{b=1}^B$ , define features:

- $\varphi(x) = \frac{1}{\sqrt{B}} (f_1(x), \dots, f_B(x))^\top$
- $\Phi = [f_1 \dots f_B] \in \mathbb{R}^{n \times B}$  (feature matrix)
- $\hat{K} = \Phi \Phi^\top$  (approximate kernel matrix)

**GP regression:** Posterior mean and variance:  $\mathbb{E}[f(x) \mid y] = \varphi(x)^\top (\Phi^\top \Phi + \sigma^2 I_B)^{-1} \Phi^\top y$ ,  $\text{Var}[f(x) \mid y] = \sigma^2 \varphi(x)^\top (\Phi^\top \Phi + \sigma^2 I_B)^{-1} \varphi(x)$

**Key advantage:** Inverts  $B \times B$  matrix instead of  $n \times n$  when  $B \ll n$ .

## 7 Generative Models

### 7.0.1 Linear Autoencoders

**Setup.** Encoder  $C \in \mathbb{R}^{k \times d}$ , decoder  $D \in \mathbb{R}^{d \times k}$ , data  $X \in \mathbb{R}^{d \times n}$  (centered cols):  $\min_{C, D} \|X - DCX\|_F^2$ .

**Optimal Solution (PCA).** Let  $S = XX^\top$  with eigendecomposition  $S = Q\Lambda^2 Q^\top$ ,  $\lambda_1 \geq \dots \geq \lambda_d \geq 0$ . Optimal reconstruction via rank- $k$  projection:  $\hat{X} = U_k^* U_k^{*\top} X$  where  $U_k^* = Q_{[:,1:k]}$  are top- $k$  eigenvectors of  $S$  (equiv, top- $k$  left singular vectors of  $X$ ).

- Any  $C = U_k^* U_k^{*\top}, D = A^{-1}(U_k^*)^\top$  is optimal ( $\forall A$ )
- Reduces to truncated SVD:  $\hat{X} = U^* \Lambda_k V^\top$  with  $\Lambda_k = \text{Diag}(\lambda_1, \dots, \lambda_k, 0, \dots, 0)$
- Convex objective with no spurious local minima (gradient descent finds global optimum)
- Singular vectors may not be uniquely identified

### 7.0.2 Factor analysis

**Latent Variable Models** are a generic way to describe generative models. Latent variable  $z \sim p(z)$ , conditional models for observables  $x, p(x|z)$ , observed data model:  $p(x) = \int p(x|z)p(z) dz$ .

**Mixture models:** simple discrete models:  $z \in [K], p(z)$  mixing proportions,  $p(x|z)$  condit. densities.

$$x \sim \mathcal{N}(\mu, \Sigma): p(x; \mu, \Sigma) = \frac{\exp[-\frac{1}{2}(x - \mu)^\top \Sigma^{-1}(x - \mu)]}{\sqrt{(2\pi)^n \det(\Sigma)}}$$

### 7.0.3 Linear Factor Analysis

- Latent prior:  $z \sim \mathcal{N}(0, I)$ ,  $z \in \mathbb{R}^m$
  - Observation:  $x = \mu + Wz + \eta$ ,  $\eta \sim \mathcal{N}(0, \Sigma)$
  - Independence:  $\eta \perp z$
  - Typically  $m < n$  (fewer factors than features)
- Marginal distribution:**  $x \sim \mathcal{N}(\mu, WW^\top + \Sigma)$
- $WW^\top$ : shared variance (low-rank, explained by latent factors)
  - $\Sigma$ : unique variance (diagonal, observation-specific)
- Non-identifiability:  $(WQ)(WQ)^\top = WQQ^\top W^\top = WW^\top$  for any orthogonal  $Q$ . Factors only identifiable up to rotations/reflections.  $\Rightarrow$  Use factor rotations (varimax, etc.) for interpretability.

**MLE estimation:**  $\theta = (\mu, W) \xleftarrow{\text{max}} \log p(X; \mu, W)$

- $\hat{\mu} = \frac{1}{s} \sum_{i=1}^s x_i$  (closed form)
- No closed form for  $W \rightarrow$  use GD or EM algorithm

**Posterior (encoder):**  $p(z|x) = \frac{p(x|z)p(z)}{p(x)}$ .

$$\mu_{z|x} = W^\top (WW^\top + \Sigma)^{-1}(x - \mu).$$

$$\Sigma_{z|x} = I - W^\top (WW^\top + \Sigma)^{-1}W.$$

**Probabilistic PCA:** Special case  $\Sigma = \sigma^2 I$ . Optimal  $i$ -th column:  $w_i = \rho_i u_i$ ,  $\rho_i^2 = \max\{0, \lambda_i - \sigma^2\}$ .  $W = U_m L_m$  where  $(\lambda_i, u_i)$  is  $i$ -th eigenpair of data covariance.

As  $\sigma \rightarrow 0$ :  $\mu_{z|x} \rightarrow W^\dagger(x - \mu)$  (standard PCA). If  $W$  has orthogonal columns, then  $W^\dagger = W^\top$ .

## 7.1 Variational Autoencoders

$z \in \mathbb{R}^d$  is learned embedding of  $x$ . For generation,  $z \sim \mathcal{N}(0, I)$ , decoder  $p_\theta(x|z)$  maps latent to data.

**Problem:**  $p_\theta(x) = \int p(z)p_\theta(x|z) dz$  intractable. **Solution:** Maximize ELBO instead:  $\log p_\theta(x) \geq \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)] - D_{\text{KL}}(q_\phi(z|x) \parallel p(z))$

- Reconstruction:** Encode  $x \rightarrow z$ , decode back
- KL term:** Keep encoder output close to prior  $p(z) \sim \mathcal{N}(0, I)$ , ensures generation using latents

Encoder  $q_\phi(z|x) = \mathcal{N}(\mu_\phi(x), \text{diag}(\sigma_\phi^2(x)))$ .

**KL closed form:**  $D_{\text{KL}}(q_\phi(z|x) \parallel p(z)) = \frac{1}{2} \sum_{j=1}^d (\sigma_{\phi,j}^2 + \mu_{\phi,j}^2 - 1 - \log \sigma_{\phi,j}^2)$ .

$\text{KL}(\mathcal{N}(\mu_0, \sigma_o^2) \parallel \mathcal{N}(\mu_1, \sigma_1^2)) = \frac{1}{2} \left( \frac{\sigma_1^2}{\sigma_o^2} + \frac{(\mu_0 - \mu_1)^2}{\sigma_o^2} - 1 + \log \frac{\sigma_1^2}{\sigma_o^2} \right)$

$$\text{KL}(p\parallel q)) = \mathbb{E}_p [\log \frac{p(x)}{q(x)}]$$

- Fwd KL:**  $q_1^* = \arg \min_{q \in Q} \text{KL}(p \parallel q)$
- Rev KL:**  $q_2^* = \arg \min_{q \in Q} \text{KL}(q \parallel p)$

Rev KL: Mode-seeking ( $p = 0 \Rightarrow q = 0$ ), FwdKL: Mean-seeking ( $p \neq 0 \Rightarrow q \neq 0$ ). MLE minimizes fwd KL to empirical  $\mathcal{D}$ .

$$\log p_\theta(x|z) = -\frac{1}{2\sigma^2} \|x - \mu_\theta(z)\|_2^2 - \frac{d}{2} \log(2\pi\sigma^2).$$

**Reparameterization trick:**  $z = \mu_\phi(x) + \sigma_\phi(x) \odot \varepsilon$ ,  $\varepsilon \sim \mathcal{N}(0, I)$ , enables backprop through sampling.

$$\log p_\theta(x) - \text{ELBO} = D_{\text{KL}}(q_\phi(z|x) \parallel p_\theta(z|x)), \text{ tight when } q_\phi = \text{true posterior.}$$

### Monte

**Carlo estimation:**  $E_{q_\phi(z|x)}[\log p_\theta(x|z)] \approx -\frac{1}{2\sigma^2 K} \sum_{k=1}^K \|x - \mu_\theta(z_k)\|_2^2 - \frac{d}{2} \log(2\pi\sigma^2)$ .

**Generative Classifiers** Given  $y \in \{0, 1\}$ ,  $p(y=1) = p(y=0) = \frac{1}{2}$ ,  $p(x|y) = \mathcal{N}(x; \mu_y, I_d)$ , where  $x \sim \mathcal{N}(\mu_0, \mu_1 \in \mathbb{R}^d)$ ,  $p(y=1|x) = \frac{p(y=1)p(x|y=1)}{p(y=1)p(x|y=0) + p(y=0)p(x|y=0)} = \frac{\frac{1}{2}(2\pi)^{-\frac{d}{2}} \exp(-\frac{1}{2}\|x - \mu_1\|^2)}{\frac{1}{2}(2\pi)^{-\frac{d}{2}} \exp(-\frac{1}{2}\|x - \mu_0\|^2) + \frac{1}{2}(2\pi)^{-\frac{d}{2}} \exp(-\frac{1}{2}\|x - \mu_0\|^2)} = \frac{1 + \exp(\frac{1}{2}\|x - \mu_1\|^2 - \frac{1}{2}\|x - \mu_0\|^2)}{1 + \exp(-(\frac{1}{2}\|x - \mu_1\|^2 + \frac{1}{2}\|x - \mu_0\|^2))}$  equiv to logistic regression where  $p(y=1|x) = \sigma(w^\top x + b)$  with  $w = \mu_1 - \mu_0$ ,  $b = \frac{1}{2}(\|\mu_0\|^2 - \|\mu_1\|^2)$ .

**ELBO for Hierarchical VAEs:** Model  $x$  by decoding from latents  $z = (z_1, \dots, z_L)$ .  $p_\theta(x, z) = p_\theta(z_1) \prod_{i=1}^{L-1} p_\theta(z_i|z_{i+1}) p(z_L)$ . Inference top-down:  $q_\phi(z|x) = q_\phi(z_L|x) \prod_{i=1}^{L-1} q_\phi(z_i|z_{i+1})$ .

**ELBO for HVAE:**  $\mathcal{L}(x) = \mathbb{E}_{z|x \sim q_\phi} [\log \frac{p_\theta(x,z)}{q_\phi(z|x)}] = \mathbb{E}_{z|x \sim q_\phi} [\log p_\theta(x|z_1) + \log \frac{p_\theta(z_1|z_2)}{q_\phi(z_1|z_2)} + \dots + \sum_{i=2}^{L-1} \log \frac{p_\theta(z_i|z_{i+1})}{q_\phi(z_i|z_{i+1})} + \log \frac{p_\theta(z_L)}{q_\phi(z_L|z_{L-1})}]$ .

**Change of variables** spherical to 3D euclidian  $(x, y, z) \mapsto (r \cos \theta \cos \phi, r \cos \theta \sin \phi, r \sin \theta)$ . Lengths of the three sides of an infinitesimal cuboid whose diagonally opposite vertices are at  $r, \theta, \phi$  and  $(r + dr, \theta + d\theta, \phi + d\phi)$  are  $(dr, rd\theta, r \cos \theta d\phi)$ . Volume is  $r^2 \cos \theta dr d\theta d\phi$ . Determinant of jacobian  $\left| \frac{\partial(x,y,z)}{\partial(r,\theta,\phi)} \right| = r^2 \cos \theta$ .

Density on spherical coordinates  $p(r, \theta, \phi) \rightarrow$  density on Euclidian coordinates is  $p(x, y, z) = p(r, \theta, \phi) \left| \frac{\partial(x,y,z)}{\partial(r,\theta,\phi)} \right|^{-1}$ . Infinitesimal probability mass of the cuboid above is equal to the mass of a Euclidian cuboid of size  $(dx, dy, dz)$  at  $(x, y, z)$ .

## 7.2 Normalizing Flows

Transform simple distribution  $z \sim \mathcal{N}(0, I)$  through invertible map  $T$  to get complex  $x = T(z)$ . Exact likelihood (no ELBO like VAEs), easy sampling.

### Change of Variables Formula:

$$p_x(x) = p_z(T^{-1}(x)) \cdot \det J_{T^{-1}}(x), \quad |\det J_{T^{-1}}(x)| = \frac{1}{|\det J_T(T^{-1}(x))|}.$$

**Diffeomorphism:**  $T$  is bijective, differentiable, with differentiable inverse. Guarantees  $\det J_T \neq 0$ .

**Computational problem:** Computing  $\det J$  is  $O(d^3)$  for dense Jacobian. **Solution:** Design  $T$  s.t. Jacobian is **triangular**, then only  $O(d)$ !

### Two architectures with triangular Jacobians:

	MAF	IAF
Fast / parallel	Density eval	Sampling
Slow / sequential	Sampling	Density

**Coupling layers:** Trick that makes both directions fast, at the cost of being less expressive per layer.

### 7.3 Autoregressive Models

$$p(x) = \prod_{i=1}^d p(x_i | x_{<i}).$$

## 7.4 Generative Adversarial Networks

Likelihood-free generative model: train via adversarial game between two networks: **Generator**  $G_\theta$  maps latent  $z \sim p_z$  (typically Gaussian) to fake samples; **Discriminator**  $D_\varphi$ : outputs prob that input is real.

### GAN

**Objective:**  $\min_\theta \max_\varphi \mathbb{E}_{x \sim p_r} [\log D_\varphi(x)] + \mathbb{E}_{z \sim p_z} [\log (1 - D_\varphi(G_\theta(z)))]$

- Discriminator** maximizes: correctly classify real (high  $D$ ) and fake (low  $D$ )
- Generator** minimizes: fool discriminator (make  $D(G(z))$  high)

Common alternative objective for the generator is to maximize  $\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$  instead of minimizing  $\mathbb{E}_{z \sim p(z)} [\log (1 - D(G(z)))]$  to help mitigate vanishing gradient problem when discriminator becomes good, i.e.  $D(G(z)) \rightarrow 0$ . For a fixed generator  $G$ , **optimal discriminator**  $D^*$  is given by  $D^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_G(x)}$ .

If **discriminator optimal**, GAN objective reduces to  $V(D^*, G) = 2D_{\text{JS}}(p_{\text{data}} \parallel p_G) - \log 4$ .

### 7.4.1 Theoretical Foundation

Binary classification with  $p(y=1) = p(y=0) = \frac{1}{2}$ :

- $y=1$ : sample from real  $p_{r(x)}$
- $y=0$ : sample from generator  $p_{\theta(x)}$

**Bayes Optimal Classifier** (prob that  $x$  is real):  $p_\theta(x) = P(y=1|x) = \frac{p_{r(x)}}{p_{r(x)} + p_\theta(x)}$ .

### Generator Logistic Objective = JS Divergence:

$$\ell^*(\theta) = \mathbb{E}_{\tilde{p}_\theta(x,y)} [y \ln q_\theta(x) + (1-y) \ln(1 - q_\theta(x))] = \text{JS}(p_r \parallel p_\theta) - \ln 2.$$

**Jensen-Shannon Divergence:**  $\text{JS}(p_r \parallel p_\theta) = \frac{1}{2} D_{\text{KL}}(p_r \parallel p_m) + \frac{1}{2} D_{\text{KL}}(p_\theta \parallel p_m)$ ,  $p_m = \frac{p_r + p_\theta}{2}$ .

**Bounded:**  $0 \leq \text{JS}(p_r \parallel p_\theta) \leq \log 2$ .

### Jensen Inequality

If  $\varphi$  convex:  $\varphi(\mathbb{E}[X]) \leq \mathbb{E}[\varphi(X)]$ . If concave, other way around.

### 7.4.2 Training

Alternating SGD (heuristic, may diverge!). Training is **Saddle-point problem**, notoriously unstable!

**JS Divergence Saturates** when distributions don't overlap. If  $p_r$  and  $p_\theta$  have disjoint supports: discriminator perfect, no gradient for generator!

### Wasserstein Distance (Earth Mover's Distance):

$$W(p_r, p_\theta) = \inf_{\gamma \in \Pi(p_r, p_\theta)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$

Minimum total "work" to transport mass from  $p_r$  to  $p_\theta$ . Provides meaningful gradients even without overlap.

### Dual (Kantorovich-Rubinstein):

$$W(p_r, p_\theta) = \sup_{\|f\|_L \leq 1} \mathbb{E}_{x \sim p_r} [f(x)] - \mathbb{E}_{x \sim p_\theta} [f(x)].$$

Maximize gap between avg score of real vs fake samples w.r.t. Lipschitz constraint. Max achievable gap = Wasserstein distance.

**WGAN** uses critic  $f_w$  (not classical discriminator!):

$$\min_\theta \max_w \mathbb{E}_{x \sim p_r} [f_w(x)] - \mathbb{E}_{z \sim p_z} [f_w(G_\theta(z))].$$

### Enforcing Lipschitz:

- Weight clipping** (original): crude, problematic
- Gradient penalty**: add  $\lambda \mathbb{E}_{\hat{x}} [\| \nabla_{\hat{x}} f_w(\hat{x}) \|_2 - 1]^2$

**Mode Collapse:** Generator produces only few samples that fool discriminator, ignoring full distribution diversity.

## 7.5 Diffusion Models

**Forward process (fixed):** Gradually add Gaussian noise over  $T$  steps until data becomes pure noise.

Fwd step:  $q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{1-\beta_t}x_{t-1}, \beta_t I)$ . Full fwd process:  $q(x_{1:T} | x_0) = \prod_{t=1}^T q(x_t | x_{t-1})$ .

**Noise schedule:**  $\{\beta_t \in (0, 1)\}_{t=1}^T$  noise added at each step. Define:  $\alpha_t = 1 - \beta_t$  and  $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$ .

### Direct sampling (reparameterization trick):

$$q(x_t | x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I). \quad x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\varepsilon, \quad \varepsilon \sim \mathcal{N}(0, I).$$

**Reverse process (learned):** Train NN to denoise step by step:  $p_\theta(x_{t-1} | x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \sigma_t^2 I)$ .

For small  $\beta_t$ , the reverse  $q(x_{t-1} | x_t)$  is also Gaussian.

$\log p_\theta(x_0)$  intractable, so derive **Variational Lower Bound (VLB)**:  $-\log p_\theta(x_0) \leq \mathcal{L}_{\text{VLB}} =$

$$\mathbb{E}_q \left[ \log \frac{q(x_{1:T} | x_0)}{p_\theta(x_{1:T})} \right]. \quad \text{Decomposition into 3 terms: } \mathcal{L}_{\text{VLB}} = \underbrace{D_{\text{KL}}(q(x_T | x_0) \parallel p(x_T))}_{L_T} +$$

$$\underbrace{\sum_{t=2}^T \mathbb{E}_{q(x_t | x_0)} [D_{\text{KL}}(q(x_{t-1} | x_t, x_0) \parallel p_\theta(x_{t-1} | x_t))]}_{L_{t-1}} - \underbrace{\mathbb{E}_{q(x_1 | x_0)} [\log p_\theta(x_0 | x_1)]}_{L_0}$$

- $L_T$ : Is  $q(x_T | x_0) \approx \mathcal{N}(0, I)$ ? Not optimized.
- $L_{t-1}$ : Match learned reverse to true reverse
- $L_0$ : Reconstruction term

**Tractable Reverse Posterior**  $q(x_{t-1} | x_t, x_0)$  is Gaussian with closed form (product of Gaussians):  $q(x_{t-1} | x_t, x_0) = \mathcal{N}(x_{t-1}; \mu_{q,t}(x_t, x_0), \sigma_t^2 I)$ ,

$$\mu_{q,t}(x_t, x_0) = \frac{1}{\sqrt{\alpha_t}} \left( x_t - \frac{1 - \alpha_t}{\sqrt{1 - \alpha_t}} \varepsilon_0 \right), \quad \sigma_t^2 = \frac{1}{1 - \alpha_{t-1}} \beta_t.$$

## 7.6 Noise Prediction Parameterization

Predict the noise  $\varepsilon_\theta(x_t, t)$  instead of mean directly. Parameterize learned mean to mirror true posterior:  $\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left( x_t - \frac{1 - \alpha_t}{\sqrt{1 - \alpha_t}} \varepsilon_\theta(x_t, t) \right)$ . Since both distributions are Gaussian with same variance:  $D_{\text{KL}}(\mathcal{N}(\mu_q, \sigma_q^2) \parallel \mathcal{N}(\mu_p, \sigma_p^2)) = \frac{1}{2\sigma^2} \|\mu_q - \mu_p\|^2$ . This simplifies  $L_{t-1}$  to comparing noise:  $L_{t-1} = \mathbb{E}_{x_0, \varepsilon_0} \left[ \frac{(1 - \alpha_t)}{2\alpha_t(1 - \alpha_t)\sigma_t^2} \|\varepsilon_0 - \varepsilon_\theta(x_t, t)\|^2 \right]$ .

$$\mathcal{L}_{\text{simple}} = \mathbb{E}_{t \sim [1, T], x_0, \varepsilon_0} [\|\varepsilon_0 - \varepsilon_\theta(x_t, t)\|^2].$$

Training	Sampling
<ol style="list-style-type: none"> <li>Sample real image <math>x_0 \sim q(x_0)</math></li> <li>Sample random timestep <math>t \sim \text{Uniform}\{1, \dots, T\}</math></li> <li>Sample noise <math>\varepsilon \sim \mathcal{N}(0, I)</math></li> <li>Compute noisy image: <math>x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\varepsilon</math></li> <li>Grad step on <math>\nabla_\theta \mathcal{L} - \varepsilon_\theta(x_t, t)^2</math></li> </ol>	<ol style="list-style-type: none"> <li>Sample <math>x_T \sim \mathcal{N}(0, I)</math></li> <li>For <math>t = T, \dots, 1</math>:             <ul style="list-style-type: none"> <li><math>z \sim \mathcal{N}(0, I)</math> if <math>t &gt; 1</math>, else <math>z = 0</math></li> <li><math>x_{t-1} = \frac{1}{\sqrt{\alpha_t}}(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \alpha_t}}\varepsilon_\theta(x_t, t)) + \sigma_t z</math></li> </ul> </li> <li>Return <math>x_0</math></li> </ol>

Cosine noise schedule performs better than linear.

Used architecture is U-Net. **Input:** Noisy image  $x_t + \text{timestep } t$ ; **Output:** Predicted noise  $\varepsilon_\theta(x_t, t)$ .

Model **conditional distribution**  $p_\theta(x_{0:T} | y)$  where  $y$  is condition (class, text, image). Extend denoiser to take  $y$  as input.

**Latent Diffusion Models (LDM)** run diffusion in **compressed latent space** instead of pixel space.

## 8 Tricks

**Short Connections** in DNs: Add less deep paths to a very deep network. **Residual connections**:

shortcut and add back in. Skip connections: concatenate.

For  $y_1 = \alpha f(x, \Theta_1) + x$ ,  $x \in \mathbb{R}^d$ ,  $y_i \in \mathbb{R}^d$ ,  $\Theta_i \in \mathbb{R}^{d \times d}$  and  $y_{i \geq 2} = \alpha f(y_{i-1}, \Theta_i) + y_{i-1}$ ,  $l = L(y_n)$  it holds that  $\frac{\partial y_k}{\partial \Theta_{k-1}} = \alpha \frac{\partial f(y_{k-1}, \Theta_{k-1})}{\partial y_{k-1}} + I_d$  and  $\frac{\partial y_k}{\partial \Theta_k} = \alpha \frac{\partial f(y_{k-1}, \Theta_k)}{\partial \Theta_k} = \alpha \frac{\partial L(y_n)}{\partial y_n} \frac{\partial y_{n-1}}{\partial y_{n-1}} \dots \frac{\partial y_{k+1}}{\partial y_k} \alpha \frac{\partial f(y_{k-1}, \Theta_k)}{\partial \Theta_k}$ . By applying the chain rule we have  $\frac{\partial \Theta_k}{\partial \Theta_k} = \frac{\partial L(y_n)}{\partial y_n} \frac{\partial y_{n-1}}{\partial y_{n-1}} \dots \frac{\partial y_{k+1}}{\partial y_k} \alpha \frac{\partial f(y_{k-1}, \Theta_k)}{\partial \Theta_k}$ . Set  $\alpha^2 = \frac{a}{n}$  for  $a > 0$  s.t.  $\lim_{n \rightarrow \infty} \mathbb{E}\|y\|^2 < \infty$ .

## 8.1 Weight Decay & Early Stopping

L2 regularization

$$\mathcal{R}_\Omega(\theta; \mathcal{S}) = \mathcal{R}(\theta; \mathcal{S}) + \Omega(\theta), \quad \Omega_\mu(\theta) = \frac{\mu}{2} \|\theta\|^2, \quad \mu \geq 0.$$

Only penalize weights, not biases. **GD upd w/ WD:**  $\Delta\theta = -\eta \nabla \mathcal{R}(\theta) - \eta \nabla \Omega_\mu(\theta) = -\eta \nabla \mathcal{R}(\theta) - \eta \mu \theta$ .

Geometric interpretation (local quadratic approx): Regularized optimum:  $\theta_\mu^* = (H + \mu I)^{-1} H \theta^*$ , where  $H = Q^\top \Lambda Q$  gives  $\theta_\mu^* = Q \text{diag}\left(\frac{\lambda_i}{\lambda_i + \mu}\right) Q^\top \theta^*$ .

$\lambda_i \gg \mu: \frac{\lambda_i}{\lambda_i + \mu} \approx 1 \rightarrow$  weak shrinkage (important dirs).

$\lambda_i \ll \mu: \frac{\lambda_i}{\lambda_i + \mu} \approx 0 \rightarrow$  strong shrinkage (flat dirs).

Adaptively shrinks based on loss geometry, preserves important dirs, removes unnecessary complexity.

**Early stopping:** Rather than training to convergence, stop when validation performance plateaus. Analysis shows that this is approximately equivalent to L2 regularization. GD trajectories can be approximated as  $\theta(k) = [I - (I - \eta \Lambda)^k] \theta^*$ . For small step sizes, behaves like weight decay when  $k = \frac{1}{\eta \mu}$ .

$L^1$  regularized second-order approximation of an arbitrary loss function around optimal  $\theta^*$  is  $R_{L^1}(\theta) \approx R(\theta^*) + \frac{1}{2}(\theta - \theta^*)^\top H(\theta - \theta^*) + \lambda \|\theta\|_1$ . Assuming  $H = \text{diag}(h_1, \dots, h_d)$ , we get  $R_{L^1}(\theta) \approx \sum_{i=1}^d \left[ \frac{h_i}{2} (\theta_i - \theta_i^*)^2 + \lambda |\theta_i| \right] + \text{const}$  so we need to minimize  $f(a) = \frac{1}{2}(a - b)^2 + \beta |a|$  and this gives  $a^* = \text{sgn}(b) \max\{0, |b| - \beta\}$ , so  $\theta_i = \text{sgn}(\theta_i^*) \max\{0, |\theta_i^*| - \frac{\lambda}{h_i}\}$ . For  $L^2$  regularization we get  $\theta_i = \frac{h_i}{h_i + \lambda} \theta_i^*$ . **Connecting L<sup>2</sup> w early stopping:**  $R(w) \approx R(w^*) + \frac{1}{2}(w - w^*)^\top H(w - w^*)$ .  $\nabla R(w) = H(w - w^*)$ . GD update:  $w' = w^{t-1} - \eta H(w^{t-1} - w^*)$  gives  $w^t - w^* = (I_d - \eta H)(w^{t-1} - w^*)$ . Using  $H = Q \Lambda Q^\top$ :  $Q^\top (w^t - w^*) = (I_d - \eta \Lambda) Q^\top (w^{t-1} - w^*)$ . If  $w^0 = 0$ ,  $Q^\top (w^t - w^*) = (I_d - \eta \Lambda)^t Q^\top (0 - w^*) \Rightarrow Q^\top w^t = [I_d - (I_d - \eta \Lambda)^t] Q^\top w^*$ . Optimal  $w$  under  $L^2$  reg gives  $Q^\top w = [I_d - \lambda(\Lambda + \lambda I_d)^{-1}] Q^\top w^*$ . Matching both gives  $t \approx \frac{1}{\eta \lambda}$ . **Weight normalization** is like Batch-Norm, with the covariance matrix replaced by the identity matrix.

## 8.2 Ensemble Methods

**Bagging:** Create  $K$  bootstrap samples of Data (sampling with replacement), train separate models, and average predictions:  $p(y|x) = \frac{1}{K} \sum_{k=1}^K p(y|x; \theta_k)$

**Dropout:** Randomly drop units during training with probability  $1 - \pi$ . Creates an exponential ensemble of sub-networks sharing weights. Test time: Scale weights by  $\pi$  to approximate the ensemble average.

## 8.3 Normalization

**Batch Norm:** Normalize activations across mini-batch:  $\tilde{z} = \frac{z - \mu_{\text{batch}}}{\sigma_{\text{batch}}}$ ,  $\hat{z} = \alpha \tilde{z} + \beta$ ,  $\mu_{\text{batch}} = \frac{1}{b} \sum_{i=1}^b z_i$ ,  $\sigma_{\text{batch}} = \sqrt{\frac{1}{b} \sum_{i=1}^b (z_i - \mu_{\text{batch}})^2}$ .

**Layer Norm:** Normalize features in a layer instead; particularly effective for RNNs (batch statistics are less stable).

## 8.4 Data / Task Augmentation

Augment Data by applying valid transformations.

**Semi-supervised Learning:** Train jointly on labeled and unlabeled data w combined loss. Pre-training & Fine-tuning. Multi-task Learning. Self-supervised Learning: Create free supervision from data.

## 9 Recurrent Neural Networks

**Evolution:**  $z_t = F[\theta](z_{t-1}, x_t)$ , with  $z_0 = 0$ . Optional output:  $y_t = G[\theta](z_t)$ .

**Simple RNN:**  $z_t = \phi(W z_{t-1} + U x_t)$  where  $W \in \mathbb{R}^{m \times m}$ ,  $U \in \mathbb{R}^{m \times n}$

**Backpropagation Through Time** (param sharing):  $\frac{\partial R}{\partial w_{ij}} = \sum_t \frac{\partial R}{\partial z_i^t} \cdot \dot{\phi}_i^t \cdot z_{t-1}^t$ .  
 $\frac{\partial R}{\partial u_{ik}} = \sum_t \frac{\partial R}{\partial z_i^t} \cdot \dot{\phi}_i^t \cdot x_k^t \cdot \dot{\phi}_k^t = \phi'(F_i(z^{t-1}, x^t))$ .

**Gradient flow backward through time:**

$$\nabla_{x_t} \mathcal{R} = [\prod_{r=t+1}^s W^\top S(z^r)] \cdot J_G \cdot \nabla_y \mathcal{R}$$

**Spectral analysis:**  $\|\prod W^\top S(z^r)\|_2 \leq [\sigma_{\max}(W)]^{s-t}$

**Root cause:** Repeated matmul through time.

$\Rightarrow$  Simple RNNs cannot learn long dependencies.

**Deep RNNs** stack layers vertically:

$$z^{t,\ell} = \varphi(W_\ell z^{t-1,\ell} + U_\ell z^{t,\ell-1}) \text{ where } z^{t,0} = x_t$$

For RNN with  $z_{t+1} = \varphi(U z_t + V x_{t+1})$ ,  $L = \sum_{t=1}^T \ell(\hat{y}_t, y_t)$ , where  $\hat{y}_t$  depends on  $z_t$ . Then  $\frac{\partial L}{\partial U} = \sum_{t=1}^T \frac{\partial L}{\partial z_t} \cdot (\varphi'_t \cdot z_t)$ ,  $\frac{\partial L}{\partial V} = \sum_{t=1}^T \frac{\partial L}{\partial z_t} \cdot (\varphi'_t \cdot x_{t+1})$ .

**Weight Sharing in RNNs (LSTM):**

$$\frac{\partial L}{\partial W} = \sum_{t=1}^T \frac{\partial L}{\partial W_t}$$

**Proof idea:** Introduce dummy parameters  $\tilde{W}_i = f(W)$  for each time step. By chain rule:  $\frac{\partial L}{\partial W} = \sum_i \frac{\partial L}{\partial \tilde{W}_i} \frac{\partial \tilde{W}_i}{\partial W}$ . With constraint  $\tilde{W}_i = W$ , we have  $\frac{\partial \tilde{W}_i}{\partial W} = I$ , giving the sum.

Initialization of bias in RNNs: Use 1.

## 9.1 Long Short-Term Memory (LSTM)

- $C_t$ : cell state (internal memory, protected highway)
- $z_t$ : hidden state (external output, filtered view)

$$C_t = \underbrace{\sigma(F \tilde{x}^t) \odot C_{t-1}}_{\text{forget}} + \underbrace{\sigma(G \tilde{x}^t) \odot \tanh(V \tilde{x}^t)}_{\text{input}}$$

$$z_t = \underbrace{\sigma(H \tilde{x}^t) \odot \tanh(C_t)}_{\text{output}}, \text{ where } \tilde{x}^t = [x_t, z_{t-1}]$$

## 9.2 Gated Recurrent Unit (GRU)

**Single state  $z_t$ .** Input:  $\tilde{x}^t = [x_t, z_{t-1}]$ .

$$u_t = \sigma(U \tilde{x}^t), \quad r_t = \sigma(R \tilde{x}^t),$$

$$z_t = u_t \odot z_{t-1} + (1 - u_t) \odot \tanh(W[r_t \odot z_{t-1}, x_t])$$

Often comparable to LSTM with fewer resources. Gating creates identity paths  $\rightarrow$  better gradient flow.

## 9.3 Linear Recurrent Models

RNNs not parallelizable during training. LRU has linear dynamics:  $z_{t+1} = A z_t + B x_t$ . Diagonalize to  $A = P \Lambda P^{-1}$ ,  $\lambda_i \in \mathbb{C}$ , change basis  $\zeta_t = P^{-1} z_t$ . Then:  $\zeta_{t+1} = \Lambda \zeta_t + C x_t$ . Each dimension evolves independently (no channel mixing). Compensate with expressive output:  $y_t = \text{MLP}(\text{Re}(G z_t))$ .

**Stability:** Require  $\max |\lambda_j| \leq 1$  (spectral radius  $\leq 1$ ).

**Parameterization:**  $\lambda_i = \exp(-\exp(\nu_i) + i \varphi_i)$  ensures  $|\lambda_i| \in (0, 1)$  automatically,  $|\lambda_i| \approx 1$ : Long-term memory,  $|\lambda_i| \approx 0$ : Short-term patterns.

**Provably universal** as sequence-to-sequence map.

### 9.3.1 Connectionist Temporal Classification

**Problem:** Unsegmented sequences (e.g., speech).

**Solution:** RNN outputs prob distribution over vocabulary at each time step. Model all alignments with blank symbol “-”:  $p(\ell|x) = \sum_{\pi \in B^{-1}(\ell)} \prod_t y_{\pi_t}$ .

$B$  removes blanks and repeated symbols.

## 9.4 Sequence Learning

**Teacher Forcing:**  $p(y^t)$  depends on  $y^{1:t-1}$  only through  $z^t$ , means during autoregressive generation, model doesn't see its own predictions.

**Solution:** Add feedback connections from  $y^{t-1}$  to  $z^t$ :  $z^t = \text{RNN}(z^{t-1}, x^t, y^{t-1})$ , now model conditions on its own previous predictions  $\rightarrow$  more coherent gen.

**Professor Forcing:** Train two networks (teacher-forced + free-running), discriminator matches hidden states  $\rightarrow$  improved generalization.

**Exposure bias:** Model relies on itself where inputs come from the previous output because of the non-availability of the ground truth.

**Seq2Seq:** Input and output sequences have different lengths: Use encoder-decoder framework.

Gradients in bi-directional RNNs are computed by making a forward and backward run, then at timestep  $t$  we combine (concatenate/add) and continue with the backpropagation. This happens at every bidirectional layer.

## 11 Ethics

**Adversarial examples** (given  $f(x) = y$  correctly):

- Untargeted:  $\|\delta\| \leq \varepsilon$  s.t.  $f(x + \delta) \neq y$ .
  - Optimize  $\max_{\|\delta\| \leq \varepsilon} L(f(x + \delta), y)$ .
- Targeted:  $\|\delta\| \leq \varepsilon$  s.t.  $f(x + \delta) = t \neq y$ .
  - Optimize  $\min_{\|\delta\| \leq \varepsilon} L(f(x + \delta), t)$ .

**Linear Binary** ( $y \in \{-1, 1\}$ ,  $f(x) = w^\top x + b$ ):  
**Correct:**  $y(w^\top x + b) > 0$ .

Adv. flips when  $y w^\top \delta \leq -y(w^\top x + b)$  cross hyperplane. **L<sub>2</sub> optimal:**  $\delta^* = \frac{-w^\top x + b}{\|w\|_2^2} w$ ,  $\|\delta^*\|_2 = \frac{|w^\top x + b|}{\|w\|_2}$ .

**L<sub>∞</sub> optimal:**  $\delta = -\varepsilon \text{ sign}(yw)$ .

**Multiclass:**  $f_k(x) = w_k^\top x + b_k$ , use  $\text{argmax}_k f_k(x)$ .

**Margin** to class  $j$ :  $m_j(x) = (w_y - w_j)^\top x + (b_y - b_j)$ .

$f_j(x) = f_j(x) \Leftrightarrow m_j(x) = 0$ . **Correct** if  $f_j(x) > f_j(x) \forall j \neq y$ , **adversarial** if  $\exists j \neq y$  s.t.  $f_j(x + \delta) < f_j(x + \delta)$ . Distance to boundary:  $\frac{m_j(x)}{\|w_y - w_j\|_2}$ .

**Adversarial attacks for NNs:** Approximate boundary by  $f(x + \delta) \approx f(x) + \nabla f(x)^\top \delta$ .

**FGSM** is a one-step  $L_\infty$  attack:  $\delta = \varepsilon \text{ sign}(\nabla_x L(f(x), y))$ . **PGD** is multi-step  $\delta_{t+1} = \text{Proj}_{\|\delta\| \leq \varepsilon}(\delta_t + \alpha \text{ sign}(g_t))$ .

**Distributionally Robust Optimization:**

$\min_f \sup_{Q \in U(P)} E_Q[L(f(x))]$ , where  $U$  means close. Can use upper bound on Wasserstein distance e.g.

**Robust training**  $\min_f \mathbb{E}[\max_{\delta \in \mathcal{D}} L(f(x + \delta), y)]$ .

Adversarial training can be viewed as robustness to distribution shift measured by Wasserstein distance.

**Interpretability:** Local - explain pred for specific  $x$ , Global - explain model behaviour on avg over data.

**Local:** Ceteris paribus (vary  $x_j$ , fix  $x_{-j}$ ). Sensitivity ( $\partial_{x_j} f(x)$ ), missing info ( $f(x) - \mathbb{E}[f(X) | X_{-j} = x_{-j}]$ ).

**Global:** Mutual info ( $I(X_j; Y | [X_{-j}])$ ), Predictive util (train  $f$  w/ and w/o  $x_j$ ). For log-loss predictive util  $\approx$  conditional mutual information.

SHAP attributes predictions, while SAGE attributes risk reduction.

A protected attribute,  $Y$  target outcome,  $\hat{Y}$  prediction. Demographic Parity:  $\hat{Y} \perp A$ ; Equalized Odds:  $\hat{Y} \perp A | Y$ ; Equality of Opportunity:  $\hat{Y} \perp A | Y = 1$ .