

# 1 Connectionism

## 1.1 Perceptron

**Threshold Unit**  $f[w, b](x) = \text{sign}(x \cdot w + b)$  with  
**Dec. Boundary**  $x \cdot w + b = 0 \Leftrightarrow \frac{x \cdot w}{\|w\|} + \frac{b}{\|w\|} = 0$ .

**Geometric Margin**  $\gamma[w, b](x, y) = \frac{y(x \cdot w + b)}{\|w\|}$ .

**Maximum Margin Classifier**

$(w^*, b^*) \in \operatorname{argmax}_{w, b} \gamma[w, b](\mathcal{S})$ ,  
with  $\gamma[w, b](\mathcal{S}) := \min_{(x, y) \in \mathcal{S}} \gamma[w, b](x, y)$ .

### Perception Learning

If  $f[w, b](x) \neq y$ : update  $w += yx$ , and  $b += y$ .

$w_t \in \text{span}(x_1, \dots, x_s) \Rightarrow w_t \in \text{span}(x_1, \dots, x_s) \forall t$ .

### Convergence

1. If  $\exists w^*, \|w^*\| = 1$ , s.t.  $\gamma[w^*](\mathcal{S}) = \gamma > 0 \Rightarrow w_t \cdot w^* \geq t\gamma$ .

2. Let  $R = \max_{x \in \mathcal{S}} \|x\|$ . Then  $\|w_t\| \leq R\sqrt{t}$ .

$\cos \angle(w^*, w_t) = \frac{w^* \cdot w_t}{\|w^*\| \|w_t\|} \geq \frac{t\gamma}{\sqrt{t}R} = \sqrt{t} \frac{\gamma}{R} \leq 1 \Rightarrow t \leq \frac{R^2}{\gamma^2}$ .

### Cover's Theorem for $\mathcal{S} \subset \mathbb{R}^n$ , $|\mathcal{S}| = s$

$C(\mathcal{S}, n)$ : # of ways to separate  $\mathcal{S}$  in  $n$  dimensions.

Position of pts does not matter (general position).

$C(s+1, n) = 2 \sum_{i=0}^{n-1} \binom{s}{i}$ ,  $C(s, n) = 2^s$  for  $s \leq n$ .

Phase transition at  $s = 2n$ . For  $s < 2n$  empty version space is the exception, otherwise the rule.

## 1.1.1 Hopfield Networks

**Hopfield Model**  $E(X) = -\frac{1}{2} \sum_{i \neq j} w_{ij} X_i X_j + \sum_i b_i X_i$ , where  $X_i \in \{\pm 1\}$ .  $w_{ij} = w_{ji}$ ,  $w_{ii} = 0$ .

### Hebbian Learning

Choose patterns  $\{x\}_{t=1}^s \in \{\pm 1\}^n$ , build weights once using them:  $w_{ij} = \frac{1}{n} \sum_{t=1}^s x_i^t x_j^t$ ,  $w_{ii} = 0$ . For inference, update  $X$  iteratively:  $X^{t+1} = \text{sign}\left(\sum_j w_{ij} X_j^t + b_i\right)$  asynchronously. Capacity for random, uncorrelated patterns:  $s_{\max} \approx 0.138n$ .

## 1.2 Feedforward Networks

### 1.2.1 Linear Models

#### Linear regression (MSE)

$$L[w](X, y) = \frac{\|Xw - y\|^2}{2n}, \nabla L = \frac{X^\top Xw - X^\top y}{n}$$

#### Moore-Penrose inverse solution

$w^* = X^*y \in \operatorname{argmin}_w L[w](X, y)$ , where  $X^* = \lim_{\delta \rightarrow 0} (X^\top X + \delta I)^{-1} X^\top$  Moore-Penrose inverse.

#### Stochastic gradient descent update

$$w_{t+1} = w_t + \eta (y_{i_t} - w_t^\top x_{i_t}) x_{i_t}, i_t \sim \mathcal{U}([1, n]).$$

#### Gaussian noise model

$y_i = w^\top x_i + \varepsilon_i$ ,  $\varepsilon_i \sim N(0, \sigma^2)$ , LSQ equivalent to NLL of gaussian noise model.

### Ridge regression

$$h_\lambda[w] = h[w] + \frac{\lambda}{2} \|w\|^2, w^* = (X^\top X + \lambda I)^{-1} X^\top y.$$

### Logistic function

$$\sigma(z) = \frac{1}{1+e^{-z}}, \sigma(z) + \sigma(-z) = 1.$$

$$\sigma' = \sigma(1 - \sigma), \sigma'' = \sigma(1 - \sigma)(1 - 2\sigma)$$

### Cross entropy loss for $y \in \{0, 1\}$

$$\ell(y, z) = -y \log \sigma(z) - (1 - y) \log(1 - \sigma(z)) \\ = -\log \sigma((2y - 1)z).$$

**Logistic regression with CE loss:**  $L[w] = \frac{1}{n} \sum_{i=1}^n \ell_i(y_i, w^\top x_i), \nabla \ell_i = [\sigma(w^\top x_i) - y_i] x_i$ .

### Generic feedforward layers

$$F : \underbrace{\mathbb{R}^{m(n+1)}}_{\text{parameters}} \times \underbrace{\mathbb{R}^n}_{\text{input}} \rightarrow \underbrace{\mathbb{R}^m}_{\text{output}}, F[\theta](x) = \varphi(Wx + b).$$

### Composition of layers

$$G = F^L[\theta^L] \circ \dots \circ F^1[\theta^1].$$

### Layer activations

$$x^l = F^l \circ \dots \circ F^1(x) = F^l(x^{l-1}), x^0 = x, x^L = F(x)$$

**Softmax**  $(z)_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$ , CE-loss in terms of logits

$$\ell(y, z) = -z_y + \log \sum_j e^{z_j}.$$

**Residual layer**  $F[W, b](x) = x + (\varphi(Wx + b) - \varphi(0))$ , therefore  $F[0, 0] = \text{id}$ . Link that propagates  $x$  forward is called a **skip connection**.

### 1.2.2 Sigmoid Networks

**Sigmoid activation**  $\sigma(z) = \frac{1}{1+e^{-z}} = 1 - \sigma(-z)$ .

$$\sigma'(z) = \sigma(z)(1 - \sigma(z)).$$

### Hyperbolic tangent activation

$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}} = 2\sigma(2z) - 1.$$

$$\tanh'(z) = 1 - \tanh^2(z).$$

### Smooth function approximation

Polynomials, ridge functions ( $\varphi(a^\top x + b)$ ) and MLPs with  $C^\infty$  activations are universal approximators.

### Barron's Theorem: Approximation error

For  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  with  $C_f = \int \|\omega\| |\hat{f}(\omega)| d\omega < \infty$ ,  $\exists$  width- $m$  MLP  $g_m$  s.t.:  $\int_B |f - g_m|^2 dx \leq O(\frac{1}{m})$

### 1.2.3 ReLU( $z$ ) = max(0, $z$ ) networks

ReLU networks are universal approximators.

### Zalavsky's Theorem: Activation patterns

$m$  ReLU neurons in  $\mathbb{R}^n$ . Each neuron's hyperplane  $\{w_i^\top x = 0\}$  partitions  $\mathbb{R}^n$  into  $R(m)$  connected regions of constant activation pattern.  $R(m) \leq \sum_{i=0}^{\min(n, m)} \binom{m}{i} \ll 2^m$ .

### Montufar: Connected regions in ReLU network

$$R(m, L) \geq R(m) \left[ \frac{m}{n} \right]^{n(L-1)}, L: \text{layers}, m: \text{width}.$$

## 1.3 Gradient-Based Learning

### 1.3.1 Backpropagation

#### Parameter derivatives for ridge function layers

$$\frac{\partial x_i^l}{\partial w_{ij}^l} = \dot{\varphi}_i^l x_j^{l-1},$$

$$\dot{\varphi}_i^l := \dot{\varphi}_i^l \left( (w_i^l)^\top x^{l-1} + b_i^l \right)$$

$$\frac{\partial x_i^l}{\partial b_i^l} = \dot{\varphi}_i^l$$

### Loss derivatives

$$\frac{\partial h[\theta](x, y)}{\partial w_{ij}^l} = \frac{\partial h^{l[\theta]}(x^l, y)}{\partial x_i^l} \frac{\partial x_i^l}{\partial w_{ij}^l} = \delta_i^l \dot{\varphi}_i^l x_j^{l-1},$$

$$\frac{\partial h[\theta](x, y)}{\partial b_i^l} = \frac{\partial h^{l[\theta]}(x^l, y)}{\partial x_i^l} \frac{\partial x_i^l}{\partial b_i^l} = \delta_i^l \dot{\varphi}_i^l$$

with  $\delta_i^l = \frac{\partial h}{\partial x_i^l} \dot{\varphi}_i^l$

### 1.3.2 Gradient Descent

#### Gradient descent update

$$\theta_{t+1} = \theta_t - \eta \nabla h(\theta_t)$$

#### Gradient flow ODE

$$d \frac{\theta}{dt} t = -\nabla h(\theta)$$

### L-smoothness

$$\|\nabla h(\theta_1) - \nabla h(\theta_2)\| \leq L \|\theta_1 - \theta_2\| \text{ (forall } \theta_1, \theta_2)$$

$$\lambda_{\max}(\nabla^2 h) \leq L$$

$$\ell(w) - \ell(w') \leq \nabla \ell(w')^\top (w - w') + \frac{L}{2} \|w - w'\|^2$$

$$\ell''(x) \leq L$$

### Polyak-Lojasiewicz condition

$$\frac{1}{2} \|\nabla h(\theta)\|^2 \geq \mu(h(\theta) - \min h) \text{ (forall } \theta)$$

### Convergence rate

$$\eta = \frac{1}{L}$$

$$\Delta t = \frac{2L}{\varepsilon^2} (h(\theta_0) - \min h) \text{ for } \varepsilon\text{-critical point}$$

$$\Delta h(\theta_t) - \min h \leq (1 - \frac{\mu}{L})^t (h(\theta_0) - \min h)$$

### 1.3.3 Acceleration and Adaptivity

#### Heavy ball momentum update

$$\theta_{t+1} = \theta_t - \eta \nabla h(\theta_t) + \beta(\theta_t - \theta_{t-1})$$

#### Nesterov acceleration

$$\tilde{\theta}_{t+1} = \theta_t + \beta(\theta_t - \theta_{t-1})$$

$$\theta_{t+1} = \tilde{\theta}_{t+1} - \eta \nabla h(\tilde{\theta}_{t+1})$$

More theoretical grounding than heavy ball

### AdaGrad updates

$$\theta_{i,t+1} = \theta_{i,t} - \eta_i^t \frac{\partial h}{\partial \theta_i}(\theta_t),$$

$$\gamma_i^t = \gamma_i^{t-1} + \left( \frac{\partial h}{\partial \theta_i}(\theta_t) \right)^2,$$

$$\eta_i^t = \frac{\eta}{\sqrt{\gamma_i^t + \delta}}$$

### Adam updates

$$g_i^t = \beta g_i^{t-1} + (1 - \beta) \frac{\partial h}{\partial \theta_i}(\theta_t)$$

$$\gamma_i^t = \alpha \gamma_i^{t-1} + (1 - \alpha) \left( \frac{\partial h}{\partial \theta_i}(\theta_t) \right)^2$$

$$\theta_{i,t+1} = \theta_{i,t} - \eta_i^t g_i^t, \eta_i^t := \frac{\eta}{\sqrt{\gamma_i^t + \delta}}$$

### RMSprop

Adam without momentum term

## 1.3.4 Stochastic Gradient Descent

### Stochastic gradient descent update

$$\theta_{t+1} = \theta_t - \eta \nabla h(\theta_t)(x_{i_t}, y_{i_t})$$

### SGD variance

$$V[\theta](S) = \frac{1}{s} \sum_{i=1}^s \|\nabla h[\theta](S) - \nabla h[\theta](x_i, y_i)\|^2$$

### SGD convergence rate

$$E[h(|(\theta)_t|)] - \min h \leq O\left(\frac{1}{\sqrt{t}}\right) \text{ (general)}$$

$$E[h(|(\theta)_t|)] - \min h \leq O\left(\log \frac{t}{\delta}\right) \text{ (strongly convex)}$$

$$E[h(|(\theta)_t|)] - \min h \leq O\left(\frac{1}{t}\right) \text{ (additionally smooth)}$$

## 1.3.5 Function Properties

### Convexity

$$\ell(\lambda w + (1 - \lambda)w') \leq \lambda \ell(w) + (1 - \lambda) \ell(w')$$

$$\ell''(x) \geq 0 \text{ forall } x$$

### Convexity and differentiability

$$\ell(w) \geq \ell(w') + \nabla \ell(w')^\top (w - w')$$

Implies convexity for differentiable functions and vice versa

### Strong convexity and differentiability

$$\ell(w) \geq \ell(w') + \nabla \ell(w')^\top (w - w') + \frac{\mu}{2} \|w - w'\|^2$$

$$\ell''(x) \geq \mu \text{ forall } x$$

## 1.4 Convolutional Networks

### 1.4.1 Convolutions

#### Convolution definition

$$(f * g)(u) := \int_{-\infty}^{\infty} g(u-t) f(t) dt = \int_{-\infty}^{\infty} f(u-t) g(t) dt$$

#### Fourier transform convolution property

$$F(f * g) = F(f) * F(g)$$

#### Discrete convolution

$$(f * g)[u] := \sum_{t=-\infty}^{\infty} f[t] g[u-t]$$

#### Cross-correlation

$$(g \star f)[u] := \sum_{t=-\infty}^{\infty} g[t] f[u+t]$$

#### Toepplitz matrices

$$(f * g) = \text{Toeplitz-Matrix}(g)f$$

### 1.4.2 Convolutional Networks

#### Conventions

Padding: Add zeros around input

Stride: Step size of convolution

#### Max-Pooling

Take maximum value in windows (size  $r$ )

#### ConvNets for Images

$$y[r][s, t] = \sum_u \sum_{\Delta s, \Delta t} w[r, u][\Delta s, \Delta t] * x[u][s + \Delta s, t + \Delta t]$$

$r$ : output channel,  $u$ : input channel

## Number of parameters of a convolutional layer

$$D = (|r| * |u|) * (|\Delta s| * |\Delta t|)$$

fully connected · window size

## 1.4.3 Natural Language Processing with ConvNets

### Word embedding

$$\Omega : w \mapsto x_w \in \mathbb{R}^n$$

### Conditional log-bilinear model

Prediction of output word  $\mu$  given word  $w$  in neighborhood

$$P(\mu | w) = \frac{\exp(x_w^\top y_\mu)}{\sum_\mu} \exp(x_w^\top y_\mu)$$

$$h(\{x_w\}, \{y_\mu\}) = \sum_{(w,\mu)} \ell_{w\mu}$$

$$\ell_{w,\mu} = -x_w^\top y_\mu + \ln \sum_\mu \exp(x_w^\top y_\mu)$$

### Negative sampling

$$\tilde{\ell}_{w,\mu} = -\ln \sigma(x_w^\top y_\mu) - \beta E_{\mu \sim D} \ln(1 - \sigma(x_w^\top y_\mu))$$

## 1.5 Recurrent Networks

### 1.5.1 Simple Recurrent Networks

#### Time evolution equation

$$z_t := F[\theta](z_{t-1}, x_t), z_0 := 0 \text{ (forall } t)$$

#### Output map

$$\hat{y}_t := G[\xi](z_t)$$

#### RNN parameterization

$$F[U, V](z, x) := \varphi(Uz + Vx)$$

$$G[W](z) := \psi(Wz), W \in \mathbb{R}^{q \times m}$$

#### Backpropagation through time

$$\frac{\partial h}{\partial z_i^s} = \sum_{s=t}^T \delta_k^s \sum_{j=1}^m \frac{\partial y_k^s}{\partial z_j^s} \frac{\partial z_j^s}{\partial z_i^s}$$

$$\frac{\partial h}{\partial v_{kj}} = \dot{\psi}_k^s w_{kj}$$

$$\frac{\partial h}{\partial u_{ij}} = \sum_{t=1}^T \frac{\partial h}{\partial z_i^t} \dot{\varphi}_i^t x_j^t$$

#### Spectral norm

$$\|A\|_2 = \max_{x: \|x\|=1} \|Ax\|_2 = \sigma_1(A)$$

#### Gradient norms

$$\frac{\partial z^T}{\partial z^0} = \dot{\Phi}^T U * \dots * \dot{\Phi}^1 U$$

The norm of gradients either:

1. Vanishes exponentially if  $\sigma_1(U) < \frac{1}{\|\alpha\|}$ :  $\left\| \frac{\partial z^t}{\partial z^0} \right\|_2 \leq (|\alpha| \sigma_1(U))^t \rightarrow \infty$
2. Explodes if  $\sigma_1(U)$  is too large

#### Bidirectional RNNs

$$\hat{y}_t = \psi(Wz_t + \tilde{W}\tilde{z}_t)$$

### 1.5.2 Gated Memory

#### LSTM

$$z_t := \sigma(F\tilde{x}_t) * z_{t-1} + \sigma(G\tilde{x}_t) * \tanh(V\tilde{x}_t)$$

$$\tilde{x}_t := \text{mat}(x_t; h_t), h_{t+1} = \sigma(H\tilde{x}_t) * \tanh(Uz_t)$$

### GRU

$$\begin{aligned} z_t &= (1 - \sigma) * z_{t-1} + \sigma * \tilde{z}_t, \\ \sigma &:= \sigma(G[x_t, z_{t-1}]) \\ \tilde{z}_t &:= \tanh(V[r_t * z_{t-1}, x_t]) \\ r_t &:= \sigma(H[z_{t-1}, x_t]) \end{aligned}$$

### 1.5.3 Linear Recurrent Models

#### Linear state evolution

$$z_{t+1} = Az_t + Bx_t$$

#### Diagonal form

$$A = P\Lambda P^{-1}, \Lambda := \text{diag}(\lambda_1, \dots, \lambda_m), \lambda_i \in \mathbb{C}$$

#### Stability condition

$$\max_j |\lambda_j| \leq 1$$

#### Initialization

$$\lambda_i = \exp(-\exp(\kappa_i) + i\varphi_i),$$

$$e^{\kappa_i} = -\ln r_i$$

$$\varphi_i \sim \text{Uni}[0; 2\pi], r_i \sim \text{Uni}[I], I \subset [0; 1]$$

#### Advantages

- (i) clear modeling of long/short range dependencies
- (ii) no channel mixing required
- (iii) parallelizable training

## 1.6 Attention and Transformers

### 1.6.1 Attention

#### Attention mixing

$$\xi_s := \sum_t \alpha_{st} Wx_t, \alpha_{st} \geq 0, \sum_t \alpha_{st} = 1$$

$$A = (a_{st}) \in \mathbb{R}^{T \times T}, \text{s.t. } \Xi = WXA^\top$$

#### Query-key matching

$$Q = U_Q X, K = U_K X$$

$$(U_Q, U_K \in \mathbb{R}^{q \times n})$$

$$Q^\top K = X^\top U_Q^\top U_K X \text{ rank } \leq q$$

$$(Q^\top K \in \mathbb{R}^{T \times T})$$

#### Softmax attention

$$A = \text{softmax}(\beta Q^\top K),$$

$$a_{st} = \frac{e^{\beta [Q^\top K]_{st}}}{\sum_r} e^{\beta [Q^\top K]_{sr}}$$

$$\text{usually } \beta = \frac{1}{\sqrt{q}}$$

#### Feature transformation

$$X \mapsto \Xi \mapsto F(\Xi),$$

$$F(\theta)(\Xi) = (F(\xi_1), \dots, F(\xi_T))$$

#### Positional encoding

$$p_{tk} = \text{cases}(\sin(t\omega_k), k \text{ even}; \cos(t\omega_k), k \text{ odd}),$$

$$\omega_k = C^{\frac{k}{d}}$$

#### Transformer architecture

Self-attention: attend to its own values in the past

Cross-attention: E.g. decoder attends to encoder output (query from decoder, key and value from encoder)

## Vision transformer patch embedding

$$\mathbb{R}^{p \times p \times q} \ni \text{patch}_t \mapsto x_t := V \text{ vec}(\text{patch}_t) \in \mathbb{R}^n \text{ with } V \in \mathbb{R}^{n \times (qp^2)}$$

#### GELU activation

$$\varphi(z) = z \text{ Prob}(z \leq Z), Z \sim \text{N}(0, 1)$$

## 1.7 Geometric Deep Learning

### 1.7.1 Sets and Points

#### Function over sets

$$\{x_1, \dots, x_M\} \subset \mathbb{R}, f: 2^{\mathbb{R}} \rightarrow Y$$

#### Order-invariance property

$$f(x_1, \dots, x_M) = f(x_{\pi(1)}, \dots, x_{\pi(M)}) \text{ forall } \pi \in S_M$$

#### Equivariance property

$$f(x_1, \dots, x_M) = (y_1, \dots, y_M) \Rightarrow$$

$$f(x_{\pi(1)}, \dots, x_{\pi(M)}) = (y_{\pi(1)}, \dots, y_{\pi(M)})$$

#### Permutation invariant sum

$$\sum_{m=1}^M x_m = \sum_{m=1}^M x_{\pi(m)}, \text{forall } M, \text{forall } \pi \in S_M$$

#### Deep Sets model

$$f(x_1, \dots, x_M) = \rho\left(\sum_{m=1}^M \varphi(x_m)\right)$$

#### Max pooling variant

$$f(x_1, \dots, x_M) = \rho\left(\max_{m=1}^M \varphi(x_m)\right)$$

#### Equivariant map construction

$$\begin{aligned} \rho: \mathbb{R} \times \mathbb{R}^N &\rightarrow Y, \\ (x_m, \sum_{k=1}^M \varphi(x_k)) &\mapsto y_m \end{aligned}$$

### 1.7.2 Graph Convolutional Networks

#### Feature and adjacency matrices

$$X = \text{mat}(x_1^\top; \dots; x_M^\top), A = (a_{nm}) \text{ with } a_{nm} = \text{cases}(1, \text{if } \{v_n, v_m\} \in E; 0, \text{otherwise})$$

#### Permutation matrix constraints

$$\begin{aligned} P &\in \{0, 1\}^{M \times M} \text{ s.t.} \\ \sum_{n=1}^M p_{nm} &= \sum_{n=1}^M p_{mn} = 1 \text{ (forall } m) \end{aligned}$$

#### Graph invariance definition

$$f(X, A) \neq f(PX, PAP^\top), \text{forall } P \in \Pi_M$$

#### Graph equivariance definition

$$f(X, A) \neq Pf(PX, PAP^\top), \text{forall } P \in \Pi_M$$

#### Node neighborhood features

$$X_m := \{x_n : \{v_n, v_m\} \in E\}, \quad \{\text{cdot}\} = \text{multiset}$$

#### Message passing scheme

$$\varphi(x_m, X_m) = \varphi(x_m, m_{X_m} \psi(x))$$

$m$  is a permutation-invariant operation

#### Normalized adjacency matrix

$$|(A) = D^{-\frac{1}{2}}(A + I)D^{-\frac{1}{2}}$$

$$D = \text{diag}(d_1, \dots, d_M), d_m = 1 + \sum_{n=1}^M a_{nm}$$

## GCN layer

$$X^+ = \sigma(|(A)XW), W \in \mathbb{R}^{M \times N}$$

### Two-layer GCN

$$Y = \text{softmax}(|(A)|(A)XW^0)W^1$$

### 1.7.3 Spectral Graph Theory

#### Laplacian operator

$$\Delta f := \sum_{n=1}^N \frac{\partial^2 f}{\partial x_n^2}, f: \mathbb{R}^N \rightarrow \mathbb{R}$$

#### Graph Laplacian

$$L = D - A, (Lx)_n = \sum_{m=1}^M a_{nm}(x_n - x_m)$$

#### Normalized Laplacian

$$\tilde{L} = I - D^{-\frac{1}{2}}AD^{-\frac{1}{2}} = D^{-\frac{1}{2}}(D - A)D^{-\frac{1}{2}}$$

#### Graph Fourier transform

$$L = D - A = U\Lambda U^\top,$$

$$\Lambda := \text{diag}(\lambda_1, \dots, \lambda_M), \lambda_i \geq \lambda_{i+1}$$

#### Convolution

$$x * y = U((U^\top x) \text{ odot } (U^\top y))$$

#### Filtering operation

$$G_{\theta(L)}x = UG_{\theta(\Lambda)}U^\top x$$

#### Polynomial kernels

$$U(\sum_{k=0}^K \alpha_k \Lambda^k)U^\top = \sum_{k=0}^K \alpha_k L^k$$

#### Polynomial kernel network layer

$$x_i^{l+1} = \sum_j p_{ij}(L)x_j^l + b_i,$$

$$p_{ij}(L) = \sum_{k=0}^K \alpha_{ijk} L^k$$

### 1.7.4 Attention GNNs

#### Attention coupling matrix

$$Q = (q_{ij}),$$

$$q_{ij} = \text{softmax}(\rho(u^\top (Vx_i; Vx_j; x_{ij})))$$

$$\text{s.t. } \sum_i A_{ij} q_{ij} = 1$$

#### Attention propagation

$$X^+ = \sigma(QXW)$$

#### Weisfeiler-Lehman test

## 1.8 Tricks of the Trade

### 1.8.1 Initialization

#### Random initialization

$$\theta_i^0 \sim \mathcal{N}(0, \sigma_i^2), \text{ or}$$

$$\theta_i^0 \sim \text{Uniform}(-\sqrt{3}\sigma_i; \sqrt{3}\sigma_i)$$

#### LeCun initialization

$$w_{ij} \sim \text{Uniform}[-a; a], a := \frac{1}{\sqrt{n}}, b_i = 0$$

Stabilizes variance

#### Glorot initialization

$$w_{ij} \sim \text{Uniform}[-\sqrt{3}\gamma; \sqrt{3}\gamma],$$

$$\gamma := \frac{2}{n+m}$$

Stabilizes variance of gradients in backpropagation

## He initialization

$w_{ij} \sim N(0, \gamma)$  or  $w_{ij} \sim \text{Uniform}[-\sqrt{3}\gamma; \sqrt{3}\gamma]$ ,

$$\gamma := \frac{2}{n}$$

In ReLU networks typically only  $\frac{n}{2}$  units active

## Orthogonal initialization

$$\frac{1}{\sqrt{m}}W \sim \text{Uniform}(O(m))$$

$$\text{s.t. } W^\top W = WW^\top = mI$$

## 1.8.2 Weight Decay

### L2 regularization

$$\Omega_{\mu(\theta)} = \frac{\mu}{2}\|\theta\|^2, \mu \geq 0$$

### Gradient descent with weight decay

$$\Delta\theta = -\eta\nabla E(\theta) - \eta\nabla\Omega_{\mu(\theta)} = -\eta\nabla E(\theta) - \eta\mu\theta$$

### Weight decay for multiple layers

$$\theta = (\text{vec}(W^1), \text{vec}(W^2), \dots, \text{vec}(W^L)),$$

$$\Omega_{\mu(\theta)} = \sum_{l=1}^L \mu_l \|W^l\|_F^2$$

### Local loss landscape

$$\theta_\mu^* = (H + \mu I)^{-1} H \theta^*, H = Q^\top \Lambda Q$$

$$(\Lambda + I)^{-1} \Lambda = \text{diag}\left(\frac{\lambda_i}{\lambda_i + \mu}\right)$$

The minimum  $\theta^*$  is shrunk along directions with small eigenvalues

### Generalization

$$\mu = \frac{\sigma^2}{u^2}, u: \text{teacher "sign"al}$$

Optimal weight decay inverse proportional to the "sign"al-to-noise ratio

## 1.8.3 Dropout

### Probability $\varphi_i$ of keeping a unit

#### Dropout as Ensembling

$$p(y | x) = \sum_{b \in \{0,1\}^R} p(b)p(y | x; b)$$

$$\text{with } p(b) = \prod_{i=1}^R \varphi_i^{b_i} (1 - \varphi_i)^{1-b_i}$$

#### Weight scaling for inference

$$\tilde{w}_{ij} \leftarrow \varphi_j w_{ij}$$

## 1.8.4 Normalization

### Batch normalization

$E$  and  $V$  from minibatches or population statistics

$$|(f) = \frac{f - E[f]}{\sqrt{V[f]}}, E[|(f)] = 0, V[|(f)] = 1$$

$$|(f)[\mu, \gamma] = \mu + \gamma |(f)$$

### Weight normalization

$$f(v, \gamma)(x) = \varphi(w^\top x), w := \frac{\gamma}{\|v\|_2} v$$

Gradient descent with respect to decoupled  $\gamma$  and  $v$ :

$$\frac{\partial E}{\partial \gamma} = \nabla_w E * \frac{v}{\|v\|_2}$$

$$\nabla_v E = \frac{\gamma}{\|v\|} \left( I - \frac{ww^\top}{\|w\|^2} \right) \nabla_w E$$

### Layer normalization

$$\tilde{f}_i = \frac{f_i - E[f]}{\sqrt{V[f]}}$$

$$E[f] = \frac{1}{m} \sum_{i=1}^m f_i$$

$$V[f] = \frac{1}{m} \sum_{i=1}^m (f_i - E[f])^2$$

Using population averages across units in a layer

## 1.8.5 Model Distillation

### Tempered cross entropy loss for distillation

$$\ell(x) = \sum_{y=1}^K \frac{\exp\left[\frac{F_y(x)}{T}\right]}{\sum_{\mu=1}^K \exp\left[\frac{F_\mu(x)}{T}\right]} \left[ \frac{1}{T} G_{y(x)} - \ln \sum_{\mu=1}^K \exp\left[\frac{F_\mu(x)}{T}\right] \right]$$

$T > 0$ ,  $F_y$ : teacher logits,  $G_y$ : student logits

### Gradient of distillation loss

$$\frac{\partial \ell}{\partial G_y} = \frac{1}{T} \left[ \frac{e^{\frac{F_y}{T}}}{\sum_\mu e^{\frac{F_\mu}{T}}} - \frac{e^{\frac{G_y}{T}}}{\sum_\mu e^{\frac{G_\mu}{T}}} \right]$$

## 1.9 Theory

### 1.9.1 Neural Tangent Kernel

#### Linearized DNN taylor approximation

$$h(\beta)(x) = f(x) + \beta * \nabla f(x)$$

with  $\beta \approx \theta - \theta_0$ ,  $f(x) := f(\theta_0)(x)$

#### Kernel of gradient feature maps

$$k(x, \xi) = \nabla f(x) * \nabla f(\xi), \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$$

#### Dual representation

$$h(\alpha)(x) = f(x) + \sum_{i=1}^s \alpha_i \nabla f(x_i) * \nabla f(x)$$

#### Squared loss

$$E(\alpha) = \frac{1}{2s} \sum_{i=1}^s \left( \sum_{j=1}^s \alpha_j \nabla f(x_j) * \nabla f(x_i) + f(x_i) - y_i \right)$$

#### Optimal solution of linearized DNN

$$K = [k(x_i, x_j)]_{i,j=1}^n \in \mathbb{R}^{n \times n}$$

$$\alpha^* = K^+(y - f),$$

$$h^{*(x)} = k(x)K^+(y - f)$$

#### Neural Tangent Kernel NTK

$$k(\theta)(x, \xi) := \nabla f(\theta)(x) * \nabla f(\theta)(\xi)$$

#### Quadratic loss

$$E(\theta) = \frac{1}{2} \|f(\theta) - y\|^2, y := (y_1, \dots, y_s)^\top$$

#### Gradient flow ODE

$$\dot{\theta} := d\frac{\theta}{dt} = \sum_{i=1}^s (y_i - f_{i(\theta)}) \nabla f_{i(\theta)}$$

#### Functional gradient flow

$$\dot{f}_j = \nabla f_j * \theta = \sum_{i=1}^s (y_i - f_i) k(\theta)(x_i, x_j)$$

$$f = K(\theta)(y - f)$$

#### Infinite width limit

$$w_{ij}^l = \frac{\sigma_w}{\sqrt{m_l}} \varepsilon_{ij}^l,$$

$$b_i^l = \frac{\sigma_b}{\sqrt{m_l}} \beta_i^l,$$

$$\varepsilon_{ij}^l, \beta_i^l \sim N(0, 1)$$

$$k(\theta) \rightarrow k_\infty \text{ for } m_l \rightarrow \infty$$

Initial NTK converges to deterministic limit

#### NTK constancy

$$d\frac{k(\theta(t))}{dt} t = 0$$

$$f_{\infty(x)} = k(x)K^+(y - f), k = k_\infty$$

NTK remains constant when training in infinite width limit

### Vanishing curvature

$$\frac{\|\nabla^2 f(\theta_0)\|_2}{\|\nabla f(\theta_0)\|_2^2} \ll 1$$

### Near-constancy

$$\|k(\theta_0) - k(\theta_t)\|_F \in O\left(\frac{1}{m}\right), m = m_1 = \dots = m_L$$

## 1.9.2 Bayesian DNNs

### Bayesian predictive distribution

$$f(x) = \int f(\theta)(x) p(\theta | S) d\theta$$

#### Bayes rule

$$p(\theta | S) = \frac{p(\theta)p(S | \theta)}{p(S)},$$

$$p(S) = \int p(\theta)p(S | \theta) d\theta$$

#### Parameter priors (Gaussian)

$$p(\theta) = \prod_{i=1}^d p(\theta_i), \theta_i \sim N(0, \sigma_i^2)$$

$$-\log p(\theta) = \frac{1}{2\sigma^2} \|\theta\|^2 + \text{const}$$

Essentially a weight decay term

#### Likelihood (Gaussian noise)

$$-\log p(S | \theta) = \frac{1}{2\gamma^2} \|y - f(\theta)\|^2 + \text{const.}$$

with  $y_i = f^{*(x_i)} + \nu_i, \nu_i \sim N(0, \gamma^2)$

#### Posterior

$$-\log p(\theta | S) = E(\theta) + \text{const},$$

$$E(\theta) = \frac{1}{2\gamma^2} \|y - f\|^2 + \frac{1}{2\sigma^2} \|\theta\|^2$$

#### Bayesian ensembling (post hoc)

$$f(\Theta)(x) = \sum_{i=1}^n \frac{\exp[-E(\theta_i)]}{\sum_{j=1}^n \exp[-E(\theta_j)]} f(\theta_i)(x)$$

Relative posterior weighting

#### Markov chain monte carlo (MCMC)

$$\theta_0, \theta_1, \theta_2, \dots,$$

$$\theta_{t+1} | \theta_t \sim \Pi$$

$$p(\theta_t | S) \Pi(\theta_2 | \theta_1) = p(\theta_2 | S) \Pi(\theta_1 | \theta_2)$$

#### Metropolis-Hastings

$$\Pi(\theta_1 | \theta_2) = \tilde{\Pi}(\theta_1 | \theta_2) A(\theta_1 | \theta_2)$$

$$A(\theta_1 | \theta_2) = \min\left\{1, \frac{p(\theta_1 | S) \tilde{\Pi}(\theta_2 | \theta_1)}{p(\theta_2 | S) \tilde{\Pi}(\theta_1 | \theta_2)}\right\}$$

Modified transition probability with acceptance step  $A$

#### Hamiltonian monte carlo

$$E(\theta) = -\sum_{x,y} \log p(y | x; \theta) - \log p(\theta)$$

$$H(\theta, v) = E(\theta) + \frac{1}{2} v^\top M^{-1} v$$

with  $p(\theta, v)$  propto  $\exp[-H(\theta, v)]$

$$\dot{v} = -E(\theta), \dot{\theta} = v$$

$$\theta_{t+1} = \theta_t + \eta v_t$$

$$v_{t+1} = v_t - \eta \nabla E(\theta_t)$$

#### Langevin dynamics

$$\dot{\theta} = v$$

$$dv = -\nabla E(\theta) dt - B v dt + N(0, 2B dt)$$

$$\theta_{t+1} = \theta_t + \eta v_t$$

$$v_{t+1} = (1 - \eta\gamma)v_t - \eta \int \nabla \tilde{E}(\theta) + \sqrt{2\gamma\eta} N(0, I)$$

## 1.9.3 Gaussian Processes

### Gaussian process

$$(f(x_1), \dots, f(x_s)) \sim N$$

$$\sum_{i=1}^s \alpha_i f(x_i) \sim N, \text{ for all } \alpha \in \mathbb{R}^s$$

### Mean and covariance functions

GPs are completely defined by first and second order statistics

$$\mu(x) := E_{x, \xi}[f(x)]$$

$$k(x, \xi) := E_{x, \xi}[f(x)f(\xi)] - \mu(x)\mu(\xi)$$

$$K_{\mu\nu} = k(x_\mu, x_\nu), K \in \mathbb{R}^{s \times s}$$

### Example kernels

$$k(x, \xi) = x^\top \xi, k(x, \xi) = e^{-\gamma \|x - \xi\|^2}$$

### GP in DNN

Treating parameters as random variables. Each unit in a DNN becomes a random function.

### Linear Layer

$$w \sim N(0, \frac{\sigma^2}{n} I_{n \times n})$$

$$E[y_i y_j] = \frac{\sigma^2}{n} x_i^\top x_j$$

### Deep layers

$$W^{l+1} X^l, l \geq 1$$

No longer normal as products break normality, but near-normal for high dimensional inputs.

### Non-linear activations

$$\mu(x^{l+1}) = E[\varphi(W^l x^l)]$$

### Kernel recursion

$$K_{\mu\nu}^l = E[\varphi(x_{i\mu}^{l-1}) \varphi(x_{i\nu}^{l-1})]$$

$$= \sigma^2 E[\varphi(f_\mu) \varphi(f_\nu)]$$

$$f \sim GP(0, K^{l-1})$$

### Kernel regression

Mean of bayesian predictive distribution

$$f^{*(x)} = k(x)^\top K^+ y$$

$$E[(f(x) - f^{*(x)})^2] = K(x, x) - k(x)^\top K^+ k(x)$$

## 1.9.4 Statistical Learning Theory

### VC learning theory

$$L_t = -\frac{\|m(x_t, x_0, t) - m_{\theta(x_t, t)}\|^2}{2\sigma_t^2} + \text{const.}$$

$$\text{VC-dim}(F) := \max_s \sup_{|S|=s} 1\|F(S)\| = 2^s$$

### VC inequality

$$P(\sup_F |\hat{E}(f) - E(f)| > \varepsilon) \leq 8 |F(s)| e^{-s \frac{\varepsilon^2}{32}}$$

### Double descent

Beyond the interpolation point, models start to learn and eventually may level out at a lower generalization error.

### Generalization gap

$$\Delta := \max(0, E - \hat{E})$$

$E$ : expected population error,  $\hat{E}$ : empirical error

## KL divergence

$$D_{\text{KL}}(p \parallel q) = \int p(x) \log\left(\frac{p(x)}{q(x)}\right) dx = E_{x \sim p} [\ln\left(\frac{p(x)}{q(x)}\right)]$$

## PAC-Bayesian theorem

For fixed  $E$  and any  $Q$  over  $s$  samples:

$$E_Q[E(f)] - E_Q[\hat{E}(f)] \leq \sqrt{\frac{2}{s} [\text{KL}(Q \parallel P) + \ln\left(\frac{1}{2\sqrt{s}\varepsilon}\right)]}$$

Ensures general rate  $\tilde{O}\left(\frac{1}{\sqrt{s}}\right)$

## PAC-Bayesian bound

$$Q := N(\theta, \text{diag}(\sigma_i^2))$$

$$\text{KL}(Q \parallel P) = \sum_i \log\left(\frac{\lambda_i}{\sigma_i} + \frac{\sigma_i^2 + \theta_i^2}{2\lambda_i^2} - \frac{1}{2}\right)$$

$$E_{\text{PAC}}(Q) :=$$

$$E_Q[\hat{E}] + \sqrt{\frac{2}{s} [\text{KL}(Q \parallel P) + \ln\left(\frac{1}{2\sqrt{s}\varepsilon}\right)]}$$

Favours minima robust to parameter perturbations

## PAC-bayesian learning implementation

$$\theta_{t+1} = \theta_t - \eta \nabla E_Q[\hat{E}] = \theta_t - \eta \nabla \hat{E}(\hat{\theta}),$$

with  $\hat{\theta} \sim Q(\theta, \sigma)$

Gradient loss on perturbed parameters

## Reparameterization trick

$$\hat{\theta} = \theta + \text{diag}(\sigma_i)\varepsilon, \varepsilon \sim N(0, I)$$

Backpropagation to  $\theta$  and  $\sigma_i$

## 1.10 Generative Models

### 1.10.1 Variational Autoencoders

#### Linear autoencoder

$$x \mapsto z = Cx, C \in \mathbb{R}^{m \times n}$$

$$z \mapsto \hat{x} = Dz, D \in \mathbb{R}^{n \times m}$$

$$E(C, D)(x) = \frac{1}{2}\|x - \hat{x}\|^2 = \frac{1}{2}\|x - DCx\|^2$$

$$DCX = \hat{X} = U\Sigma_m V^\top$$

$$\Sigma_m = \text{diag}(\sigma_1, \dots, \sigma_m, 0, \dots, 0)$$

For centered data equivalent to PCA, but generally has non-global minima

#### Linear factor analysis

##### Probability Model

$$p_{X(z)} = \int p_{Z(z)} p_{X|Z}(x \mid z) dz$$

$Z$ : latent variables,  $X$ : observed variables

#### Linear observation model

$$x = \mu + Wz + \nu \text{ with } \nu \sim N(0, \Sigma)$$

$$x \sim N(\mu, WW^\top + \Sigma) \text{ for } z \sim N(0, I)$$

#### Posterior mean and covariance

$$\mu_{z|x} = W^\top(WW^\top + \Sigma)^{-1}(x - \mu)$$

$$\Sigma_{z|x} = I - W^\top(WW^\top + \Sigma)^{-1}W$$

#### Pseudoinverse limit

$$W^\top(WW^\top + \sigma^2 I)^{-1} \rightarrow W^+ \in \mathbb{R}^{m \times n}$$

$$\mu_{z|x} \rightarrow W^+(x - \mu), \Sigma_{z|x} \rightarrow 0$$

## Maximum likelihood estimation

$$\mu, W \max \rightarrow \log p_{\mu, W}(S)$$

## Optimality condition for $W$

$$w_i = \rho_i u_i, \rho_i = \max\{0, \sqrt{\lambda_i - \sigma^2}\}$$

With  $(\lambda_i, u_i)$  eigenvalues and eigenvectors of covariance matrix.

For  $\sigma = 0$  equivalent to PCA.

## Variational autoencoder (VAE)

$$z \sim N(0, I)$$

$$x = F(\theta)(z) = (F^L @ \dots @ F^1)(z)$$

## Evidence lower bound (ELBO)

$$\begin{aligned} \log p_\theta(x) &= \log \int p_\theta(x \mid z)p(z) dz \\ &= \log \int q(z) \left[ \frac{p_\theta(x \mid z)p(z)}{q(z)} \right] dz \\ &\geq \int q(z) \log p_\theta(x \mid z) dz - \int q(z) \log\left(\frac{q(z)}{p(z)}\right) dz \\ &=: L(\theta, q)(x) \\ \theta \max &\rightarrow L(\theta, q)(S) = \sum_{i=1}^s L(\theta, q)(x_i) \end{aligned}$$

## Inference network

$$z \sim N(\mu(x), \Sigma(x))$$

$$z = \mu + \Sigma^{\frac{1}{2}}\varepsilon, \varepsilon \sim N(0, I)$$

$$\nabla_\mu E[f(z)] = E[\nabla_z f(z)]$$

$$\nabla_\Sigma E[f(z)] = \frac{1}{2}E[\nabla_z^2 f(z)]$$

Integration by parts derivation

### 1.10.2 Generative Adversarial Networks

#### GAN objective

$$V(G, D) = E_{x_r \sim p_{\text{data}}} D(x_r) + E_{z \sim p_z} (1 - D(G(z)))$$

#### Discriminator Mixture Model

$$\tilde{p}_{\theta(x,y)} = \frac{1}{2}(yp(x) + (1-y)p_{\theta(x)}),$$

$$y \in \{0, 1\},$$

$p$ : true probability,  $p_\theta$ : model probability

#### Bayes-optimal classifier

$$q_{\theta(x)} := P\{y = 1 \mid x\} = \frac{p(x)}{p(x) + p_{\theta(x)}}$$

To detect fake samples,  $y = 1$  for real samples,  $y = 0$  for fake samples

#### Logistic likelihood

$$\theta \min \rightarrow \ell^{*(\theta)} := E_{\tilde{p}_\theta}[y \ln q_{\theta(x)} + (1-y) \ln(1 - q_{\theta(x)})]$$

#### Jensen-Shannon as effective objective

$$\begin{aligned} \ell^* &= E_{\tilde{p}_\theta}[y \ln q_{\theta(x)} + (1-y) \ln(1 - q_{\theta(x)})] \\ &= -\frac{1}{2}H(p) - \frac{1}{2}H(p_\theta) + H\left(\frac{1}{2}(p + p_\theta)\right) - \ln 2 \\ &= \text{JS}(p, p_\theta) - \ln 2. \end{aligned}$$

#### Discriminator model

$$q_\varphi : x \mapsto [0; 1], \varphi \in \Phi$$

#### Objective bounds

$$\ell^{*(\theta)} \geq \sup_{\varphi \in \Phi} \ell(\theta, \varphi)$$

$$\ell(\theta, \varphi) := E_{\tilde{p}_\theta}[y \ln q_{\varphi(x)} + (1-y) \ln(1 - q_{\varphi(x)})]$$

## Saddle point optimization

$$\theta^* := \operatorname{argmin}_{\theta \in \Theta} (\sup_{\varphi \in \Phi} \ell(\theta, \varphi))$$

$\varphi$ : Generator,  $\theta$ : Discriminator

## Alternating gradient descent/ascent

$$\theta_{t+1} = \theta_t - \eta \nabla_\theta \ell(\theta_t, \varphi_t)$$

$$\varphi_{t+1} = \varphi_t + \eta \nabla_\varphi \ell(\theta_{t+1}, \varphi_t)$$

## Extra-gradient steps

$$\theta_{t+1} = \theta_t - \eta \nabla_\theta \ell(\theta_{t+0.5}, \varphi_t)$$

with  $\theta_{t+0.5} := \theta_t - \eta \nabla_\theta \ell(\theta_t, \varphi_t)$

$$\varphi_{t+1} = \varphi_t + \eta \nabla_\varphi \ell(\theta_t, \varphi_{t+0.5})$$

with  $\varphi_{t+0.5} := \varphi_t + \eta \nabla_\varphi \ell(\theta_t, \varphi_t)$

## Deconvolutional DNN

Upside-down ConvNet for image generation

### 1.10.3 Denoising Diffusion

#### Markov chains

$$x_{0:t-1} \perp x_{t+1:\infty} \mid x_t \text{ (forall } t)$$

$$p(x_t \mid x_{t-1}) = p(x_1 \mid x_0) \text{ (forall } t),$$

$$p(x_{s:t}) = p(x_t) \prod_{\tau=s+1}^t p(x_{\tau-1} \mid x_\tau)$$

$$p(x_{s:t}) = p(x_s) \prod_{\tau=s+1}^t p(x_\tau \mid x_{\tau-1}),$$

$$\pi(x_{t+1}) = \int \pi(x_t) p(x_{t+1} \mid x_t) dx_t$$

#### Denoising diffusion

##### Forward (noise generation)

$$\pi^* = \nu_0 \mapsto \nu_1 \mapsto \dots \mapsto \nu_{T-1} \mapsto \nu_T = \pi$$

##### Backward (denoising)

$$\pi = \mu_T^{\theta} \mapsto \mu_{T-1}^{\theta} \mapsto \dots \mapsto \mu_1^{\theta} \mapsto \mu_0^{\theta} \approx \pi^*$$

#### Gaussian example

$$\pi \approx N(0, I),$$

$$x_t \mid x_{t-1} \sim N(\sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$

#### Forward SDE

$$dx_t = -\frac{1}{2}\beta_t x_t dt + \sqrt{\beta_t} dw_t$$

#### Backward SDE

$$dx_t = -\frac{1}{2}\beta_t x_t -$$

$$\beta_t \nabla_{x_t} \log q_{t(x_t)} dt + \sqrt{\beta_t} d\langle \omega \rangle_t$$

score · wiener process

#### ELBO bound

$$x_t = \sqrt{1 - \beta_t}x_{t-1} + \sqrt{\beta_t}\varepsilon_t, \varepsilon_t \sim N(0, I)$$

$$\ln p_{\theta(x_0)} = \ln \int q(x_{1:T} \mid x_0) \left( \frac{p_{\theta(x_0:T)}}{q(x_{1:T} \mid x_0)} \right) dx_{1:T}$$

$$\geq E\left[\ln\left(\frac{p_{\theta(x_0:T)}}{q(x_{1:T} \mid x_0)}\right) \mid x_0\right]$$

$$= \sum_{t=0}^T L_t$$

$$L_t := \text{cases}\left(E\left[\ln p_{\theta(x_0 \mid x_1)}\right], t = 0; -D(q(x_T \mid x_0) \parallel \pi), t = T; -D(q(x_{t-1} \mid x_t, x_0) \parallel p_{\theta(x_{t-1} \mid x_t)}), \text{else}\right)$$

#### Backward model assumption

$$x_{t-1} \mid x_t \sim N(m(x_t, t), \Sigma(x_t, t))$$

#### Entropy bounds

$$H(x_t) \geq H(x_{t-1}) \Rightarrow H(x_t \mid x_{t-1}) \geq H(x_{t-1} \mid x_t)$$

## Noise schedules

$$\begin{aligned} |(\alpha)_t| &= \prod_{\tau=1}^t (1 - \beta_\tau), |(\beta)_t| = 1 - |(\alpha)_t| \\ x_t &\approx N\left(\sqrt{|(\alpha)_t}x_0, |(\beta)_t|I\right) t \rightarrow \infty \rightarrow N(0, I) \end{aligned}$$

## Forward trajectory target

$$\begin{aligned} x_{t-1} \mid x_t, x_0 &= N(m(x_t, x_0, t), \tilde{\beta}_t I) \\ m(x_t, x_0, t) &= \left(\frac{\sqrt{|(\alpha)_{t-1}| \beta_t}}{1 - |(\alpha)_t|}\right)x_0 + \left(\frac{(1 - |(\alpha)_{t-1}|)\sqrt{1 - \beta_t}}{1 - |(\alpha)_t|}\right)x_t \\ \text{with } \tilde{\beta}_t &= \frac{1 - |(\alpha)_{t-1}| \beta_t}{1 - |(\alpha)_t|} \end{aligned}$$

## Fixed isotropic covariance

$$\Sigma(x_t, t) = \sigma_t^2 I, \text{ where } \sigma_t^2 \in \{\beta_t, \tilde{\beta}_t\}$$

## Simplified ELBO

$$L_t = -\frac{\|m(x_t, x_0, t) - m_{\theta(x_t, t)}\|^2}{2\sigma_t^2} + \text{const.}$$

## Reparameterization

$$\begin{aligned} x_t &= \sqrt{|(\alpha)_t}x_0 + \sqrt{1 - |(\alpha)_t|}\varepsilon \Rightarrow x_0 = \frac{1}{\sqrt{|(\alpha)_t}}x_t - \frac{\sqrt{1 - |(\alpha)_t|}\varepsilon}{\sqrt{|(\alpha)_t}} \\ m(x_t, x_0, t) &= \frac{1}{\sqrt{\alpha_t}} \left[ x_{t(x_0, \varepsilon)} - \frac{\beta_t}{\sqrt{1 - |(\alpha)_t|}}\varepsilon \right] \\ \text{with } \varepsilon &\sim N(0, I) \end{aligned}$$

## Expected squared error

$$\begin{aligned} E_{q[L_t \mid x_0]} &= E_{\varepsilon \sim \rho_t} \left[ \left\| \varepsilon - \varepsilon_{\theta(\sqrt{|(\alpha)_t}x_0 + \sqrt{1 - |(\alpha)_t|\varepsilon}, t)} \right\|^2 \mid x_0 \right] \\ \text{with } \rho_t &= \frac{\beta_t^2}{2\sigma_t^2 \alpha_t (1 - |(\alpha)_t)} \end{aligned}$$

## Final simplified criterion

$$\begin{aligned} h(\theta)(x) &= \\ \frac{1}{T} \sum_{t=1}^T E\left[\left\| \varepsilon - \varepsilon_{\theta(\sqrt{|(\alpha)_t}x + \sqrt{1 - |(\alpha)_t|\varepsilon}, t)} \right\|^2\right] & \end{aligned}$$

## 1.11 Ethics

### 1.11.1 Adversarial Examples

#### Adversarial perturbation

$$f(x + \nu) \neq f(x) \text{ s.t. } \|\nu\|_p \leq \varepsilon$$

#### p-norm definitions

$$\|x\|_p = \left(\sum_i |x_i|^p\right)^{\frac{1}{p}}$$

$$\|x\|_\infty = \max_i |x_i|, \|x\|_0 = |\{i : x_i \neq 0\}|$$

#### Optimal perturbation (linear binary classification)

$$\nu \propto \text{sign}(f_1(x) - f_2(x))(w_2 - w_1)$$

for  $f_i = w_i^\top x + b_i$

#### Optimal perturbation (multiclass)

$$\nu = \operatorname{argmin}_{i>1} \frac{f_1(x) - f_i(x)}{\|w_1 - w_i\|_2^2} (w_i - w_1)$$

#### DeepFool iterative optimization

Iterate:  $\operatorname{argmin}_{\Delta\nu} \|\Delta\nu\|_2$  s.t.

$$(\nabla f_1(x) - \nabla f_2(x))^\top \Delta\nu < f_1(x) - f_2(x)$$

#### Robust training

$$\ell(f(x), y) \rightarrow \max_{\nu: \|\nu\|_p \leq \varepsilon} \ell(f(x + \nu), y)$$

**Projected gradient ascent ( $p = 2$ )**

$$\nu_{t+1} = \varepsilon \Pi[\nu_t + \alpha \nabla_x \ell(f(x + \nu_t), y)]$$
$$\Pi[z] := \frac{z}{\|z\|_2}$$

**Projected gradient ascent ( $p = \infty$ )**

$$\nu_{t+1} = \varepsilon \Pi[\nu_t + \alpha \operatorname{sign}(\nabla_x \ell(f(x + \nu_t), y))]$$
$$\Pi[z] := \frac{z}{\|z\|_\infty}$$

**Fast Gradient Sign Method (FGSM)**

$$\nu = \varepsilon \operatorname{sign}(\nabla_x \ell(f(x), y))$$

## 2 Computer Vision

### 2.1 The Digital Image & Sensors

#### Charge Coupled Device (CCD)

Photons

- Blooming:** Oversaturated photosites cause vertical channels to "flood" (bright vertical line)

#### Image Noise

Additive Gaussian noise:

#### Color camera concepts:

- Prism (split light, 3 sensors, needs good alignment, good color separation)

## 2.2 Image Segmentation

Pixel-wise classification problem, to group pixels in an image that share common properties.

Segmentation of  $I$ : Find  $R_1, \dots, R_n$  such that  $I = \bigcup_{i=1}^N R_i$  with  $R_i \cap R_j = \emptyset \quad \forall i \neq j$ .

#### Thresholding

Segment image into 2 classes.  
 $B(x, y) = 1$  if  $I(x, y) \geq T$  else 0, finding  $T$  with trial and error, compare results with ground truth.

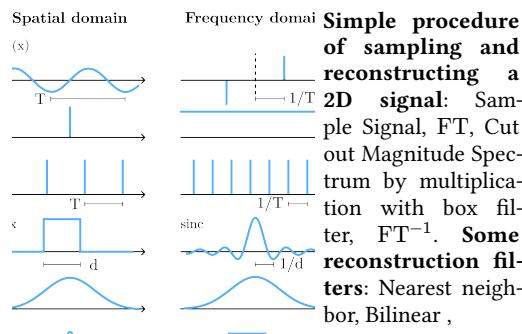
#### Important Kernels

	Low-pass/ High-pass		
Laplacian	Prewitt <sub>x</sub>	Mean / Box	High-pass
Gaussian	Sobel <sub>x</sub>	Diff <sub>x</sub>	Diff <sub>y</sub>
$\frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$	mat(-1, 0, 1; mat([0, -2], 1), mat([-1], 1)) <sup>T</sup>		

**Dirac delta:**  $\delta(x) = \text{cases}(0 \text{ if } x \neq 0, \text{undefined else})$  with  $\int_{-\infty}^{\infty} \delta(x) dx = 1$ .  $\mathcal{F}[\delta(x - x_0)](u) = e^{-i2\pi u x_0}$ .  $\delta(u) = \int_{\mathbb{R}} e^{-i2\pi x u} dx$ .

**Sampling**  $f$  at points  $x_n$ :  $f_s(x) = \sum_n f(x_n) \delta(x - x_n)$ .

**Property**  $f(x)$   $F(u)$   
 Linearity  $\alpha f_1(x) + \beta f_2(x)$   $\alpha F_1(u) + \beta F_2(u)$   
 Duality  $F(x)$   $f(-u)$



Gaussian reconstruction filter (equiv. to convolving sampled signal w/ Gaussian kernel.  $\text{sinc}(x) = \frac{\sin(\pi x)}{\pi x}$ )

#### Image restoration

Image degradation is applying kernel  $h$  to some image  $I$ . The inverse  $\tilde{h}$  should compensate:  $I \xrightarrow{h(x)} J \xrightarrow{\tilde{h}(x)} I$ .

Determine with  $\mathcal{F}[\tilde{h}](u, v) \cdot \mathcal{F}[h](u, v) = 1$ . Or  $\tilde{h} = \mathcal{F}^{-1}\left[\frac{1}{\mathcal{F}[h]}\right]$

Cancellation of frequencies & noise amplification → Regularize using  $\tilde{\mathcal{F}}[\tilde{h}](u, v) = \mathcal{F}[h] / (|\mathcal{F}[h]|^2 + \varepsilon)$ .

**Motion blur:**  $h(x, y) = \frac{1}{2l} [\theta(x + l) - \theta(x - l)] \delta(y)$