

1 Connectionism

1.1 Perceptron

Threshold Unit $f[w, b](x) = \text{sign}(x \cdot w + b)$ with **Decision Boundary** $x \cdot w + b \stackrel{!}{=} 0$. $-\frac{b}{\|w\|}$ is the signed distance of the hyperplane from 0.
Geometric Margin $\gamma[w, b](x, y) = \frac{y(x \cdot w + b)}{\|w\|}$.
Maximum Margin Classifier $(w^*, b^*) \in \arg\max_{w, b} \gamma[w, b](\mathcal{S})$, with margin $\gamma[w, b](\mathcal{S}) := \min_{(x, y) \in \mathcal{S}} \gamma[w, b](x, y)$.

Perception Learning

if $f[w, b](x) \neq y$: $w \leftarrow w + yx, b \leftarrow b + y$.
Aims to find some solution; assumes version space is non-empty. Does not aim for small err. $w_0 \in \text{span}(x_1, \dots, x_s) \Rightarrow w_t \in \text{span}(x_1, \dots, x_s) \forall t$

Convergence - Novikoff's Theorem

- If $\exists w^*, \|w^*\| = 1$, s.t. $\gamma[w^*](\mathcal{S}) = \gamma > 0 \Rightarrow w_t \cdot w^* \geq \gamma t$.
- Let $R = \max_{x \in \mathcal{S}} \|x\|$. Then $\|w_t\| \leq R\sqrt{t}$.
 $\cos \angle(w^*, w_t) = \frac{w^* \cdot w_t}{\|w^*\| \|w_t\|} \geq \frac{\gamma t}{\sqrt{tR}} = \sqrt{t}\frac{\gamma}{R} \leq 1 \Rightarrow t \leq \frac{R^2}{\gamma^2}$.

Cover's Theorem for $\mathcal{S} \subset \mathbb{R}^n$, $|\mathcal{S}| = s$
 $C(\mathcal{S}, n)$: # of ways to separate \mathcal{S} in n dimensions. Position of pts does not matter (as long as they are in general position). $C(s+1, n) = 2 \sum_{i=0}^{n-1} \binom{s}{i}$, $C(s, n) = 2^s$ for $s \leq n$. Phase transition at $s = 2n$. For $s < 2n$ empty version space is the exception, otherwise the rule.

1.2 Hopfield Networks

$$E(X) = -\frac{1}{2} \sum_{i \neq j} w_{ij} X_i X_j + \sum_i b_i X_i, \quad \text{where } X_i \in \{\pm 1\}. \quad w_{ij} = w_{ji}, w_{ii} = 0.$$

Hebbian Learning
Choose patterns $\{x^t\}_{t=1}^s \in \{\pm 1\}^n$, compute weights once using them: $w_{ij} = \frac{1}{n} \sum_{t=1}^s x_i^t x_j^t$; $w_{ii} = 0$. For inference, update X iteratively: $X_{i+1}^t = \text{sign} \left(\sum_j w_{ij} X_j^t + b_i \right)$ asynchronously. Capacity for random, uncorrelated patterns: $s_{\max} \approx 0.138n$. Requiring pattern to be retrieved with high probability: $s \leq \frac{n}{2 \log_2 n}$.

If $X = \text{diag}(1, \dots, 1)$, no reconstruction happens. Under async update step, any Hopfield network is guaranteed to converge.

2 Feedforward Networks

2.1 Linear Models

Linear regression (MSE) $L[w](X, y) = \frac{\|Xw - y\|^2}{2n}$, $\nabla_w L = \frac{X^\top Xw - X^\top y}{n}$.

Moore-Penrose inverse solution

$w^* = X^*y \in \arg\min_w L[w](X, y)$, where $X^* = \lim_{\delta \rightarrow 0} (X^\top X + \delta I)^{-1} X^\top$ Moore-Penrose inv.

Stochastic gradient descent update $w_{t+1} = w_t + \eta (y_{i_t} - w_t^\top x_{i_t}) x_{i_t}, i_t \sim \mathcal{U}([1, n])$.

Gaussian noise model $y_i = w^\top x_i + \varepsilon_i$, $\varepsilon_i \sim \mathcal{N}(0, \sigma^2)$, LSQ equivalent to NLL of gaussian noise model.

Ridge regression $h_\lambda[w] = h[w] + \frac{\lambda}{2} \|w\|^2$, $w^* = (X^\top X + \lambda I)^{-1} X^\top y$.

Logistic function

$$\sigma(z) = \frac{1}{1+e^{-z}}, \sigma(z) + \sigma(-z) = 1. \quad \sigma' = \sigma(1-\sigma), \sigma'' = \sigma(1-\sigma)(1-2\sigma)$$

Cross entropy loss for $y \in \{0, 1\}$

$$\ell(y, z) = -y \log \sigma(z) - (1-y) \log(1-\sigma(z)) = -\log \sigma(2y-1)z.$$

Logistic regression with Cross Entropy loss:

$$L[w] = \frac{1}{n} \sum_{i=1}^n \ell_i(y_i, w^\top x_i), \quad \nabla \ell_i = [\sigma(w^\top x_i) - y_i] x_i.$$

2.2 Feedforward Networks

Generic feedforward layers

$$F : \underbrace{\mathbb{R}^{m(n+1)} \times \underbrace{\mathbb{R}^n}_{\text{parameters}}}_{\text{Input}} \rightarrow \underbrace{\mathbb{R}^m}_{\text{Output}}, F[\theta](x) = \varphi(Wx + b).$$

Softmax $(z)_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$, CE-loss in terms of logits $\ell(y, z) = \frac{1}{\ln(2)} \left[-y_z + \ln \sum_j e^{z_j} \right]$. CE between two pmfs: $l(p; q) = -\sum_i p_i \log q_i$. CE with hard labels is NLL-loss.

$$\frac{d}{dx_j} \text{softmax}(x)_i = \text{sm}(x)_i (\delta_{ij} - \text{sm}(x)_j)$$

Residual layer $F[W, b](x) = x + (\varphi(Wx + b) - \varphi(0))$, therefore $F[0, 0] = \text{id}$. Link that propagates x forward is called a **skip connection**. Composing residual layers: number of paths grows exponentially, can include projections for flexibility of changing dimensionality.

2.3 Sigmoid Networks

Sigmoid activation $\sigma(z) = \frac{1}{1+e^{-z}}$.

Hyperbolic tangent activation $\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}} = 2\sigma(2z) - 1$. $\tanh'(z) = 1 - \tanh^2(z)$.

Smooth function approximation

Polynomials, ridge functions (functions of the form $\varphi(a^\top x + b)$) and MLPs with C^∞ activations are universal approximators.

Weierstrass: Polynomials are universal approximators of $C(\mathbb{R})$ on any given compact I .

Barron's Theorem: For $f : \mathbb{R}^d \rightarrow \mathbb{R}$ with $C_f = \int \|\omega\| |\hat{f}(\omega)| d\omega < \infty$, \exists width- m MLP g_m s.t.: $\int_B |f - g_m|^2 dx \leq O(\frac{1}{m})$.

2.4 ReLU(z) = max(0, z) networks

Zalavsky's Thoerem: Activation patterns

m ReLU neurons in \mathbb{R}^n . Each neuron's hyperplane $\{w_i^\top x = 0\}$ partitions \mathbb{R}^n into $R(m)$ connected regions of constant activation pattern. $R(m) \leq \sum_{i=0}^{\min(n, m)} \binom{m}{i} \ll 2^m$.

Montufar: Connected linear regions in ReLU network $R(m, L) \geq R(m) \lfloor \frac{m}{n} \rfloor^{n(L-1)}$, L : layers, m : width.

Shekhtman: Piecewise linear functions are dense in $C([0, 1])$. **Lebesgue**: Piecewise linear function with m pieces can be written $g(x) = ax + b + \sum_{i=1}^{m-1} c_i (x - x_i)_+$; $m+1$ parameters, a, b, c_i .

ReLU networks with 1 hidden layer are universal approximators.

Wang and Sun: Every continuous piecewise linear function $g : \mathbb{R}^n \rightarrow \mathbb{R}$ can be written as a signed sum of k -Hinges with $k \leq n+1$. A k -Hinge is a function $g(x) = \max_{j=1}^k \{w_j^\top x + b_j\}$, generalizes ReLU, known as Maxout unit.

Linear Autoencoder: Optimal $A = DE$, s.t. frobenius norm reconstruction err of AX is minimized, is $D = U_k$, $E = U_k^\top$, not jointly convex in E and D , but individually. $\hat{X}^* = \arg\min_{\text{rank } (\hat{X})=k} \|X - \hat{X}\|_F^2 = U \Sigma_k V^\top$ (SVD).

If f μ -strongly convex, L -smooth, GD iterates x_t with $0 < \eta \leq \frac{1}{L}$ converge to unique minimizer x^* at rate $\|x_t - x^*\|^2 \leq (1-\eta\mu)^t \|x_0 - x^*\|^2$. If f convex, diff and L -smooth, with $\eta \leq \frac{1}{L}$, $f(x_t) - f(x^*) \leq \frac{1}{2\eta t} \|x_0 - x^*\|^2$. **Non-Convex case**: If f diff, L -smooth, with minimum f^* , GD iterates with $\eta \leq \frac{1}{L}$ satisfy $\min_{i=0}^t \|\nabla f(x_i)\|^2 \leq \frac{2(f(x_0) - f^*)}{\eta(t+1)}$.

If f diff and L -smooth: $f(x) - f(x^*) \geq \frac{1}{2L} \|\nabla f(x)\|^2$.

Polyak-Lojasiewicz condition

$$\frac{1}{2} \|\nabla f(x)\|^2 \geq \mu(f(x) - \min f) \quad (\forall x). \quad \mu\text{-strong convex} \Rightarrow \mu\text{-PL}.$$

GD Convergence Rates & Learning Rates

L-smooth only: $\eta^* = \frac{1}{L}$. To reach ε -stationary point ($\|\nabla f\| \leq \varepsilon$) needs at most $\frac{2L}{\varepsilon^2} (f(x_0) - \min f)$ steps.

μ -PL + L-smooth: Use $\eta^* = \frac{2}{L+\mu}$. Convergence: $f(x_t) - f(x^*) \leq (1 - \frac{\mu}{L})^t (f(x_0) - f(x^*))$.

3 Gradient-Based Learning

Forward mode is more memory efficient, but backward mode is more runtime efficient. Fwd is $O(\# \text{params})$, reverse is $O(d_{\text{out}})$.

Numerator layout: For $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, $(\frac{\partial y}{\partial x})_{ij} = \frac{\partial y_i}{\partial x_j} \in \mathbb{R}^{m \times n}$ and $f : \mathbb{R}^{n_1 \times n_2} \rightarrow \mathbb{R}$, $\nabla f(\hat{X})_{ij} = \frac{\partial f}{\partial \hat{X}_{ij}} \in \mathbb{R}^{n_1 \times n_2}$.

3.1 Backpropagation

$$x^\ell = \varphi(W^\ell x^{\ell-1} + b^\ell), \quad \frac{\partial \mathcal{L}}{\partial W^\ell} = \delta^\ell (x^{\ell-1})^\top, \quad \delta^\ell = \delta^\ell \circ \varphi'(W^\ell x^{\ell-1} + b^\ell), \quad \frac{\partial \mathcal{L}}{\partial x^\ell} = (W^{\ell+1})^\top \delta^{\ell+1}.$$

$\frac{\partial L}{\partial x} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial x}$. For $x \in \mathbb{R}^n$ different z : $\frac{\partial (Wx)}{\partial x} = W$, element wise f gives: $\frac{\partial f(x)}{\partial x} = \text{diag}(f'(x))$, $\frac{\partial \|y\|^2}{\partial y} = 2(\hat{y} - y)^\top, \frac{\partial L}{\partial y} = h \cdot \frac{\partial L}{\partial \hat{x}}$.

3.2 Gradient Descent

Update: $x_{t+1} = x_t - \eta \nabla f(x_t)$.

Gradient flow ODE $\frac{dx}{dt} = -\nabla f(x)$ gives ideal trajectory to be approximated by GD.

Newton's method gives optimal step for quadratic model: $\Delta x = -[\nabla^2 f(x)]^{-1} \nabla f(x)$. $\nabla_x^2 [x^\top Ax + b^\top x + c] = A + A^\top$.

Optimal LR for Convex Quadratics

For $f(x) = \frac{1}{2} x^\top Qx$, $\eta^* = \frac{2}{\lambda_{\max}(Q) + \lambda_{\min}(Q)}$. Stability requires $\eta \leq \frac{2}{\lambda_{\max}(Q)}$. Quadratic approx.

of f : $f(x + \Delta x) \approx f(x) + \nabla f(x)^\top \Delta x + \frac{1}{2} \Delta x^\top \nabla^2 f(x) \Delta x$. Condition number of Q : $\frac{\lambda_{\max}}{\lambda_{\min}}$

L-smooth: $\|\nabla f(x) - \nabla f(y)\| \leq L \|x - y\|$.

Equivalently (if f twice diff) (for $L = 0 \Rightarrow$):

$$f(y) \leq f(x) + \nabla f(x)^\top (y - x) + \frac{L}{2} \|y - x\|^2.$$

Implies $\lambda_i(\nabla^2 f(x)) \leq L$ for all EVs λ_i of $\nabla^2 f(x)$.

Convexity ($\lambda \in [0, 1]$):

$$f(\lambda w + (1-\lambda)w') \leq \lambda f(w) + (1-\lambda)f(w').$$

μ -Strong convexity: ($\mu = 0 \Leftrightarrow$ convex + diff)

$$\Leftrightarrow f(y) \geq f(x) + \nabla f(x)^\top (y - x) + \frac{\mu}{2} \|y - x\|^2.$$

Implies $\lambda_i(\nabla^2 f(x)) \geq \mu$ for all EVs λ_i of $\nabla^2 f(x)$.

For $f : \mathbb{R} \rightarrow \mathbb{R}$, these become: $L \geq f''(x) \geq \mu \quad \forall x$.

signSGD: $\theta_{t+1} = \theta_t - \eta \text{ sign}(\nabla f_{i_k}(\theta_t))$.

x^4 is strictly convex but not strongly convex, since near 0 the growth of x^4 is slower than x^2 , violating the uniform lower bound on curvature.

With f L -smooth and μ -PL, GD with optimal step-size $\arg\min_\eta f(\theta_t - \eta \nabla f(\theta_t))$ converges globally at linear rate.

$f(w) = \frac{\|Xw - y\|^2}{2} + \lambda \|w\|^2$ satisfies PL-condition.

Muon: Orthogonalize gradient, should increase the scale of other rare directions which have small magnitude in update but are important.

$\Delta W = -\nabla L(W) \star \frac{d_{\text{out}}}{d_{\text{in}}} \cdot U^\top V$ ($\nabla L(W) = U\Sigma V^\top$) minimizes RHS of: $L(W + \Delta W) \leq L(W) + \langle \nabla W L(W), \Delta W \rangle_F + \frac{1}{2} \frac{d_{\text{out}}}{d_{\text{in}}} \|\Delta W\|_F^2$. GD Trajectory always orthogonal to level set.

4 Convolutional Networks

Convolution ($f * g)(u) = (g * f)(u)$
 $= \int_{-\infty}^{\infty} f(t)g(u-t) dt = \int_{-\infty}^{\infty} f(u-t)g(t) dt$.
 $(f * g)$ = Toeplitz-Matrix(g).
Fourier transform convolution property
 $\mathcal{F}(f * g) = \mathcal{F}(f) \cdot \mathcal{F}(g)$.

Cross-correlation:
 $(g * f)[u] = \sum_t g[t]f[u+t]$.

$$(f(t) * g(t))(y) = (f(-t) * g(t))(y).$$

$$(f(t) * g(t))(-y) = (f(-t) * g(t))(y) = (f(t) * g(-t))(y).$$

Translation-Equivariant Operators = Convolutions.

For $x \in \mathbb{R}^D, h \in \mathbb{R}^K$, $x * h = W_h x$. **Toeplitz Matrix**: $W_h \in \mathbb{R}^{(D-K+1) \times D}$
 $(W_h)_{k,j} = \begin{cases} h_{k+j-k} & \text{if } k \leq j \leq k+K-1 \\ 0 & \text{otherwise} \end{cases}$ for $h \in \mathbb{R}^K$.

$$\nabla_w v^\top \text{vec}(\sigma(x * w)) = \text{flip}[x] * [\text{mat}(v) \odot \sigma'(x * w)], \text{ flip rows and columns.}$$

Normal convolution of image $x \in \mathbb{R}^{d \times d}$, kernel $w \in \mathbb{R}^{q \times q} x * w$ requires $\mathcal{O}((d-q)^2 q^2)$. If w separable st $w = uv^\top$, $u \in \mathbb{R}^q, v \in \mathbb{R}^p, \mathcal{O}(dq(d-q))$

Output size: $H_{\text{out}} = \lfloor \frac{H_{\text{in}} + 2P - K}{S} \rfloor + 1$, Height, Input size, padding, Kernel size, stride.

For 2D convolution $y = x * w$, $\frac{\partial L}{\partial w} = x * \frac{\partial L}{\partial y}$

4.1 Convolutional Networks

Conventions for Padding: Add 0 around input.

ConvNets for Images (r out channel, u in ch)
 $y[r][s, t] = \sum_u \sum_{i,j} w[r, u][i, j] * x[u][s+i, t+j]$.

parameters of a convolutional layer
 $(|r| \times |u|) \cdot (|i| \times |j|)$: fully connected \times patch-size.

4.2 Word2Vec

Per word ω , have input embedding x_ω and output embedding y_ω . Predict context word ν given center word ω : $P(\nu | \omega) = \frac{\exp(x_\omega^\top y_\nu)}{\sum_\mu \exp(x_\omega^\top y_\mu)}$.

NLL loss: $\ell_{\omega, \nu} = -x_\omega^\top y_\nu + \ln \sum_\mu \exp(x_\omega^\top y_\mu)$.

Total: $h(\{x_\omega\}, \{y_\nu\}) = \sum_{(\omega, \nu)} \ell_{\omega, \nu}$ over observed pairs. Use only input embeddings after training.

5 Geometric Deep Learning

Group is set G with a binary operation s.t.:
 1) $(gh)f = g(hf)$, 2) $\exists e \in G$ s.t. $ef = fe = f$,
 3) $\forall g \exists g^{-1} \in G$ s.t. $gg^{-1} = g^{-1}g = e$, 4) $gh \in G \forall g, h$. **Abelian** if $gh = hg$.

5.1 Sets and Points

Order-invariance property:

$$f(x_1, \dots, x_M) = f(x_{\pi(1)}, \dots, x_{\pi(M)}) \text{ (perturbs).}$$

(Permutation) Equivariance property:

$$f(x_1, \dots, x_M) = (y_1, \dots, y_M) \Rightarrow f(x_{\pi(1)}, \dots, x_{\pi(M)}) = (y_{\pi(1)}, \dots, y_{\pi(M)})$$

Attn w/o mask is permutation-equivariant

Deep Sets model (invariant layer):

$$f(x_1, \dots, x_M) = \rho(\sum_{m=1}^M \varphi(x_m)).$$

Equivariant map construction:

$$\rho: \mathbb{R} \times \mathbb{R}^N \rightarrow Y, (x_m, \sum_{k=1}^M \varphi(x_k)) \mapsto y_m$$

5.2 Graph Convolutional Networks

Feature and adjacency matrices

$$X = \text{mat}(x_1^\top; \dots; x_M^\top), A = (a_{nm}) \text{ with } a_{nm} = 1 \Leftrightarrow \{v_n, v_m\} \in E.$$

Permutation matrix constraints

$$P \in \{0, 1\}^{M \times M} \text{ w/ single 1 in each row and col.}$$

Graph invariance definition

$$f(X, A) \stackrel{!}{=} f(PX, PAP^\top), \forall P \in \Pi_M.$$

Graph equivariance definition

$$f(X, A) = Pf(PX, PAP^\top), \forall P \in \Pi_M.$$

Node neighborhood features

$$X_m = \{\{x_n : \{v_n, v_m\} \in E\}\}, \{\cdot\} = \text{multiset}$$

Message passing scheme

$$\varphi(x_m, X_m) = \varphi(x_m, \bigoplus_{x \in X_m} \psi(x)),$$

\oplus is some permutation-invariant operation.

Normalized adjacency matrix

$$\bar{A} = D^{-\frac{1}{2}}(A + I)D^{-\frac{1}{2}}, \quad D = \text{diag}(d_1, \dots, d_M), \\ d_m = 1 + \sum_{n=1}^M a_{nm}, \lambda_{\max}(\bar{A}) = 1, \lambda_{\min}(\bar{A}) \geq -1. \\ \text{For } \bar{L} = I - \bar{A}, \lambda_{\min}(\bar{L}) = 0 \text{ and } \lambda_{\max}(\bar{L}) \leq 2.$$

If \bar{A} symmetric, $\bar{A} = \sum_i \lambda v_i v_i^\top$.

For \bar{A} , $v_1 \propto [-\sqrt{d_i}]^\top$, where d is degree+1.

One GCN layer

$$X^+ = \sigma(\bar{A}XW), W \in \mathbb{R}^{M \times N}.$$

5.2.1 Spectral Graph Theory

Laplacian operator

$$\Delta f = \sum_{n=1}^N \frac{\partial^2 f}{\partial x_n^2}, f: \mathbb{R}^N \rightarrow \mathbb{R}.$$

Graph Laplacian

$$L = D - A, (Lx)_n = \sum_{m=1}^M a_{nm}(x_n - x_m). \\ x^\top Lx = \frac{1}{2} \sum_u \sum_v A_{uv}(x_u - x_v)^2 \geq 0 \text{ (psd).}$$

Normalized Laplacian

$$\tilde{L} = I - D^{-\frac{1}{2}}AD^{-\frac{1}{2}} = D^{-\frac{1}{2}}(D - A)D^{-\frac{1}{2}}.$$

Graph Fourier transform

$$L = D - A = U\Lambda U^\top, \hat{f} = U^\top f, f = U\hat{f}. \\ \Lambda := \text{diag}(\lambda_1, \dots, \lambda_M), \lambda_i \geq \lambda_{i+1}.$$

Convolution: $x * y = U((U^\top x) \odot (U^\top y))$.

Filtering operation $G_\theta(L)x = UG_\theta(\Lambda)U^\top x$

Polynomial kernels

$$U(\sum_{k=0}^K \alpha_k \Lambda^k)U^\top = \sum_{k=0}^K \alpha_k L^k$$

Polynomial kernel network layer

$$x'^{l+1} = \sum_j p_{ij}(L)x'_j + b_i; p_{ij}(L) = \sum_{k=0}^K \alpha_{ijk} L^k$$

GNNs cannot distinguish between certain graphs that are topologically different. Unconstrained set architectures are more powerful. If **WL-test** says graphs are different, they are; but if it says they're the same, they might still be different.

6 Theory of DNNs

6.1 Statistical Learning Theory

Risk decomposition

$$f^*: \text{optimal predictor over all functions}, \\ f_H^* = \arg\min_{f \in H} \mathcal{R}(f), \hat{f}_H: \text{learned from finite data.} \\ \underbrace{\mathcal{R}(\hat{f}_H) - \mathcal{R}(f^*)}_{\text{excess risk}} = \underbrace{\mathcal{R}(\hat{f}_H) - \mathcal{R}(f_H)}_{\text{estimation error}} + \underbrace{\mathcal{R}(f_H) - \mathcal{R}(f^*)}_{\text{approximation error}}$$

Rademacher Complexity $G = \{g_h \mid h \in H\}$:

For $\sigma \in \{-1, 1\}$, measures how well G can fit random noise: $\hat{\mathcal{R}}_{D_n}(G) = \mathbb{E}_\sigma [\sup_{g \in G} \frac{1}{n} \sum_{i=1}^n \sigma_i g(z_i)]$.

$$\mathbb{E}[\sup_{h \in H} \mathcal{R}(h) - \hat{\mathcal{R}}_{D_n}(h)] \leq 2\hat{\mathcal{R}}_{D_n}(G), \\ \mathbb{E}[\mathcal{R}(\hat{h}_H)] \leq \mathcal{R}(h_H^*) + 2\hat{\mathcal{R}}_{D_n}(G).$$

Double descent:

Beyond the interpolation point, models eventually may level out at a lower generalization error.

Implicit bias towards min norm solutions:

Any convergent algorithm with iterates in $\text{span}\{x_1, \dots, x_n\}$ finds minimum norm solution.

6.1.1 A PAC-Bayesian result

P prior distribution over functions before seeing data, Q posterior after training.

PAC-Bayesian theorem

Bounds generalization gap for stochastic classifiers ($f \sim Q$): $E_Q[\mathcal{R}(f)] - E_Q[\hat{\mathcal{R}}_n(f)] \leq \sqrt{\frac{2}{n} [\text{KL}(Q||P) + \ln(2\sqrt{n}/\varepsilon)]}$

- P : prior, Q : posterior (learned). Rate $\tilde{O}(1/\sqrt{n})$
- $\text{KL}(Q||P)$: “information cost” of moving $P \rightarrow Q$
- Insight: generalization depends on **distance moved**, not parameter count

PAC-Bayesian for DNNs

$$P = \mathcal{N}(0, \lambda I), Q = \mathcal{N}(\theta, \text{diag}(\sigma_i^2))$$

$$\text{KL}(Q||P) = \sum_i \left[\log \frac{\lambda}{\sigma_i^2} + \frac{\sigma_i^2 + \theta_i^2 - 1}{2\lambda^2} - \frac{1}{2} \right]$$

Minimize directly: $E_Q[\hat{\mathcal{R}}] + \sqrt{\frac{2}{n} [\text{KL}(Q||P) + \dots]}$

\Rightarrow encourages wide/flat minima (perturbations $\theta + \varepsilon$ must also perform well)

Implementation: Reparameterization: $\tilde{\theta} = \theta + \text{diag}(\sigma_i)\eta, \eta \sim \mathcal{N}(0, I)$, Backprop through θ and σ .

6.2 Linearized DNNs and NTK

Training neural network $f(\theta)(x)$ can be approximated by **linearizing** around initialization θ_0 when parameters change slowly.

Linearization \rightarrow Kernel Regression:

Taylor approximation: $h(\beta)(x) = f(\theta_0)(x) + \beta \cdot \nabla f(\theta_0)(x), \beta = \theta - \theta_0$.

With residuals $\tilde{y}_i = y_i - f(\theta_0)(x_i)$, training becomes **linear regression** with features $\nabla f(\theta_0)(x_i): \min_{\beta} \sum_i \|\tilde{y}_i - \beta^\top \nabla f(\theta_0)(x_i)\|^2$

Neural Tangent Kernel (NTK):

Definition: $k_\theta(x, x') := \nabla f(\theta)(x) \cdot \nabla f(\theta)(x')$.

Dual representation:

$$h(\alpha)(x) = f(\theta_0)(x) + \sum_{i=1}^n \alpha_i k_{\theta_0}(x_i, x).$$

Optimization problem: $\min_{\alpha} \frac{1}{2} \|\mathbf{K}_{\theta_0} \alpha - \tilde{y}\|^2$

Optimal solution (kernel regression): $\alpha^* = \mathbf{K}_{\theta_0}^\dagger(y - f(\theta_0)), h^*(x) = k_{\theta_0}(x)^\top \alpha^*$

Functional Gradient Flow

Training dynamics in function space:

$$\dot{f}(\theta) = \mathbf{K}(\theta)(y - f(\theta))$$

- If $\mathbf{K}(\theta)$ constant \rightarrow linear ODE with closed-form solution
- If $\mathbf{K}(\theta)$ evolves \rightarrow **nonlinear dynamics**, feature learning

Infinite-Width Limit

Init: $w^{(1)} \sim \frac{\sigma_w}{\sqrt{m_\ell}} \mathcal{N}(0, 1)$. Result: As width $m \rightarrow \infty$: $k_{\theta(t)} \rightarrow k_\infty$ (constant during training).

- Kernel becomes **deterministic** (depends only on architecture/init scheme)
- Training = kernel regression with frozen k_∞
- No feature learning

Finite-width: $\|\mathbf{K}(\theta_0) - \mathbf{K}(\theta(t))\| = \mathcal{O}(\frac{1}{m})$

Why Kernel Stays Constant:

Kernel $\nabla K = \nabla^2 f(x) \nabla f(z) + \nabla^2 f(z) \nabla f(x)$, at $m = \infty$: $\nabla^2 f \rightarrow 0 \Rightarrow \nabla K \rightarrow 0$ kernel frozen.

Lazy Training (NTK Regime)	Feature Learning (Rich Regime)
$m \rightarrow \infty$, small LR	Finite width, normal LR
K const \rightarrow linear dynamics	K evolves \rightarrow nonlinear dynamics
No feature learning	Learns representations
Theoretically tractable	SOTA performance

Takeaways:

Linearization turns NN training into kernel regression with features $\nabla f(\theta_0)(x)$.

NTK $k_\theta = \nabla f(\theta)(x) \cdot \nabla f(\theta)(x')$ governs training dynamics via $\dot{f} = \mathbf{K}(y - f)$.

Infinite width \rightarrow kernel constant \rightarrow NN = kernel machine (no feature learning).

Finite width \rightarrow kernel evolves $\mathcal{O}(\frac{1}{m})$ \rightarrow enables feature learning.

NTK explains lazy regime but NOT why deep learning works \rightarrow real power is feature learning when kernel changes.

6.3 Random NNs and GPs

Marginals + Conditionals of MV Gaussians

Let $X \in \mathbb{R}^d \sim \mathcal{N}(\mu, \Sigma)$ with partition:

$$X = \begin{pmatrix} X_A \\ X_B \end{pmatrix}, \quad \mu = \begin{pmatrix} \mu_A \\ \mu_B \end{pmatrix}, \quad \Sigma = \begin{pmatrix} \Sigma_{AA} & \Sigma_{AB} \\ \Sigma_{BA} & \Sigma_{BB} \end{pmatrix}$$

Marginal: $X_A \sim \mathcal{N}(\mu_A, \Sigma_{AA})$

Conditional: $X_B \mid X_A \sim \mathcal{N}(\mu_{B|A}, \Sigma_{B|A})$

$$\mu_{B|A} = \mu_B + \Sigma_{BA} \Sigma_{AA}^{-1} (X_A - \mu_A)$$

$$\Sigma_{B|A} = \Sigma_{BB} - \Sigma_{BA} \Sigma_{AA}^{-1} \Sigma_{AB}$$

6.3.1 Bayesian Linear Regression

Least-squares: $\hat{w} = \arg \min_w \frac{1}{2n\sigma^2} \|y - Xw\|^2$.

Closed-form solution: $\hat{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top y$.

MLE interpretation: $y_i = x_i^\top w + \varepsilon_i, \varepsilon_i \sim \mathcal{N}(0, \sigma^2), \mathcal{L}(w) = \log \prod_{i=1}^n p(y_i | x_i, w), y_i | x_i, w \sim \mathcal{N}(x_i^\top w, \sigma^2)$.

$$\Sigma_{ww} = I_d, \Sigma_{yw} = \mathbf{X}, \Sigma_{yy} = \mathbf{XX}^\top + \sigma^2 I_n.$$

$$\mu_{wy} = \mathbf{X}^\top \Sigma_{yy}^{-1} \mathbf{y} = (\mathbf{X}^\top \mathbf{X} + \sigma^2 I_d)^{-1} \mathbf{X}^\top \mathbf{y},$$

$$\Sigma_{wy} = I_d - \mathbf{X}^\top \Sigma_{yy}^{-1} \mathbf{X} = \sigma^2 (\mathbf{X}^\top \mathbf{X} + \sigma^2 I_d)^{-1}.$$

Same result as ridge:

$$\hat{w} = (\mathbf{X}^\top \mathbf{X} + \sigma^2 I_d)^{-1} \mathbf{X}^\top \mathbf{y}.$$

Equiv. to GP with linear kernel: $f(\cdot) \sim \mathcal{GP}(0, k(\cdot, \cdot))$: $k(x, x') = \phi(x)^\top \phi(x'), y = f + \varepsilon, f \sim \mathcal{N}(0, K), \varepsilon \sim \mathcal{N}(0, \sigma^2 I_n)$, posterior $p(f|y) = \mathcal{N}(\mu_{f|y}, \Sigma_{f|y})$, $\mu_{f|y} = K(K + \sigma^2 I_n)^{-1} \mathbf{y}$.

For $y = f(x) + \varepsilon, f(x) = x^\top w, \varepsilon \sim \mathcal{N}(0, \sigma^2), p(y | X, w) = \mathcal{N}(X^\top w, \sigma^2 I)$. If $w \sim \mathcal{N}(0, \Sigma)$, $p(w | X, X) = \mathcal{N}(\Sigma_{\text{post}}^{-1} + \frac{1}{\sigma^2} X^\top X, \Sigma_{\text{post}})$. Maximizing $\log p(w | X, X)$ is same as minimizing least-squares with ℓ_2 penalty $\frac{1}{2} w^\top \Sigma^{-1} w$, ridge when $\Sigma = \lambda^{-1} I$. **Predictive:** $f_* | x_*, y, X = x_*^\top w \sim \mathcal{N}(x_*^\top \mu_{\text{post}}, x_*^\top \Sigma_{\text{post}} x_*)$. $y_* | x_*, y, X$ adds σ^2 . $f = x^\top w$ with $w \sim \mathcal{N}(0, \Sigma)$ is GP w/ $k(x, x') = \Sigma x^\top x'$.

6.3.2 NNGPs

Setup: Random 1-hidden-layer NN with m units: $f(x) = v_0 + \frac{1}{\sqrt{m}} \sum_{j=1}^m v_j \varphi(\theta_j^\top x)$. Random init: $v_0 \sim \mathcal{N}(0, \sigma_0^2), \mathbb{E}[v_j^2] = \sigma_v^2, \text{Cov}(\theta_j) = \Sigma_\theta$.

Result: As $m \rightarrow \infty$, $f(\cdot) \rightarrow \text{GP}(0, k)$ where $k(x, x') = \sigma_0^2 + \sigma_v^2 \mathbb{E}_\theta[\varphi(\theta^\top x) \varphi(\theta^\top x')]$

Monte Carlo approximation: Sample B random NNs $\{f_b\}_{b=1}^B$, define features:
 $\varphi(x) = \frac{1}{\sqrt{B}} (f_1(x), \dots, f_B(x))^\top$

$\Phi = [f_1 \dots f_B] \in \mathbb{R}^{n \times B}$ (feature matrix)

$\widehat{\mathbf{K}} = \Phi \Phi^\top$ (approximate kernel matrix)

GP regression: Posterior mean and variance: $\mathbb{E}[f(x) | y] = \varphi(x)^\top (\Phi^\top \Phi + \sigma^2 I_B)^{-1} \Phi^\top y$

$$\text{Var}[f(x) | y] = \sigma^2 \varphi(x)^\top (\Phi^\top \Phi + \sigma^2 I_B)^{-1} \varphi(x)$$

Key advantage: Inverts $B \times B$ matrix instead of $n \times n$ when $B \ll n$.

7 Generative Models

7.0.1 Linear Autoencoders

Setup: Encoder $C \in \mathbb{R}^{k \times d}$, decoder $D \in \mathbb{R}^{d \times k}$, data $X \in \mathbb{R}^{d \times n}$ (centered cols): $\min_{C, D} \|X - DCX\|_F^2$.

Optimal Solution (PCA). Let $S = XX^\top$ with eigendecomposition $S = Q\Lambda^2 Q^\top$, $\lambda_1 \geq \dots \geq \lambda_d \geq 0$. Optimal reconstruction via rank- k projection: $\hat{X} = U_k^* U_k^* X$ where $U_k^* = Q_{[:, 1:k]}$ are top- k eigenvectors of S (equiv, top- k left singular vectors of X).

- Any $C = A^{-1}(U_k)^\top, D = U_k^* A$ is optimal
- Reduces to truncated SVD: $\hat{X} = U_k^* \Lambda_k V^\top$ with $\Lambda_k = \text{Diag}(\lambda_1, \dots, \lambda_k, 0, \dots, 0)$
- Convex objective with no spurious local minima (gradient descent finds global optimum)
- Singular vectors may not be uniquely identified

7.0.2 Factor analysis

Latent Variable Models are a generic way to describe generative models. Latent variable $z \sim p(z)$, conditional models for observables $x, p(x|z)$, observed data model: $p(x) = \int p(x|z)p(z) dz$.

Mixture models: simple discrete models: $z \in [K], p(z)$ mixing proportions, $p(x|z)$ condit. densities.

$$p(x; \mu, \Sigma) = \frac{\exp[-\frac{1}{2}(x - \mu)^\top \Sigma^{-1} (x - \mu)]}{\sqrt{(2\pi)^n \det(\Sigma)}}$$

7.0.3 Linear Factor Analysis

- Latent prior: $z \sim \mathcal{N}(0, I)$, $z \in \mathbb{R}^m$
- Observation: $x = \mu + Wz + \eta, \eta \sim \mathcal{N}(0, \Sigma)$
- Independence: $\eta \perp z$
- Typically $m < n$ (fewer factors than features)

Marginal distr: $x \sim \mathcal{N}(\mu, WW^\top + \Sigma)$

- WW^\top : shared variance (low-rank, explained by latent factors)
- Σ : unique variance (diagonal, observation-specific)

Non-identifiability: $(WQ)(WQ)^\top = WQQ^\top W^\top = WW^\top$ for any orthogonal Q . Factors only identifiable up to rotations/reflections. \Rightarrow Use factor rotations (varimax, etc.) for interpretability.

MLE estimation: $\theta = (\mu, W) \leftarrow \log p(x; \mu, W)$

$$\cdot \hat{\mu} = \frac{1}{s} \sum_{i=1}^s x_i \text{ ($$

Probabilistic PCA: Special case $\Sigma = \sigma^2 I$. Optimal i -th column: $w_i = \rho_i u_i$, $\rho_i^2 = \max\{0, \lambda_i - \sigma^2\}$. $W = U_m L_m$, where (λ_i, u_i) is i -th eigenpair of data covariance.

As $\sigma \rightarrow 0$: $\mu_{z|x} \rightarrow W^\dagger(x - \mu)$ (standard PCA). If W has orthogonal columns, then $W^\dagger = W^\top$.

7.1 Variational Autoencoders

$z \in \mathbb{R}^d$ is learned embedding of x . For generation, $z \sim \mathcal{N}(0, I)$, decoder $p_\theta(z|x)$ maps latent to data.

Problem: $p_\theta(x) = \int p(z)p_\theta(x|z) dz$ intractable.
Solution: Maximize ELBO instead: $\log p_\theta(x) \geq \underbrace{\mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)]}_{\text{reconstruction}} - \underbrace{D_{\text{KL}}(q_\phi(z|x) \parallel p(z))}_{\text{regularization}}$

- **Reconstruction:** Encode $x \rightarrow z$, decode back
- **KL term:** Keep encoder output close to prior $p(z) \sim \mathcal{N}(0, I)$, ensures generation using latents

Encoder $q_\phi(z|x) = \mathcal{N}(\mu_\phi(x), \text{diag}(\sigma_\phi^2(x)))$.

KL closed form: $D_{\text{KL}}(q_\phi(z|x) \parallel p(z)) = \frac{1}{2} \sum_{j=1}^d (\sigma_{\phi,j}^2 + \mu_{\phi,j}^2 - 1 - \log \sigma_{\phi,j}^2)$.

$\text{KL}(\mathcal{N}(\mu_0, \sigma_0^2) \parallel \mathcal{N}(\mu_1, \sigma_1^2)) = \frac{1}{2} \left(\frac{\sigma_0^2}{\sigma_1^2} + \frac{(\mu_0 - \mu_1)^2}{\sigma_1^2} - 1 + \log \frac{\sigma_1^2}{\sigma_0^2} \right)$

$\text{KL}(p \parallel q) = \mathbb{E}_p [\log \frac{p(x)}{q(x)}]$

- **Fwd KL:** $q_1^* = \arg \min_{q \in Q} \text{KL}(p \parallel q)$
- **Rev KL:** $q_2^* = \arg \min_{q \in Q} \text{KL}(q \parallel p)$

Rev KL: Mode-seeking ($p = 0 \Rightarrow q = 0$), FwdKL: Mean-seeking ($p \neq 0 \Rightarrow q \neq 0$). MLE minimizes fwd KL to empirical \mathcal{D} .

$$\log p_\theta(x|z) = -\frac{1}{2\sigma^2} \|x - \mu_\theta(z)\|_2^2 - \frac{d}{2} \log(2\pi\sigma^2).$$

Reparameterization trick: $z = \mu_\varphi(x) + \sigma_\varphi(x) \odot \varepsilon$, $\varepsilon \sim \mathcal{N}(0, I)$, enables backprop through sampling.

$\log p_\theta(x) - \text{ELBO} = D_{\text{KL}}(q_\phi(z|x) \parallel p_\theta(z|x))$, tight when q_ϕ = true posterior.

MC estimation: $E_{q_\phi(z|x)}[\log p_\theta(x|z)] \approx -\frac{1}{2\sigma^2 K} \sum_{k=1}^K \|x - \mu_\theta(z_k)\|_2^2 - \frac{d}{2} \log(2\pi\sigma^2)$.

Generative Classifiers Given $y \in \{0, 1\}$, $p(y=1) = p(y=0) = \frac{1}{2}$, $p(x|y) = \mathcal{N}(x; \mu_y, I_d)$, where $\mu_0, \mu_1 \in \mathbb{R}^d$, $p(y=1|x) = \frac{p(y=1)p(x|y=1) + p(y=0)p(x|y=0)}{p(y=1)p(x|y=1) + p(y=0)p(x|y=0)} = \frac{\frac{1}{2}(2\pi)^{-\frac{d}{2}} \exp(-\frac{1}{2}\|x - \mu_1\|^2)}{\frac{1}{2}(2\pi)^{-\frac{d}{2}} \exp(-\frac{1}{2}\|x - \mu_1\|^2) + \frac{1}{2}(2\pi)^{-\frac{d}{2}} \exp(-\frac{1}{2}\|x - \mu_0\|^2)} = \frac{1}{1 + \exp(\frac{1}{2}\|x - \mu_1\|^2 - \frac{1}{2}\|x - \mu_0\|^2)} = \frac{1}{1 + \exp(-[(\mu_1 - \mu_0)^\top x + \frac{1}{2}(\|\mu_0\|^2 - \|\mu_1\|^2)] / (1 + \exp(-[(\mu_1 - \mu_0)^\top x + \frac{1}{2}(\|\mu_0\|^2 - \|\mu_1\|^2)]))}$ equiv to logistic regression where $p(y=1|x) = \sigma(w^\top x + b)$ with $w = \mu_1 - \mu_0$, $b = \frac{1}{2}(\|\mu_0\|^2 - \|\mu_1\|^2)$.

ELBO for Hierarchical VAEs: Model x by decoding from latents $z = (z_1, \dots, z_L)$. $p_\theta(x, z) = p_\theta(z|z_1) \prod_{i=1}^{L-1} p_\theta(z_i|z_{i+1}) p(z_L)$. **Inference** top-down: $q_\phi(z|x) = q_\phi(z_L|x) \prod_{i=1}^{L-1} q_\phi(z_i|z_{i+1})$. **ELBO for HVAE:** $\mathcal{L}(x) = \mathbb{E}_{z|x \sim q_\phi} [\log \frac{p_\theta(x, z)}{q_\phi(z|x)}] = \mathbb{E}_{z|x \sim q_\phi} [\log p_\theta(z|z_1) + \log \frac{p_\theta(z_1|z_2)}{q_\phi(z_1|x)} + \dots + \sum_{i=2}^{L-1} \log \frac{p_\theta(z_i|z_{i+1})}{q_\phi(z_i|z_{i-1})} + \log \frac{p_\theta(z_L)}{q_\phi(z_L|z_{L-1})}]$.

Change of variables spherical to 3D euclidian $(x, y, z) \mapsto (r \cos \theta \cos \phi, r \cos \theta \sin \phi, r \sin \theta)$. Lengths of the three sides of an infinitesimal cuboid whose diagonally opposite vertices are at r, θ, ϕ and $(r+dr, \theta+d\theta, \phi+d\phi)$ are $(dr, rd\theta, r \cos \theta d\phi)$. Volume is $r^2 \cos \theta dr d\theta d\phi$. Determinant of jacobian $|\frac{\partial(x,y,z)}{\partial(r,\theta,\phi)}| = r^2 \cos \theta$. Density on spherical coordinates $p(r, \theta, \phi) \rightarrow$ density on Euclidian coordinates is $p(x, y, z) = p(r, \theta, \phi) |\frac{\partial(x,y,z)}{\partial(r,\theta,\phi)}|^{-1}$. Infinitesimal probability mass of the cuboid above is equal to the mass of a Euclidian cuboid of size (dx, dy, dz) at (x, y, z) .

7.2 Normalizing Flows

Transform simple distribution $z \sim \mathcal{N}(0, I)$ through invertible map T to get complex $x = T(z)$. Exact likelihood (no ELBO like VAEs), easy sampling.

Change of Variables Formula:

$$p_x(x) = p_z(T^{-1}(x)) \cdot \det J_{T^{-1}}(x), \quad |\det J_{T^{-1}}(x)| = \frac{1}{|\det J_T(T^{-1}(x))|}.$$

Diffeomorphism: T is bijective, differentiable, w/ differentiable inverse. Guarantees $\det J_T \neq 0$.

Computational problem: Computing $\det J$ is $O(d^3)$ for dense Jacobian. **Solution:** Design T s.t. Jacobian is **triangular**, then only $O(d)$!

Two architectures with triangular Jacobians:

	MAF	IAF
Fast / parallel	Density eval	Sampling
Slow / sequential	Sampling	Density

Coupling layers: Trick that makes both directions fast, at the cost of being less expressive per layer.

7.3 Generative Adversarial Networks

Likelihood-free generative model: train via adversarial game between two networks: **Generator** G_θ maps latent $z \sim p_z$ (typically Gaussian) to fake samples; **Discriminator** D_φ : outputs prob that input is **real**. **GAN Objective:**

$$\min_\theta \max_\varphi \underbrace{\mathbb{E}_{x \sim p_r} [\log D_\varphi(x)]}_{\text{real samples}} + \underbrace{\mathbb{E}_{z \sim p_z} [\log (1 - D_\varphi(G_\theta(z)))]}_{\text{fake samples}}$$

- **Discriminator** maximizes: correctly classify real (high D) and fake (low D)
- **Generator** minimizes: fool discriminator (make $D(G(z))$ high)

Common alternative objective for the generator is to maximize $\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$ instead of minimizing $\mathbb{E}_{z \sim p(z)} [\log (1 - D(G(z)))]$ to help mitigate vanishing gradient problem when discriminator becomes to good, i.e. $D(G(z)) \rightarrow 0$. For a fixed generator G , **optimal discriminator** D^* is given by $D^*(x) = \frac{p_{\text{data}}(x)}{p_{\text{data}}(x) + p_G(x)}$. If **discriminator optimal**, GAN objective reduces to $V(D^*, G) = 2D_{\text{JS}}(p_{\text{data}} \parallel p_G) - \log 4$.

7.3.1 Theoretical Foundation

Binary classification with $p(y=1) = p(y=0) = \frac{1}{2}$:

- $y=1$: sample from real $p_{r(x)}$
- $y=0$: sample from generator $p_{\theta(x)}$

Bayes Optimal Classifier (prob that x is real): $q_\theta(x) = P(y=1|x) = \frac{p_{r(x)}}{p_{r(x)} + p_{\theta(x)}}$.

Generator Logistic Objective = JS Div:

$$\ell^*(\theta) = \mathbb{E}_{p_{\theta}(x,y)} [y \ln q_\theta(x) + (1-y) \ln(1 - q_\theta(x))] = \text{JS}(p_r \parallel p_\theta) - \ln 2.$$

Jensen-Shannon Divergence: $\text{JS}(p_r \parallel p_\theta) = \frac{1}{2} D_{\text{KL}}(p_r \parallel p_m) + \frac{1}{2} D_{\text{KL}}(p_\theta \parallel p_m)$, $p_m = \frac{p_r + p_\theta}{2}$.

Bounded: $0 \leq \text{JS}(p_r \parallel p_\theta) \leq \log 2$.

Jensen Inequality

If φ convex: $\varphi(\mathbb{E}[X]) \leq \mathbb{E}[\varphi(X)]$. If concave, other way around.

7.3.2 Training

Alternating SGD (heuristic, may diverge!). Training is **Saddle-point problem**, unstable!

JS Divergence Saturates when distributions don't overlap. If p_r and p_θ have disjoint supports: discriminator perfect, no gradient for generator!

Wasserstein Distance (Earth Mover's Distance):

$$W(p_r, p_\theta) = \inf_{\gamma \in \Pi(p_r, p_\theta)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$

Minimum total "work" to transport mass from p_r to p_θ . Provides meaningful gradients even without overlap.

Dual (Kantorovich-Rubinstein):

$$W(p_r, p_\theta) = \sup_{f \in \mathcal{F}} \mathbb{E}_{x \sim p_r} [f(x)] - \mathbb{E}_{x \sim p_\theta} [f(x)]$$

Maximize gap between avg score of real vs fake samples w.r.t. Lipschitz constraint.

Max achievable gap = Wasserstein distance.

WGan uses critic f_w (not classical discriminator!):

$$\min_\theta \max_w \mathbb{E}_{x \sim p_r} [f_w(x)] - \mathbb{E}_{z \sim p_z} [f_w(G_\theta(z))]$$

Enforcing Lipschitz:

- **Weight clipping** (original): crude, problematic
- **Gradient penalty:**
add $\lambda \mathbb{E}_{\hat{x}} [(\|\nabla_{\hat{x}} f_w(\hat{x})\|_2 - 1)^2]$

Mode Collapse: Generator produces only few samples that fool discriminator, ignoring full distribution diversity.

7.4 Diffusion Models

Fixed Fwd process: Gradually add Gaussian noise over T steps until data becomes pure noise.

Forward step: $q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t} x_{t-1}, \beta_t I)$. Full fwd process: $q(x_{1:T} | x_0) = \prod_{t=1}^T q(x_t | x_{t-1})$.

Noise schedule: $\{\beta_t \in (0, 1)\}_{t=1}^T$ noise added at each step. Define: $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$.

Direct sampling (reparameterization trick):

$$q(x_t | x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t) I). \\ x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, I).$$

Reverse process (learned): Train NN to denoise step by step: $p_\theta(x_{t-1} | x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \sigma_\theta^2(t) I)$.

For small β_t , reverse $q(x_{t-1} | x_t)$ is also Gaussian.

$\log p_\theta(x_0)$ intractable, so derive **Variational Lower Bound (VLB)**: $-\log p_\theta(x_0) \leq \mathcal{L}_{\text{VLB}} = \mathbb{E}_q [\log \frac{q(x_{1:T} | x_0)}{p_\theta(x_{1:T} | x_0)}]$.

Decomposition into 3 terms: $\mathcal{L}_{\text{VLB}} = \underbrace{D_{\text{KL}}(q(x_T | x_0) \parallel p(x_T))}_{L_{T-1}} + \underbrace{\sum_{t=2}^T \mathbb{E}_{q(x_t | x_0)} [D_{\text{KL}}(q(x_{t-1} | x_t, x_0) \parallel p_\theta(x_{t-1} | x_t, x_0))]}_{L_{t-1}} - \underbrace{\mathbb{E}_{q(x_1 | x_0)} [\log p_\theta(x_0 | x_1)]}_{L_0}$

- L_T : Is $q(x_T | x_0) \approx \mathcal{N}(0, I)$? Not optimized.
- L_{t-1} : Match learned reverse to true reverse
- L_0 : Reconstruction term

Tractable Reverse Posterior $q(x_{t-1} | x_t, x_0)$ is Gaussian with closed form (product of Gaussians): $q(x_{t-1} | x_t, x_0) = \mathcal{N}(x_{t-1}; \mu_{q,t}(x_t, x_0), \sigma_{q,t}^2 I)$, with:

$$\mu_{q,t}(x_t, x_0) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \alpha_t}} \varepsilon_0 \right), \quad \sigma_{q,t}^2 = \frac{1}{1 - \alpha_{t-1}} \beta_t.$$

$\lambda_i \gg \mu: \frac{\lambda_i}{\lambda_i + \mu} \approx 1 \rightarrow$ weak shrinkage (important dirs).

$\lambda_i \ll \mu: \frac{\lambda_i}{\lambda_i + \mu} \approx 0 \rightarrow$ strong shrinkage (flat dirs).

Adaptively shrinks based on loss geometry, preserves important dirs, removes unnecessary complexity.

Early stopping: Rather than training to convergence, stop when validation performance plateaus. Analysis shows that this is approximately equivalent to L2 regularization. GD trajectories can be approximated as $\theta(k) = [I - (I - \eta \Lambda)^k] \theta^*$. For small step sizes, behaves like weight decay when $k = \frac{1}{\eta \mu}$.

L^1 regularized second-order approximation of an arbitrary loss function around optimal θ^* is $R_{L^1}(\theta) \approx R(\theta^*) + \frac{1}{2}(\theta - \theta^*)^\top H(\theta - \theta^*) + \lambda \|\theta\|_1$. Assuming $H = \text{diag}(h_1, \dots, h_d)$, we get $R_{L^1}(\theta) \approx \sum_{i=1}^d \left[\frac{h_i}{2} (\theta_i - \theta_i^*)^2 + \lambda |\theta_i| \right] + \text{const}$ so we need to minimize $f(a) = \frac{1}{2}(a - b)^2 +$

Cosine noise schedule performs better than linear.

Used architecture is U-Net. **Input:** Noisy image $x_t + \text{timestep } t$; **Output:** Predicted noise $\varepsilon_\theta(x_t, t)$.

Model **conditional distribution** $p_\theta(x_{0:T} | y)$ where y is condition (class, text, image). Extend denoiser to take y as input.

Latent Diffusion Models (LDM) run diffusion in **compressed latent space** instead of pixel space.

Autoregressive Models
 $p(x) = \prod_{i=1}^d p(x_i | x_{<i})$.

8 Tricks

Short Connections in DNs: Add less deep paths to a very deep network. **Residual connections**: shortcut and add back in. **Skip connections**: concatenate.

For $y_1 = \alpha f(x, \Theta_1) + x$, $x \in \mathbb{R}^d, y_1 \in \mathbb{R}^d, \Theta_i \in \mathbb{R}^{d \times d}$ and $y_{i+2} = \alpha f(y_{i-1}, \Theta_i) + y_i$, $i = L(y_1)$ it holds that $\frac{\partial y_k}{\partial y_{k-1}} = \alpha \frac{\partial f(y_{k-1}, \Theta_{k-1})}{\partial y_{k-1}} + I_d$ and $\frac{\partial y_k}{\partial \Theta_k} = \alpha \frac{\partial f(y_{k-1}, \Theta_{k-1})}{\partial \Theta_k}$. By applying the chain rule we have $\frac{\partial \ell}{\partial \Theta_k} = \frac{\partial L}{\partial y_n} \frac{\partial y_n}{\partial y_{n-1}} \dots \frac{\partial y_{k+1}}{\partial y_k} \frac{\partial y_k}{\partial \Theta_k}$. Set $\alpha^2 = \frac{a}{n}$ for $a > 0$ s.t. $\lim_{n \rightarrow \infty} \mathbb{E}[y]^2 < \infty$.

8.1 Weight Decay & Early Stopping

L2 regularization $\mathcal{R}_\Omega(\theta; \mathcal{S}) = \mathcal{R}(\theta; \mathcal{S}) + \Omega(\theta)$, $\Omega_\mu(\theta) = \frac{\mu}{2} \|\theta\|^2$, $\mu \geq 0$.

Only penalize weights, not biases. **GD upd w/ WD**: $\Delta \theta = -\eta \nabla \mathcal{R}(\theta) - \eta \nabla \Omega_\mu(\theta) = -\eta \nabla \mathcal{R}(\theta) - \eta \mu \theta$.

Geometric interpretation (local quadratic approx): Regularized optimum: $\theta_\mu^* = (H + \mu I)^{-1} H \theta^*$, where $H = Q^\top \Lambda Q$ gives $\theta_\mu^* = Q \text{ diag}(\frac{\lambda_i}{\lambda_i + \mu}) Q^\top \theta^*$.

$\lambda_i \gg \mu: \frac{\lambda_i}{\lambda_i + \mu} \approx 1 \rightarrow$ weak shrinkage (important dirs).

$\lambda_i \ll \mu: \frac{\lambda_i}{\lambda_i + \mu} \approx 0 \rightarrow$ strong shrinkage (flat dirs).

Adaptively shrinks based on loss geometry, preserves important dirs, removes unnecessary complexity.

Early stopping: Rather than training to convergence, stop when validation performance plateaus. Analysis shows that this is approximately equivalent to L2 regularization. GD trajectories can be approximated as $\theta(k) = [I - (I - \eta \Lambda)^k] \theta^*$. For small step sizes, behaves like weight decay when $k = \frac{1}{\eta \mu}$.

L^1 regularized second-order approximation of an arbitrary loss function around optimal θ^* is $R_{L^1}(\theta) \approx R(\theta^*) + \frac{1}{2}(\theta - \theta^*)^\top H(\theta - \theta^*) + \lambda \|\theta\|_1$. Assuming $H = \text{diag}(h_1, \dots, h_d)$, we get $R_{L^1}(\theta) \approx \sum_{i=1}^d \left[\frac{h_i}{2} (\theta_i - \theta_i^*)^2 + \lambda |\theta_i| \right] + \text{const}$ so we need to minimize $f(a) = \frac{1}{2}(a - b)^2 +$

$|a|$ and this gives $a^* = \text{sgn}(b) \max\{0, |b| - \beta\}$, so $\theta_i = \text{sgn}(\theta_i^*) \max\{0, |\theta_i^*| - \frac{\lambda}{h_i}\}$. For L^2 regularization we get $\theta_i = \frac{h_i}{h_i + \lambda} \theta_i^*$. **Connecting L^2 w early stopping:** $R(w) \approx R(w^*) + \frac{1}{2}(w - w^*)^\top H(w - w^*)$. $\nabla R(w) = H(w - w^*)$. GD update: $w^t = w^{t-1} - \eta H(w^{t-1} - w^*)$ gives $w^t - w^* = (I_d - \eta H)(w^{t-1} - w^*)$. Using $H = Q\Lambda Q^\top$: $Q^\top (w^t - w^*) = (I_d - \eta \Lambda) Q^\top (w^{t-1} - w^*)$. If $w^0 = 0$, $Q^\top (w^t - w^*) = (I_d - \eta \Lambda)^t Q^\top (0 - w^*) \Rightarrow Q^\top w^t = [I_d - (I_d - \eta \Lambda)^t] Q^\top w^*$. Optimal w under L^2 reg gives $Q^\top w = [I_d - \lambda(I + \lambda I_d)^{-1}] Q^\top w^*$. Matching both gives $t \approx \frac{1}{\eta \lambda}$. **Weight normalization** is like Batch-Norm, with the covariance matrix replaced by the identity matrix.

8.2 Ensemble Methods

Bagging: Create K bootstrap samples of Data (sampling with replacement), train separate models, and average predictions: $p(y|x) = \frac{1}{K} \sum_{k=1}^K p(y|x; \theta_k)$.

Dropout: Randomly drop units during training with probability $1 - \pi$. Creates an exponential ensemble of sub-networks sharing weights. Test time: Scale weights by π to approximate the ensemble average.

8.3 Normalization

Batch Norm: Normalize activations across mini-batch: $\tilde{z} = \frac{z - \mu_{\text{batch}}}{\sigma_{\text{batch}}}, \hat{z} = \alpha \tilde{z} + \beta, \mu_{\text{batch}} = \frac{1}{b} \sum_{i=1}^b z_i, \sigma_{\text{batch}} = \sqrt{\frac{1}{b} \sum_{i=1}^b (z_i - \mu_{\text{batch}})^2}$.

Layer Norm: Normalize features in a layer instead; particularly effective for RNNs (batch statistics are less stable).

8.4 Data / Task Augmentation

Augment Data by applying valid transformations. Semi-supervised Learning: Train jointly on labeled and unlabeled data w combined loss. Pre-training & Fine-tuning. Multi-task Learning. Self-supervised Learning: Create free supervision from data.

9 Recurrent Neural Networks

Evolution: $z_t = F[\theta](z_{t-1}, x_t)$, with $z_0 = 0$.

Optional output: $y_t = G[\theta](z_t)$.

Simple RNN: $z_t = \phi(Wz_{t-1} + Ux_t)$ where $W \in \mathbb{R}^{m \times m}, U \in \mathbb{R}^{m \times n}$.

Backpropagation Through Time (param sharing): $\frac{\partial R}{\partial w_{ij}} = \sum_t \frac{\partial R}{\partial z_i} \cdot \phi'_t \cdot z_j^{t-1}$.

$\frac{\partial R}{\partial u_{ik}} = \sum_t \frac{\partial R}{\partial z_i} \cdot \phi'_t \cdot x_k^t; \phi'_t = \phi'(F_i(z^{t-1}, x^t))$.

Gradient flow backward through time:

$$\nabla_{x_t} \mathcal{R} = [\prod_{r=t+1}^s W^\top S(z^r)] \cdot J_G \cdot \nabla_y \mathcal{R}$$

Spectral analysis: $\|\prod W^\top S(z^r)\|_2 \leq [\sigma_{\max}(W)]^{s-t}$ **Root cause:** Repeated matmul through time. If $\sigma_{\max}(W) < 1$, vanishing gradients; if $\sigma_{\max}(W) > 1$, exploding gradients.

⇒ Simple RNNs cannot learn long dependencies.

Deep RNNs stack layers vertically:

$$z^{t,\ell} = \varphi(W_\ell z^{t-1,\ell} + U_\ell z^{t,\ell-1}) \text{ where } z^{t,0} = x_t.$$

For RNN with $z_{t+1} = \varphi(Uz_t + Vx_{t+1})$, $L = \sum_{t=1}^T \ell(\hat{y}_t, y_t)$, where \hat{y}_t depends on z_t . Then $\frac{\partial L}{\partial U} = \sum_{t=1}^T \frac{\partial L}{\partial z_t} \cdot (\varphi'_t \cdot z_t), \frac{\partial L}{\partial V} = \sum_{t=1}^T \frac{\partial L}{\partial z_t} \cdot (\varphi'_t \cdot x_{t+1})$.

Weight Sharing in RNNs (LSTM):

$$\frac{\partial L}{\partial W} = \sum_{t=1}^T \frac{\partial L}{\partial W_t}.$$

Proof idea: Introduce dummy parameters $\tilde{W}_i = f(W)$ for each time step. By chain rule: $\frac{\partial L}{\partial W} = \sum_i \frac{\partial L}{\partial \tilde{W}_i} \frac{\partial \tilde{W}_i}{\partial W_i}$. With constraint $\tilde{W}_i = W$, we have $\frac{\partial \tilde{W}_i}{\partial W_i} = I$, giving the sum.

Initialization of bias in RNNs: Use 1.

Backpropagation for example network

- $a_t = F(x_t, h^{(t-1)}, y^{(t-1)}; \theta)$
- $h_t = \sigma(a_t)$
- $\hat{y}_t = G(h_t; \varphi)$
- $L_t = H(\hat{y}_t, y_t)$
- $L = \sum_{t=1}^T L_t$

$$\frac{\partial L}{\partial \theta} = \sum_{t=1}^T \frac{\partial L_t}{\partial \hat{y}_t} \frac{\partial \hat{y}_t}{\partial h_t} \frac{\partial h_t}{\partial a_t} \sum_{i=1}^t \left(\prod_{j=t}^{i+1} \frac{\partial h_j}{\partial h_{j-1}} \right) \frac{\partial a_i}{\partial \theta},$$

Other example for BPTT: $h_t = f_H(x_t, h_{t-1}; \theta)$ If $L = \sum_{t=1}^T L_t$, where $L_t = L(h_t)$, one gets:

$$\frac{\partial L_t}{\partial \theta} = \frac{\partial L_t}{\partial h_t} \sum_{k=1}^t \left(\frac{\partial h_t}{\partial \theta} \cdot \frac{\partial h_k}{\partial \theta} \right), \text{ where } \frac{\partial h_t}{\partial \theta} = \prod_{i=k+1}^t \frac{\partial h_i}{\partial \theta}.$$

9.1 Long Short-Term Memory (LSTM)

- C_t : cell state (internal memory, protected highway)
- z_t : hidden state (external output, filtered view)

$$C_t = \underbrace{\sigma(F\tilde{x}^t) \odot C_{t-1}}_{\text{forget}} + \underbrace{\sigma(G\tilde{x}^t) \odot \tanh(V\tilde{x}^t)}_{\text{input}}, \\ z_t = \underbrace{\sigma(H\tilde{x}^t) \odot \tanh(C_t)}_{\text{output}}, \text{ where } \tilde{x}^t = [x_t, z_{t-1}].$$

Keeping or forgetting stored content?

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

Prep new input info to be added to the memory.

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Combine stored and new information

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t \text{ (long term memory)}$$

Compute output selectively

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b)$$

$h_t = o_t \odot \tanh(C_t)$ (short term memory + outpt)

9.2 Gated Recurrent Unit (GRU)

Simplified LSTMs (3 weight mats instead of 5).

Single state z_t . Input: $\tilde{x}^t = [x_t, z_{t-1}]$.

$$u_t = \sigma(U\tilde{x}^t), r_t = \sigma(R\tilde{x}^t),$$

$$\tilde{z}_t = \tanh(W[r_t \odot z_{t-1}, x_t])$$

$$z_t = (1 - u_t) \odot z_{t-1} + u_t \odot \tilde{z}_t$$

Often comparable to LSTM with fewer resources. Gating creates identity paths → better gradient flow.

9.3 Linear Recurrent Models

RNNs not parallelizable during training. LRU has linear dynamics: $z_{t+1} = Az_t + Bx_t$. Diagonalize to $A = P\Lambda P^{-1}, \lambda_i \in \mathbb{C}$, change basis $\zeta_t = P^{-1}z_t$. Then: $\zeta_{t+1} = \Lambda\zeta_t + Cx_t$. Each dimension evolves independently (no channel mixing). Compensate with expressive output: $y_t = \text{MLP}(\text{Re}(Gz_t))$.

Stability: Require $\max|\lambda_j| \leq 1$ (spectral radius ≤ 1).

Parameterization: $\lambda_i = \exp(-\exp(\nu_i) + i\varphi_i)$ ensures $|\lambda_i| \in (0, 1)$ automatically, $|\lambda_i| \approx 1$: Long-term memory, $|\lambda_i| \approx 0$: Short-term patterns.

Provably universal as sequence-to-sequence map.

9.3.1 Connectionist Temporal Classification

Problem: Unsegmented sequences (e.g., speech).

Solution: RNN outputs prob distribution over vocabulary at each time step. Model all alignments with blank symbol “_”: $p(\ell | x) = \sum_{\pi \in \mathcal{B}^{-1}(\ell)} \prod_t y_{\pi_t}$. \mathcal{B} removes blanks and repeated symbols.

9.4 Sequence Learning

Teacher Forcing: $p(y^t)$ depends on $y^{1:t-1}$ only through z^t , means during autoregressive generation, model doesn't see its own predictions.

Solution: Add feedback connections from y^{t-1} to z^t : $z^t = \text{RNN}(z^{t-1}, x^t, y^{t-1})$, now model conditions on its own previous predictions → more coherent gen.

Professor Forcing: Train two networks (teacher-forced + free-running), discriminator matches hidden states → improved generalization.

Exposure bias: Model relies on itself where inputs come from the previous output because of the non-availability of the ground truth.

Seq2Seq: Input and output sequences have different lengths: Use encoder-decoder framework.

Gradients in bi-directional RNNs are computed by making a forward and backward run, then at timestep t we combine (concatenate/add) and continue with the backpropagation. This happens at every bidirectional layer.

To initialize points uniformly on the unit ring in \mathbb{C} between circles of radii $r_2 > r_1$, sample from $\exp(\frac{1}{2} \log(u_1(r_2^2 - r_1^2) + r_1^2) + i2\pi u_2)$ w $u_1, u_2 \sim \mathcal{U}(0, 1)$.

10 Attention and Transformers

Seq2Seq with Attention: Encoder generates hidden state sequence. Decoding RNNs output attends to encoder states and gets used as input in next step.

Attention: Learn to index, multiplicative gating to combine bottom-up and top-down information.

KV - attention map:

$$F(\xi, ((x_1, z_1), \dots, (x_s, z_s)) = [z_1, \dots, z_s]. \\ f(\xi, (x_1, \dots, x_s)) \text{ consisting of a query } \xi \text{ (what to look for?), keys } x_i \text{ (index) and values } z_i.$$

Scaled Dot-Product Attention: $f(\xi, x) = \frac{\xi \cdot x}{\sqrt{n}}$

Multi-headed attention:

$$G(\xi, (x^t, z^t)_{t=1}^s) = W \begin{bmatrix} F_1(\xi, (x^t, z^t)) \\ \vdots \\ F_h(\xi, (x^t, z^t)) \end{bmatrix},$$

where

$$F_j(\xi, (x^t, z^t)) = F\left(W_j^q \xi, (W_j^k x^t, W_j^v z^t)\right).$$

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V,$$

where

$$Q = XW_Q, \quad K = XW_K, \quad V = XW_V \quad \text{and} \quad X \in \mathbb{R}^{T \times d_{\text{model}}}, \quad W_Q, W_K, W_V \in \mathbb{R}^{d_{\text{model}} \times d_k}, \quad W_V \in \mathbb{R}^{d_{\text{model}} \times d_v}, \quad Q, K \in \mathbb{R}^{T \times d_k}, \quad V \in \mathbb{R}^{T \times d_v}.$$

Linear Binary ($y \in \{-1, 1\}, f(x) = w^\top x + b$):
Correct: $y(w^\top x + b) > 0$.
Adv. flips when $yw^\top x \leq -y(w^\top x + b)$ cross hyperplane. **L2 optimal:** $\delta^* = \frac{-w^\top x + b}{\|w\|_2^2}$. $\delta^* = \frac{|w^\top x|}{\|w\|_2^2} w = \frac{|w^\top x|}{\|w\|_2} w$.

Multiclass: $f_k(x) = w_k x + b_k$, use $\text{argmax}_k f_k(x)$. Margin to class j : $m_j(x) = (w_j - w_b)^\top x + (b_j - b_b)$.

$f_j(x) = f_j(x) \gg m_j(x) = 0$. Correct if $f_j(x) > f_j(y) \forall j \neq y$, adversarial if $\exists j \neq y$ s.t. $f_j(x + \delta) < f_j(x + \delta)$. Distance to boundary: $\frac{m_j - m_y}{\|w_j - w_y\|_2}$.

Adversarial attacks for NNs: Approximate boundary by $f(x + \delta) \approx f(x) + \nabla f(x)^\top \delta$. FGSM is a one-step ℓ_1 attack: $\delta = \epsilon \text{ sign}(\nabla_x L(f(x), y))$. PGD is multi-step $\delta_{t+1} = \text{Proj}_{\delta \in \mathcal{S}}(\delta_t + \alpha \text{ sign}(g_t))$.

Distributionally Robust Optimization: $\min_u \sup_{Q \in \mathcal{P}(P)} E_Q L(f(x), y)$, where U means close. Can use upper bound on Wasserstein distance e.g.

Robust training: $\min_u E[\max_{Q \in \mathcal{P}(P)} L(f(x), y)]$. Adversarial training can be viewed as robustness to distribution shift measured by Wasserstein distance.

Interpretability: Local - explain pred for specific x , Global - explain model behaviour on avg over data.

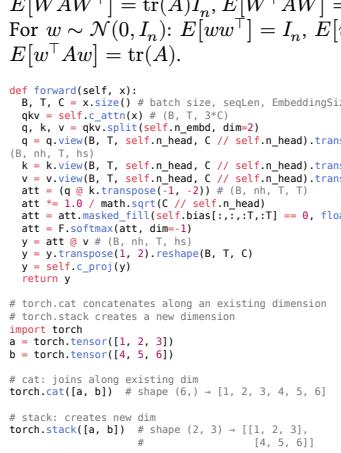
Local: Ceteris paribus (vary x_j , fix x_{-j}). Sensitivity ($\partial_{x_j} f(x)$), missing info ($f(x) - \mathbb{E}[f(X)] | X_{-j} = x_{-j}$). **Global:** Mutual info ($I(X_j; Y | X_{-j})$), Predictive util (train f w/ and w/o x_j). For log-loss predictive util \approx conditional mutual information.

SHAP attributes predictions, while SAGE attributes risk reduction.

A predicted attribute, Y target outcome, \hat{Y} prediction. Demographic Parity: $\hat{Y} \perp A$; Equalized Odds: $\hat{Y} \perp A | Y$; Equality of Opportunity: $\hat{Y} \perp A | Y = 1$.

For $W \in \mathbb{R}^{n \times d}$, $W_{i,j} \sim \mathcal{N}(0, 1)$: $E[W] = 0, E[WW^\top] = dI_n, E[W^\top W] = nI_d; E[WAW^\top] = \text{tr}(A)I_n, E[W^\top AW] = \text{tr}(A)I_d$.

For $w \sim \mathcal{N}(0, I_n)$: $E[wu^\top] = I_n, E[w^\top w] = n, E[w^\top Aw] = \text{tr}(A)$.



11 Ethics

Adversarial examples (given $f(x) = y$ correctly):

- Untargeted: $\|\delta\| \leq \epsilon$ s.t. $f(x + \delta) \neq y$.
- Optimized max $\| \delta \|_{\text{Lip}(f)} \leq L$ s.t. $f(x + \delta) = y$.
- Targeted: $\|\delta\| \leq \epsilon$ s.t. $f(x + \delta) = t \neq y$.
- Optimized min $\| \delta \|_{\text{Lip}(f)} \leq L$ s.t. $f(x + \delta) = t$.

Linear Binary ($y \in \{-1, 1\}, f(x) = w^\top x + b$):

Correct: $y(w^\top x + b) > 0$.
Adv. flips when $yw^\top x \leq -y(w^\top x + b)$ cross hyperplane. **L2 optimal:** $\delta^* = \frac{-w^\top x + b}{\|w\|_2^2}$. $\delta^* = \frac{|w^\top x|}{\|w\|_2^2} w = \frac{|w^\top x|}{\|w\|_2} w$.

Multiclass: $f_k(x) = w_k x + b_k$, use $\text{argmax}_k f_k(x)$. Margin to class j : $m_j(x) = (w_j - w_b)^\top x + (b_j - b_b)$.

$f_j(x) = f_j(x) \gg m_j(x) = 0$. Correct if $f_j(x) > f_j(y) \forall j \neq y$, adversarial if $\exists j \neq y$ s.t. $f_j(x + \delta) < f_j(x + \delta)$. Distance to boundary: $\frac{m_j - m_y}{\|w_j - w_y\|_2}$.

Adversarial attacks for NNs: Approximate boundary by $f(x + \delta) \approx f(x) + \nabla f(x)^\top \delta$. FGSM is a one-step ℓ_1 attack: $\delta = \epsilon \text{ sign}(\nabla_x L(f(x), y))$. PGD is multi-step $\delta_{t+1} = \text{Proj}_{\delta \in \mathcal{S}}(\delta_t + \alpha \text{ sign}(g_t))$.

Distributionally Robust Optimization: $\min_u \sup_{Q \in \mathcal{P}(P)} E_Q L(f(x), y)$, where U means close. Can use upper bound on Wasserstein distance e.g.

Robust training: $\min_u E[\max_{Q \in \mathcal{P}(P)} L(f(x), y)]$. Adversarial training can be viewed as robustness to distribution shift measured by Wasserstein distance.

Interpretability: Local - explain pred for specific x , Global - explain model behaviour on avg over data.

Local: Ceteris paribus (vary x_j , fix x_{-j}). Sensitivity ($\partial_{x_j} f(x)$), missing info ($f(x) - \mathbb{E}[f(X)] | X_{-j} = x_{-j}$). **Global:** Mutual info ($I(X_j; Y | X_{-j})$), Predictive util (train f w/ and w/o x_j). For log-loss predictive util \approx conditional mutual information.

SHAP attributes predictions, while SAGE attributes risk reduction.

A predicted attribute, Y target outcome, \hat{Y} prediction. Demographic Parity: $\hat{Y} \perp A$; Equalized Odds: $\hat{Y} \perp A | Y$; Equality of Opportunity: $\hat{Y} \perp A | Y = 1$.

For $W \in \mathbb{R}^{n \times d}$, $W_{i,j} \sim \mathcal{N}(0, 1)$: $E[W] = 0, E[WW^\top] = dI_n, E[W^\top W] = nI_d; E[WAW^\top] = \text{tr}(A)I_n, E[W^\top AW] = \text{tr}(A)I_d$.

For $w \sim \mathcal{N}(0, I_n)$: $E[wu^\top] = I_n, E[w^\top w] = n, E[w^\top Aw] = \text{tr}(A)$.

def forward(self, x):
 B, T, C = x.size() # batch size, seqLen, EmbeddingSize
 qkv = self.qkv_att(x) # (B, T, 3*C)
 q, k, v = qkv.split(self.n_embd, dim=2)
 q = q.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
 k = k.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
 v = v.view(B, T, self.n_head, C // self.n_head).transpose(1, 2)
 att = (q @ k.transpose(-1, -2)) * self.attn_bias # (B, n_head, T, T)
 att = att * 1.0 / math.sqrt(C // self.n_head)
 att = att.softmax(dim=-1)
 y = att @ v # (B, nh, T, hs)
 y = y.transpose(1, 2).reshape(B, T, C)
 y = self.c_proj(y)
 return y

torch.cat concatenates along an existing dimension
torch.stack creates a new dimension

import torch

a = torch.tensor([1, 2, 3])

b = torch.tensor([4, 5, 6])

cat: joins along existing dim

torch.cat([a, b]) # shape (6,) → [1, 2, 3, 4, 5, 6]

stack: creates new dim

torch.stack([a, b]) # shape (2, 3) → [[1, 2, 3], [4, 5, 6]]